



Chapter 2 State-Space Search

COMP 6721 Introduction of AI

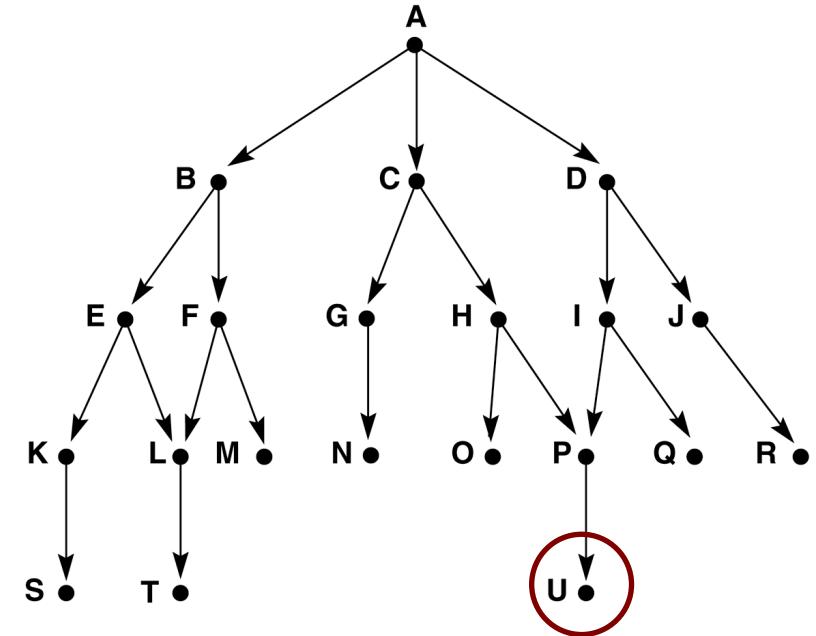
Russell & Norvig – Chapter 3 & Section 4.11

Some slides from: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Breadth-First Search Example

- ➡ BFS uses a **queue**.
- ➡ Assume U is goal State.

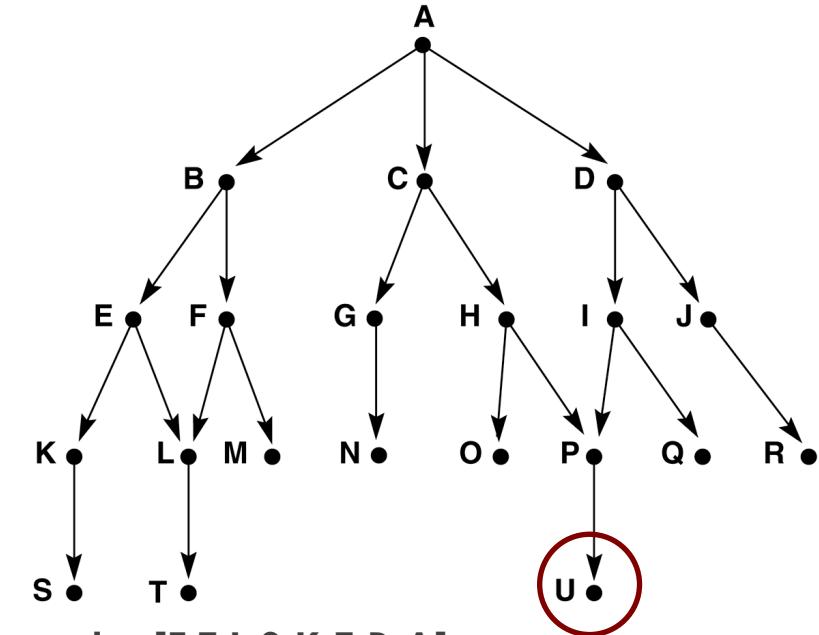
1. open = [A-null] closed = []
2. open = [B-A C-A D-A] closed = [A]
3. open = [C-A D-A E-B F-B] closed = [B A]
4. open = [D-A E-B F-B G-C H-C] closed = [C B A]
5. open = [E-B F-B G-C H-C I-D J-D] closed = [D C B A]
6. open = [F-B G-C H-C I-D J-D K-E L-E] closed = [E D C B A]
7. open = [G-C H-C I-D J-D K-E L-E M-F] **as L is already on open** closed = [F E D C B A]
8. and so on until either U is found or open = []



Depth-First Search Example

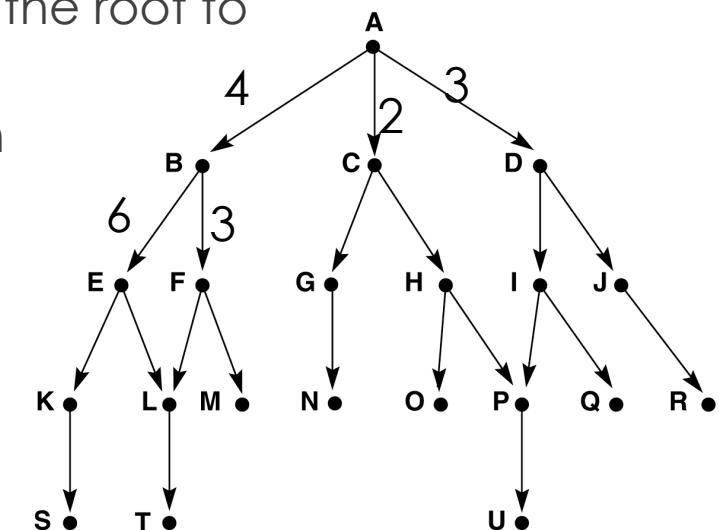
- ➡ DFS uses a **stack**.
- ➡ Assume U is goal State.

1. open = [A-null] closed = []
2. open = [B-A C-A D-A] closed = [A]
3. open = [E-B F-B C-A D-A] closed = [B A]
4. open = [K-E L-E F-B C-A D-A] closed = [E B A]
5. open = [S-K L-E F-B C-A D-A] closed = [K E B A]
6. open = [L-E F-B C-A D-A] closed = [S K E B A]
7. open = [T-L F-B C-A D-A] closed = [L S K E B A]
8. open = [F-B C-A D-A] closed = [T L S K E B A]
9. open = [M-F C-A D-A] **as L is already on closed** closed = [FTLSKEBA]
10. open = [C-A D-A] closed = [MFTLSKEBA]
11. open = [G-C H-C D-A] closed = [CMFTLSKEBA]
12. open = [N-G, H-C, D-A] closed = [GCMFTLSKEBA]
13. open = [O-H, P-H, D-A] closed = [HNGCMFTLSKEBA]
14. open = [P-H, D-A] closed = [HNGCMFTLSKEBA]
15. open = [D-A] closed = [UPOHNCGMFTLSKEBA]



Uniform Cost Search

- ▶ Breadth First Search
 - ▶ uses a priority queue sorted using the depth of the nodes from the root
 - ▶ guarantees to find the shortest solution path
- ▶ But what if all edges/moves do not have the same cost?
- ▶ Uniform Cost Search
 - ▶ uses a priority queue sorted using the cost from the root to node n – later called $g(n)$
 - ▶ guarantees to find the lowest cost solution path



Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ ***Uninformed Search***
 - ▶ **Breadth-first and Depth-first**
 - ▶ **Depth-limit Search**
 - ▶ **Iterative Deepening**
 - ▶ **Uniform Cost**
 - ▶ **Informed Search**
 - ▶ Hill Climbing
 - ▶ Best – Frist
 - ▶ Design Heuristics
 - ▶ A*

Informed Search (Heuristic Search)

- ▶ Most of the time, it is not feasible to do an exhaustive search, search space is too large
- ▶ so far, all search algorithms have been uninformed (general search)
- ▶ so need an ***informed/heuristic search***
- ▶ Idea:
 - ▶ choose "best" next node to expand
 - ▶ according to a selection function (heuristic)
- ▶ But: heuristic might fail

Heuristic – Heureka!

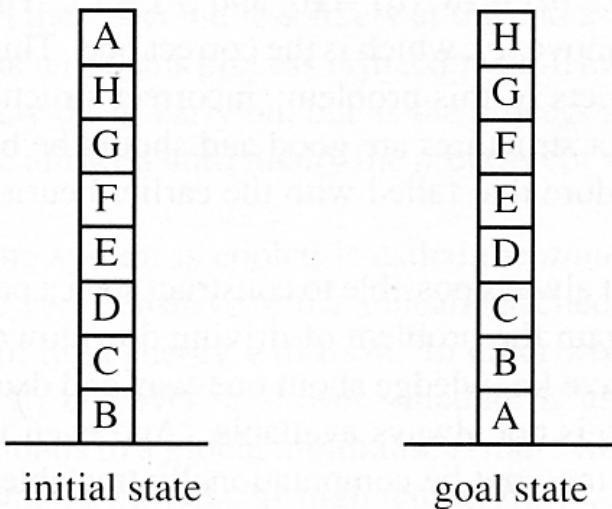


- ▶ Heuristic:
 - ▶ a rule of thumb, a good bet
 - ▶ but has no guarantee to be correct whatsoever!
- ▶ Heuristic search:
 - ▶ A technique that improves the efficiency of search, possibly sacrificing on completeness
 - ▶ Focus on paths that seem most promising according to some function
 - ▶ Need an evaluation function (heuristic function) to score a node in the search tree
- ▶ Heuristic function $h(n)$ = an **approximation** to a function that gives the true evaluation of the node n

Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ **Uninformed Search**
 - ▶ *Breadth-first and Depth-first*
 - ▶ *Depth-limit Search*
 - ▶ *Iterative Deepening*
 - ▶ *Uniform Cost*
 - ▶ Informed Search
 - ▶ **Hill Climbing**
 - ▶ Best – First
 - ▶ Design Heuristics
 - ▶ A*

Example: Hill Climbing with Blocks World

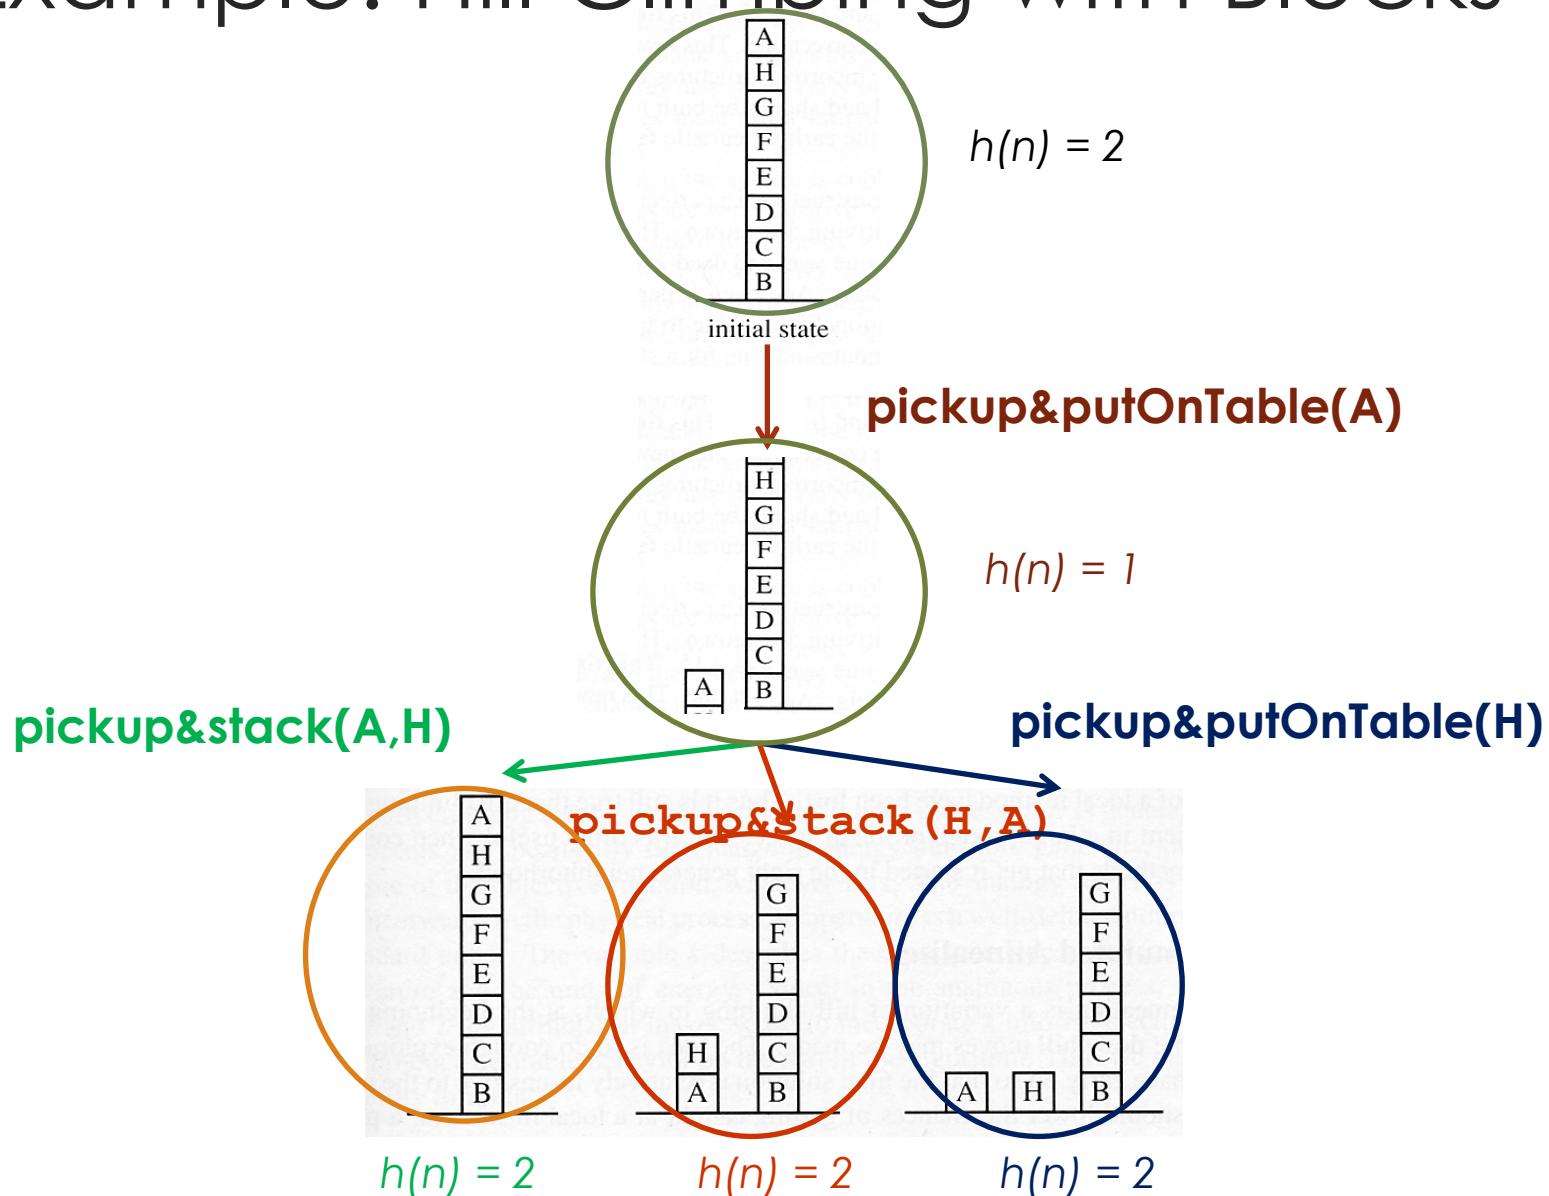


- Operators:
 - pickup&putOnTable(Block)
 - pickup&stack(Block1,Block2)

► Heuristic:

- - 1pt if a block is sitting where it is supposed to sit
- +1pt if a block is NOT sitting where it is supposed to sit
- so lower $h(n)$ is better
 - $h(\text{initial}) = 2$
 - $h(\text{goal}) = 0$

Example: Hill Climbing with Blocks World



Hill Climbing

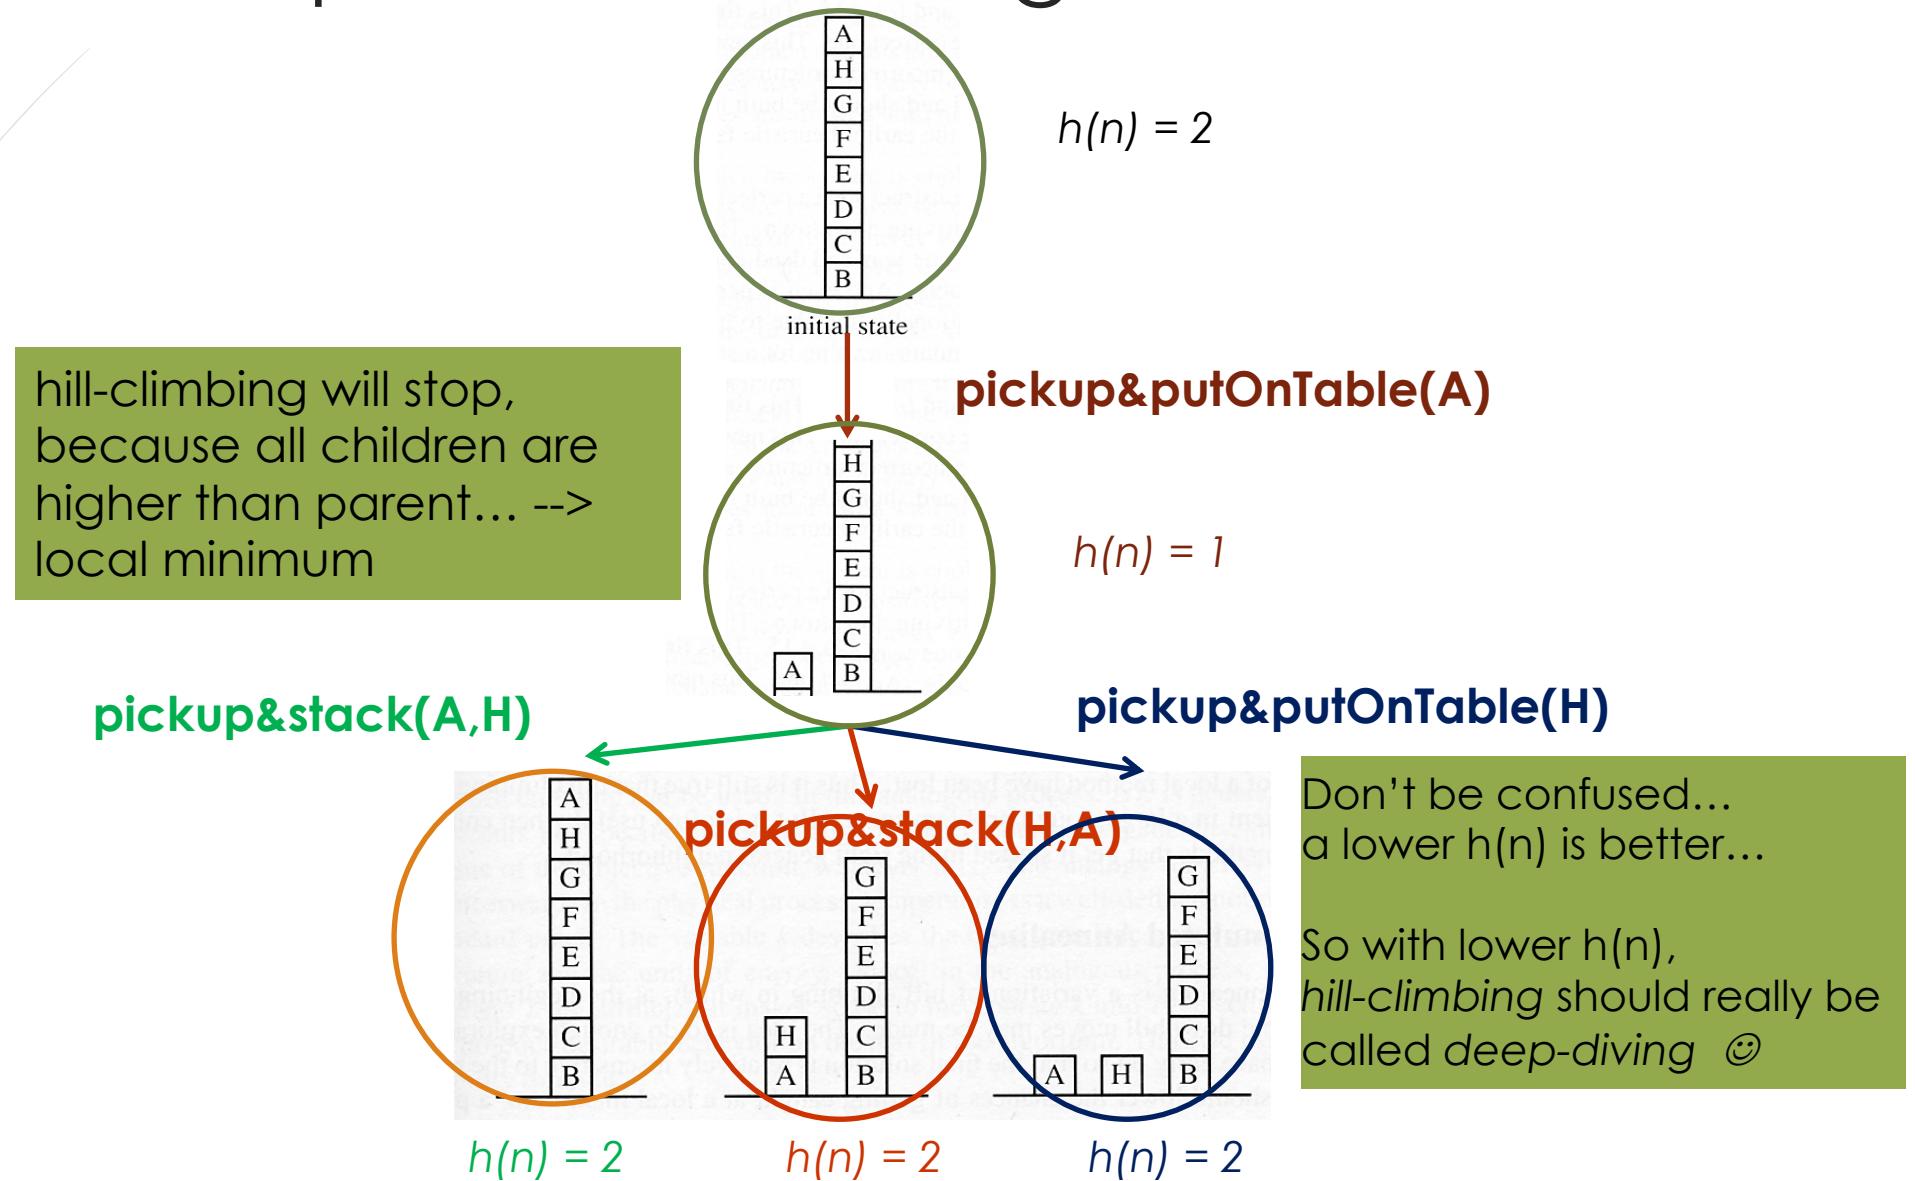
- ▶ General hill climbing strategy:
 - ▶ as soon as you find a position that is better than the current one, select it.
 - ▶ Does not maintain a list of next nodes to visit (an open list)
 - ▶ Similar to climbing a mountain in the fog with amnesia ... always go higher than where you are now, but never go back...
- ▶ Steepest ascent hill climbing:
 - ▶ instead of moving to the first position that is better than the current one
 - ▶ pick the best position out of all the next possible moves



Steepest Ascent Hill Climbing

```
currentNode = startNode;  
loop do  
    L = CHILDREN(currentNode);  
    nextEval = -INFINITY;  
    nextNode = NULL;  
  
    for all c in L  
        if (HEURISTIC-VALUE(c) < nextEval) // lower h is better  
            nextNode = c;  
            nextEval = HEURISTIC-VALUE(c);  
        if nextEval >= HEURISTIC-VALUE(currentNode)  
            return currentNode; // Return current node since no better child state exist  
    currentNode = nextNode;
```

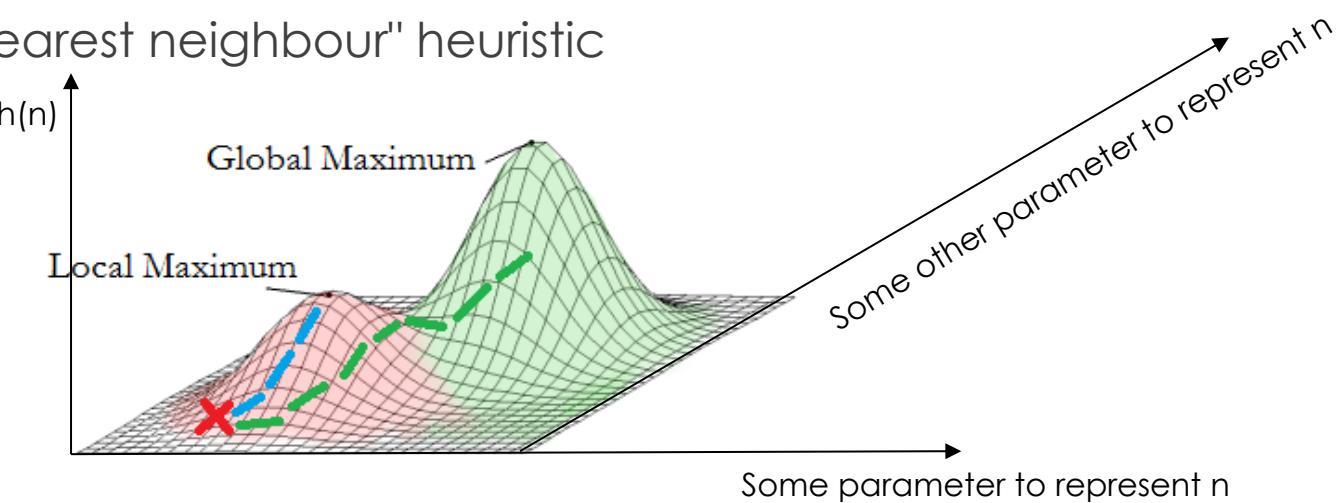
Example: Hill Climbing with Blocks World



Problems with Hill Climbing

► Foothills (or local maxima)

- reached a local maximum, not the global maximum
- a state that is better than all its neighbors but is not better than some other states farther away.
- at a local maximum, all moves appear to make things worse.
- ex: 8-puzzle: we may need to move tiles temporarily out of goal position in order to place another tile in goal position
- ex: "nearest neighbour" heuristic



Problems with Hill Climbing

► Plateau

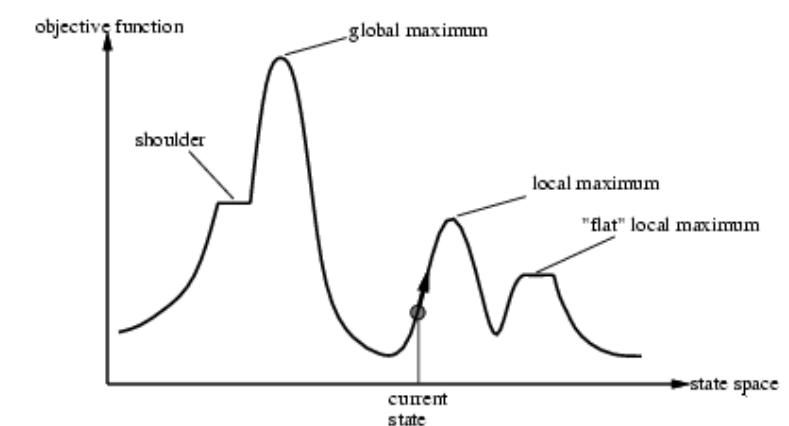
- a flat area of the search space in which the next states have the same value.
- it is not possible to determine the best direction in which to move by making local comparisons.



Solutions to Hill Climbing

- ▶ Random-restart hill-climbing
 - ▶ random initial states are generated
 - ▶ run each until it halts or makes no significant progress.
 - ▶ the best result is then chosen.

- ▶ keep going even if the best successor has the same value as current node
 - ▶ works well on a "shoulder"
 - ▶ but could lead to infinite loop on a plateau



Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ **Uninformed Search**
 - ▶ *Breadth-first and Depth-first*
 - ▶ *Depth-limit Search*
 - ▶ *Iterative Deepening*
 - ▶ *Uniform Cost*
 - ▶ Informed Search
 - ▶ *Hill Climbing*
 - ▶ *Best – Frist*
 - ▶ Design Heuristics
 - ▶ A*

Best – First Search

- ▶ Problem with hill-climbing:
 - ▶ one move is selected and all others are forgotten.
- ▶ Solution to hill-climbing:
 - ▶ use "open" as a priority queue
 - ▶ this is called best-first search
- ▶ Best-first search:
 - ▶ Insert nodes in open list so that the nodes are sorted in best (ascending/descending) $h(n)$
 - ▶ Always choose the next node to visit to be the one with the best $h(n)$ -- regardless of where it is in the search space

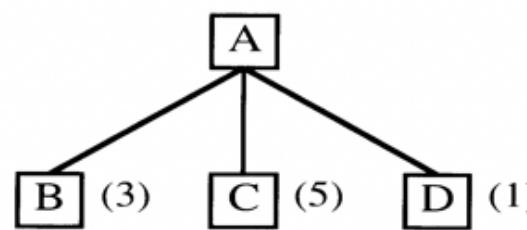
Best – First Search Example

Lower $h(n)$ is better

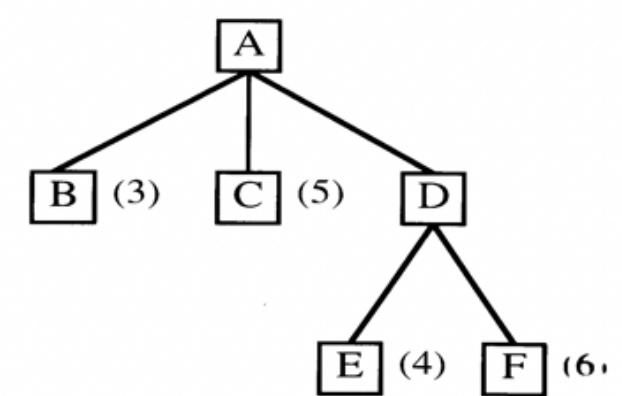
Step 1



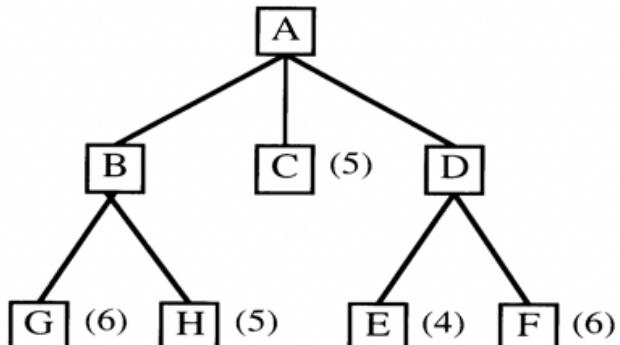
Step 2



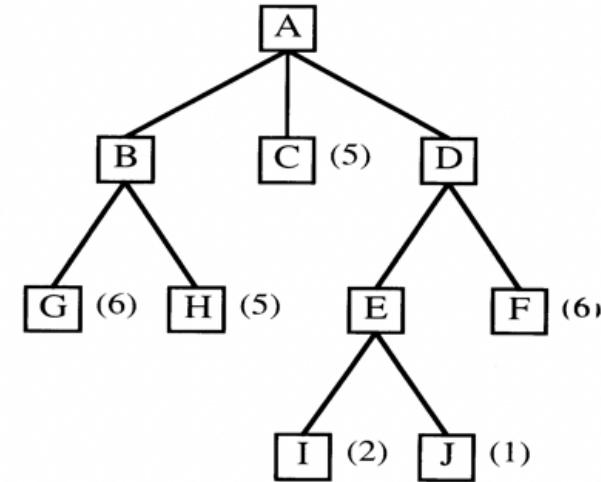
Step 3



Step 4



Step 5



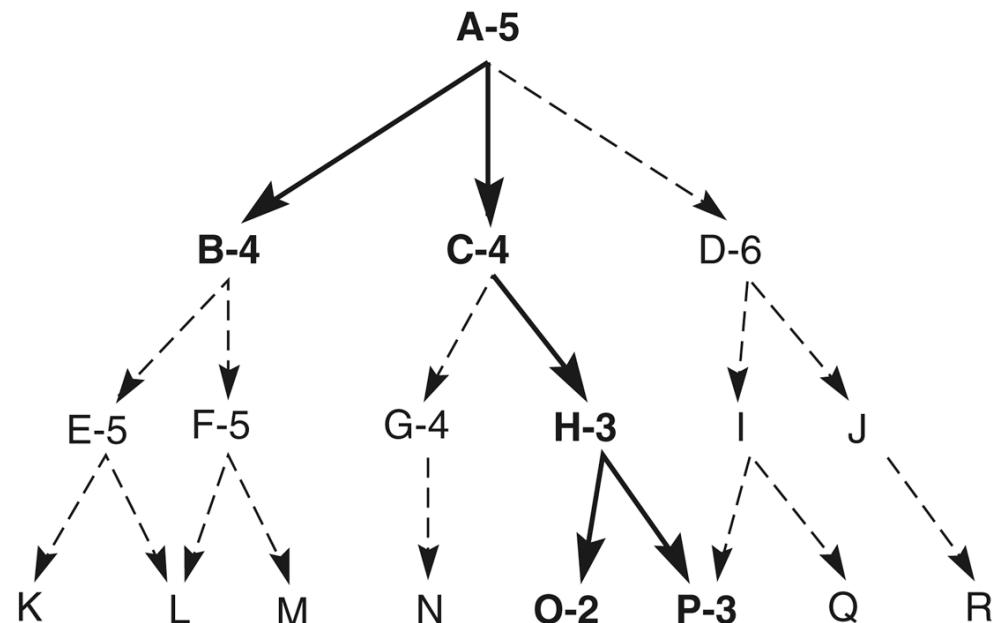
Best – First Search

- ▶ If you have a good evaluation function, best-first can find the solution very quickly
- ▶ The first solution may not be the best, but there is a good chance of finding it quickly
- ▶ It is an exhaustive search ...
 - ▶ will eventually try all possible paths

Best – First Search Example

1. open = [A-null-5] closed = []
2. open = [B-A-4 C-A-4 D-A-6] (arbitrary choice) closed [A]
3. open = [C-A-4 E-B-5 F-B-5 D-A-6] closed = [B A]
4. open = [H-C-3 G-C-4 E-B-5 F-B-5 D-A-6] closed = [C B A]
5. open = [P-H-0 O-H-2 G-C-4 E-B-5 F-B-5 D-A-6] closed = [H C B A]
6. goal P found
7. solution path: A C H P

Lower $h(n)$ is better



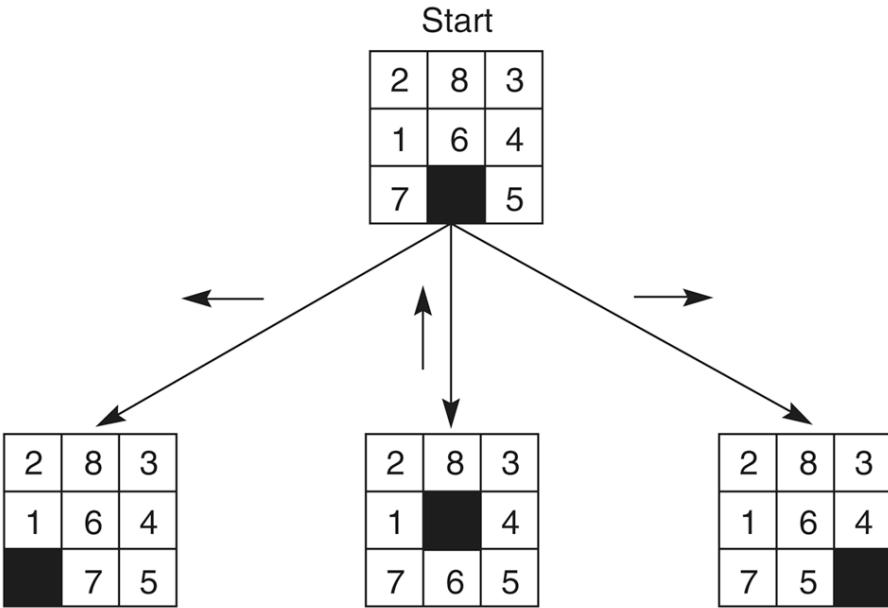
Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ **Uninformed Search**
 - ▶ *Breadth-first and Depth-first*
 - ▶ *Depth-limit Search*
 - ▶ *Iterative Deepening*
 - ▶ *Uniform Cost*
 - ▶ Informed Search
 - ▶ *Hill Climbing*
 - ▶ *Best – Frist*
 - ▶ **Design Heuristics**
 - ▶ A*

Design Heuristics

- ▶ Heuristic evaluation functions are highly dependent on the search domain
- ▶ In general: the more informed a heuristic is, the better the search performance
- ▶ Bad heuristics lead to frequent backtracking
- ▶ So how do we design a “good” heuristic?

Example: 8 – Puzzle Heuristic 1



- h_1 : Simplest heuristic
 - Hamming distance : count **number of tiles out of place** when compared with goal

5		8
4	2	1
7	3	6

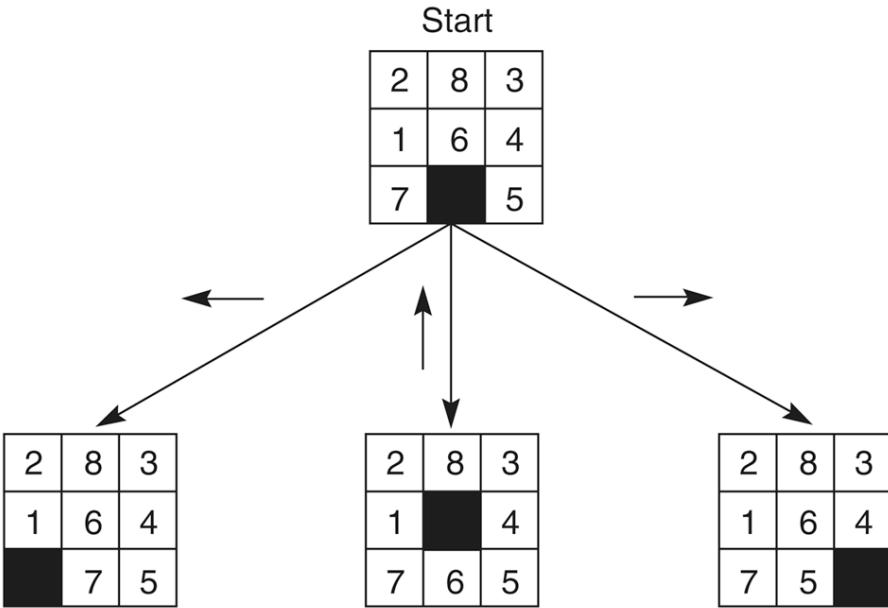
STATE n

1	2	3
4	5	6
7	8	

Goal state

- $h_1(n) = 6$
 - does not consider the distance tiles have to be moved

Example: 8 – Puzzle Heuristic 2



- h_2 : Better heuristic
- Manhattan distance: sum up all the ***distances*** by which tiles ***are out of place***

5		8
4	2	1
7	3	6

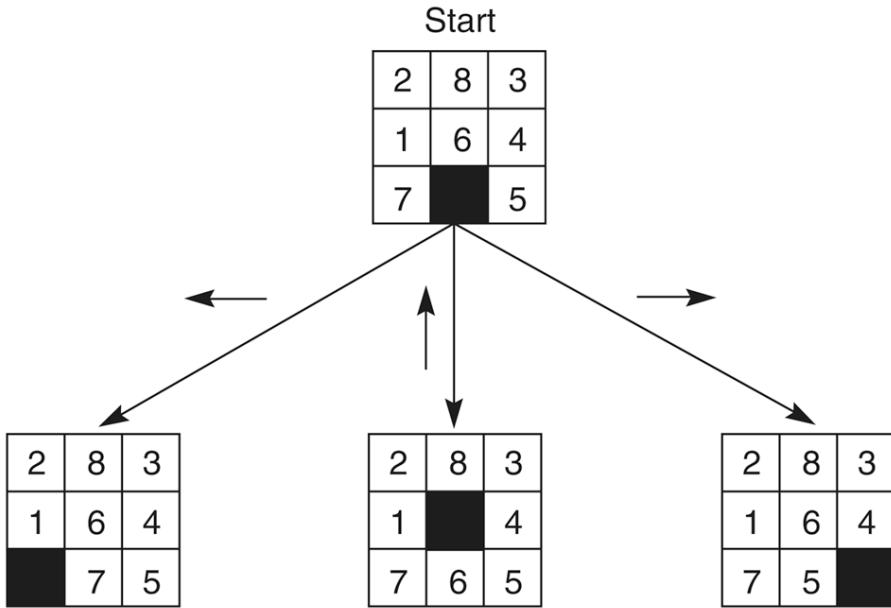
STATE n

1	2	3
4	5	6
7	8	

Goal state

$$\begin{aligned}
 \rightarrow h_2(n) &= 2+3+0+1+3+0+3+1 \\
 &= 13
 \end{aligned}$$

Example: 8 – Puzzle Heuristic 3



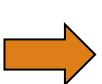
- h_3 : Even Better
 - sum of permutation inversions
 - See next slide...

Example: 8 – Puzzle Heuristic 3

- $h_3(N) = \text{sum of permutation inversions}$

5			8
4	2	1	
7	3	6	

STATE n



5		8	4	2	1	7	3	6
---	--	---	---	---	---	---	---	---

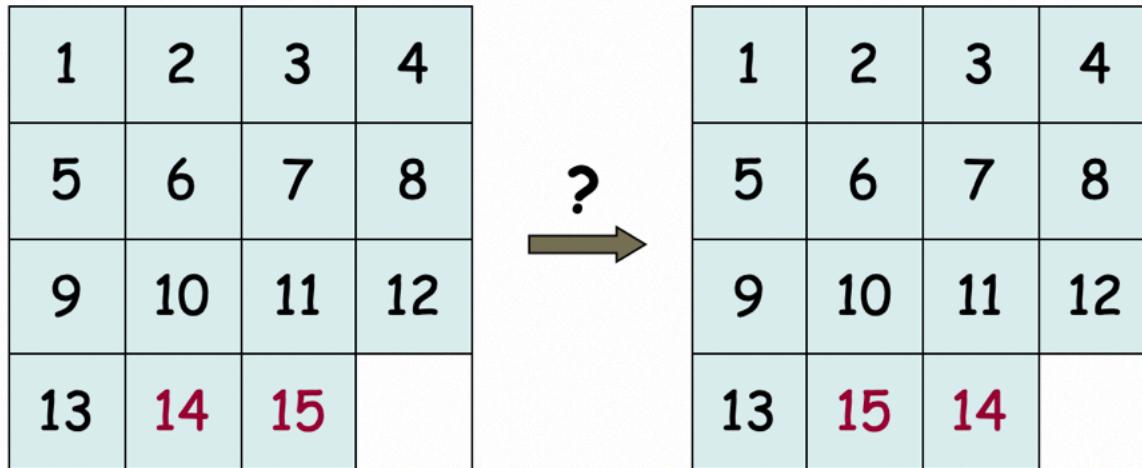
- For each numbered tile, count how many tiles on its right should be on its left in the goal state.
- $$\begin{aligned} h_3(n) &= n_5 + n_8 + n_4 + n_2 + n_1 + n_7 + n_3 + n_6 \\ &= 4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 \\ &= 16 \end{aligned}$$
- If this number is even, the puzzle is solvable.
- If this number is odd, the puzzle is not solvable.

1	2	3
4	5	6
7	8	

Goal state

Example: 15 - Puzzle

- ▶ Sam Loyd even offered \$1,000 of his own money to the first person who would solve the following problem:



$$\begin{aligned}h_3(n) &= n_1 + n_2 + n_3 + n_4 + \dots + n_{13} + n_{14} + n_{15} \\&= 0 + 0 + 0 + 0 + \dots + 0 + 1 + 0 \\&= 1\end{aligned}$$

→ the puzzle is not solvable.

Heuristics for 8 - Puzzle

5		8
4	2	1
7	3	6

STATE n

1	2	3
4	5	6
7	8	

Goal state

- $h_1(n)$ = misplaced numbered tiles (Hamming Distance)
= 6
- $h_2(n)$ = Manhattan distance
 $= 2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$
- $h_3(n)$ = sum of permutation inversions
 $= n_5 + n_8 + n_4 + n_2 + n_1 + n_7 + n_3 + n_6$
 $= 4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 = 16$

Evaluation of Heuristics

1. Admissibility:
 - ▶ “optimistic”
 - ▶ never overestimates the cost of reaching the goal
 - ▶ guarantees to find the shortest solution path to the goal (if it exists)
 - ▶ Note: does not guarantee to find the shortest search path.
 - ▶ e.g.: breadth-first is admissible -- it uses $f(n) = g(n) + 0$
2. Monotonicity:
 - ▶ “local admissibility”
 - ▶ guarantees to find the shortest path to each state encountered in the search
3. Informedness:
 - ▶ measure for the “quality” of a heuristic
 - ▶ the more informed, the better

Admissibility

- ▶ Evaluation function $f(n) = g(n) + h(n)$ for node n :
 - ▶ $g(n)$ current cost from start to node n
 - ▶ $h(n)$ estimate of the cost from n to goal
 - ▶ $f(n)$ estimate of the total cost of the solution path passing through n
- ▶ Now consider $f^*(n) = g^*(n) + h^*(n)$:
 - ▶ $g^*(n)$ cost of shortest path from start to node n
 - ▶ $h^*(n)$ actual cost of shortest path from n to goal
 - ▶ $f^*(n)$ actual cost of shortest path from start to goal through n
- ▶ In order for a heuristic to be admissible to the search problem, the estimated cost must always be lower than or equal to the actual cost of reaching the goal state.

Example: 8 - Puzzle

5		
4	2	1
7	3	6

n

1	2	3
4	5	6
7	8	

goal

- ▶ $h_1(n)$ = Hamming distance = number of misplaced tiles = 6
--> admissible
- ▶ $h_2(n)$ = Manhattan distance = 13
--> admissible

Monotonicity (Consistent)

- ▶ An admissible heuristics may temporarily reach non-goal states along a suboptimal path
- ▶ A heuristic is monotonic if it always finds the optimal path to each state the 1st time it is encountered !
- ▶ h is monotonic if for every node n and every successor n' of n :
 - ▶ $h(n) \leq c(n,n') + h(n')$
 - ▶ i.e. $f(n)$ is non-decreasing along any path

Informedness

- ▶ Intuition: number of misplaced tiles is less informed than Manhattan distance
- ▶ For two admissible heuristics h_1 and h_2
 - ▶ if $h_1(n) \leq h_2(n)$, for all states n
 - ▶ then h_2 is more informed than h_1
 - ▶ $h_1(n) \leq h_2(n) \leq h^*(n)$
 - ▶ **$h^*(n)$** actual cost of shortest path from n to goal
- ▶ More informed heuristics search smaller space to find the solution path
- ▶ However, you need to consider the computational cost of evaluating the heuristic...
- ▶ The time spent computing heuristics must be recovered by a better search

Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ **Uninformed Search**
 - ▶ *Breadth-first and Depth-first*
 - ▶ *Depth-limit Search*
 - ▶ *Iterative Deepening*
 - ▶ *Uniform Cost*
 - ▶ Informed Search
 - ▶ *Hill Climbing*
 - ▶ *Best – Frist*
 - ▶ *Design Heuristics*
 - ▶ **A***

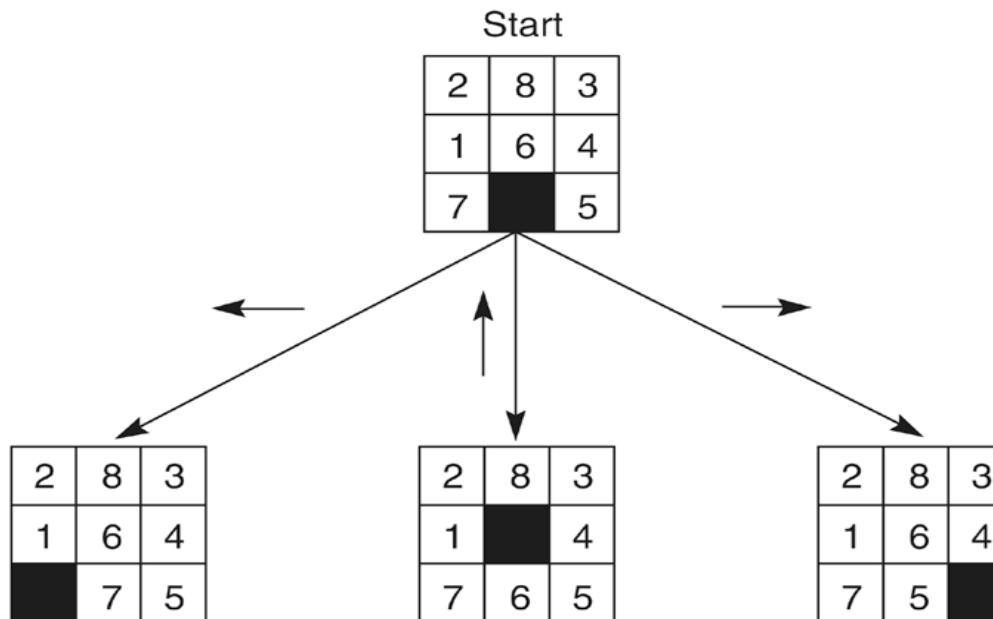
Algorithm A

- ▶ Heuristics might be wrong:
 - ▶ so search could continue down a wrong path
 - ▶ Solution:
 - ▶ Maintain depth/cost count, i.e., give preference to shorter/least expensive paths
 - ▶ Modified evaluation function f :
- $$f(n) = g(n) + h(n)$$
- ▶ $f(n)$ estimate of total cost along path through n
 - ▶ $g(n)$ actual cost of path from start to node n
 - ▶ $h(n)$ estimate of cost to reach goal from node n

Algorithm A Example: 8-Puzzle

$g(n) = 0$

$g(n) = 1$



Values of $f(n)$ for each state,

6

4

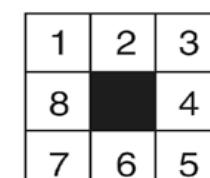
6

where:

$$f(n) = g(n) + h(n),$$

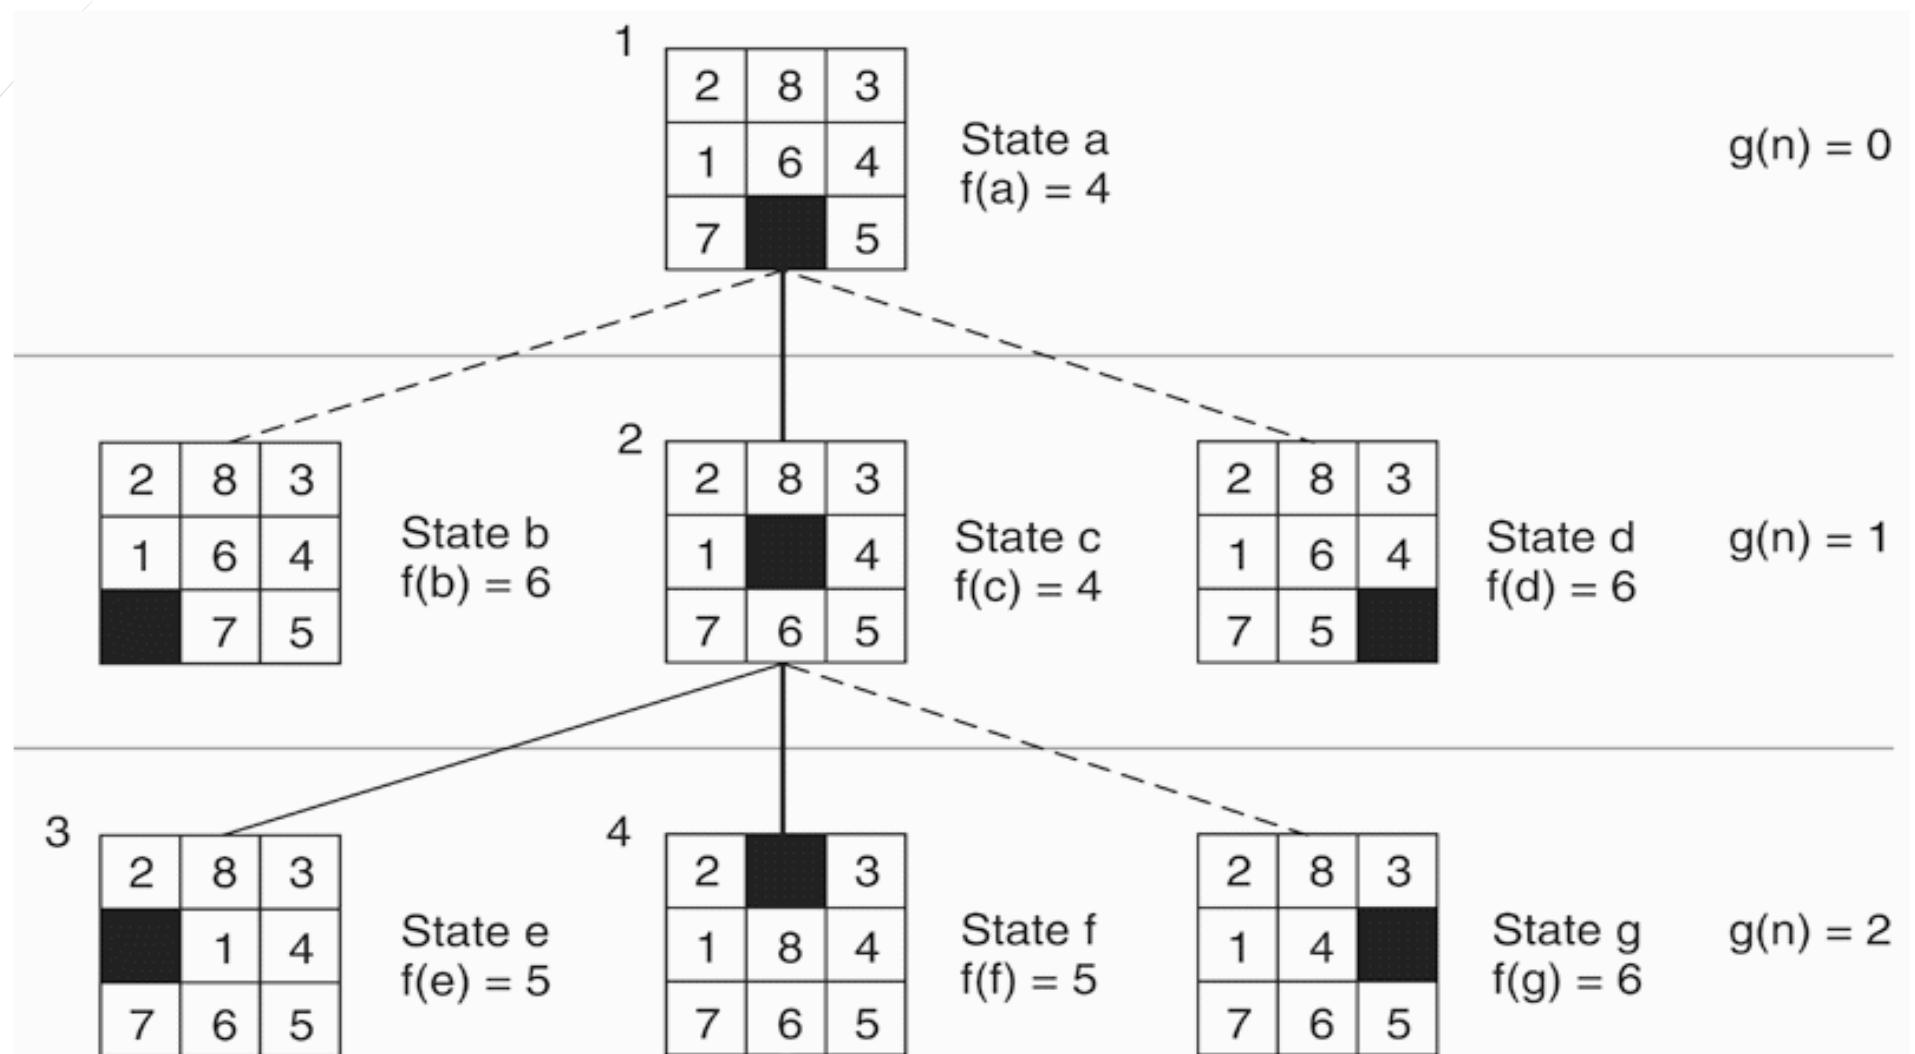
$g(n)$ = actual distance from n
to the start state, and

$h(n)$ = number of tiles out of place.

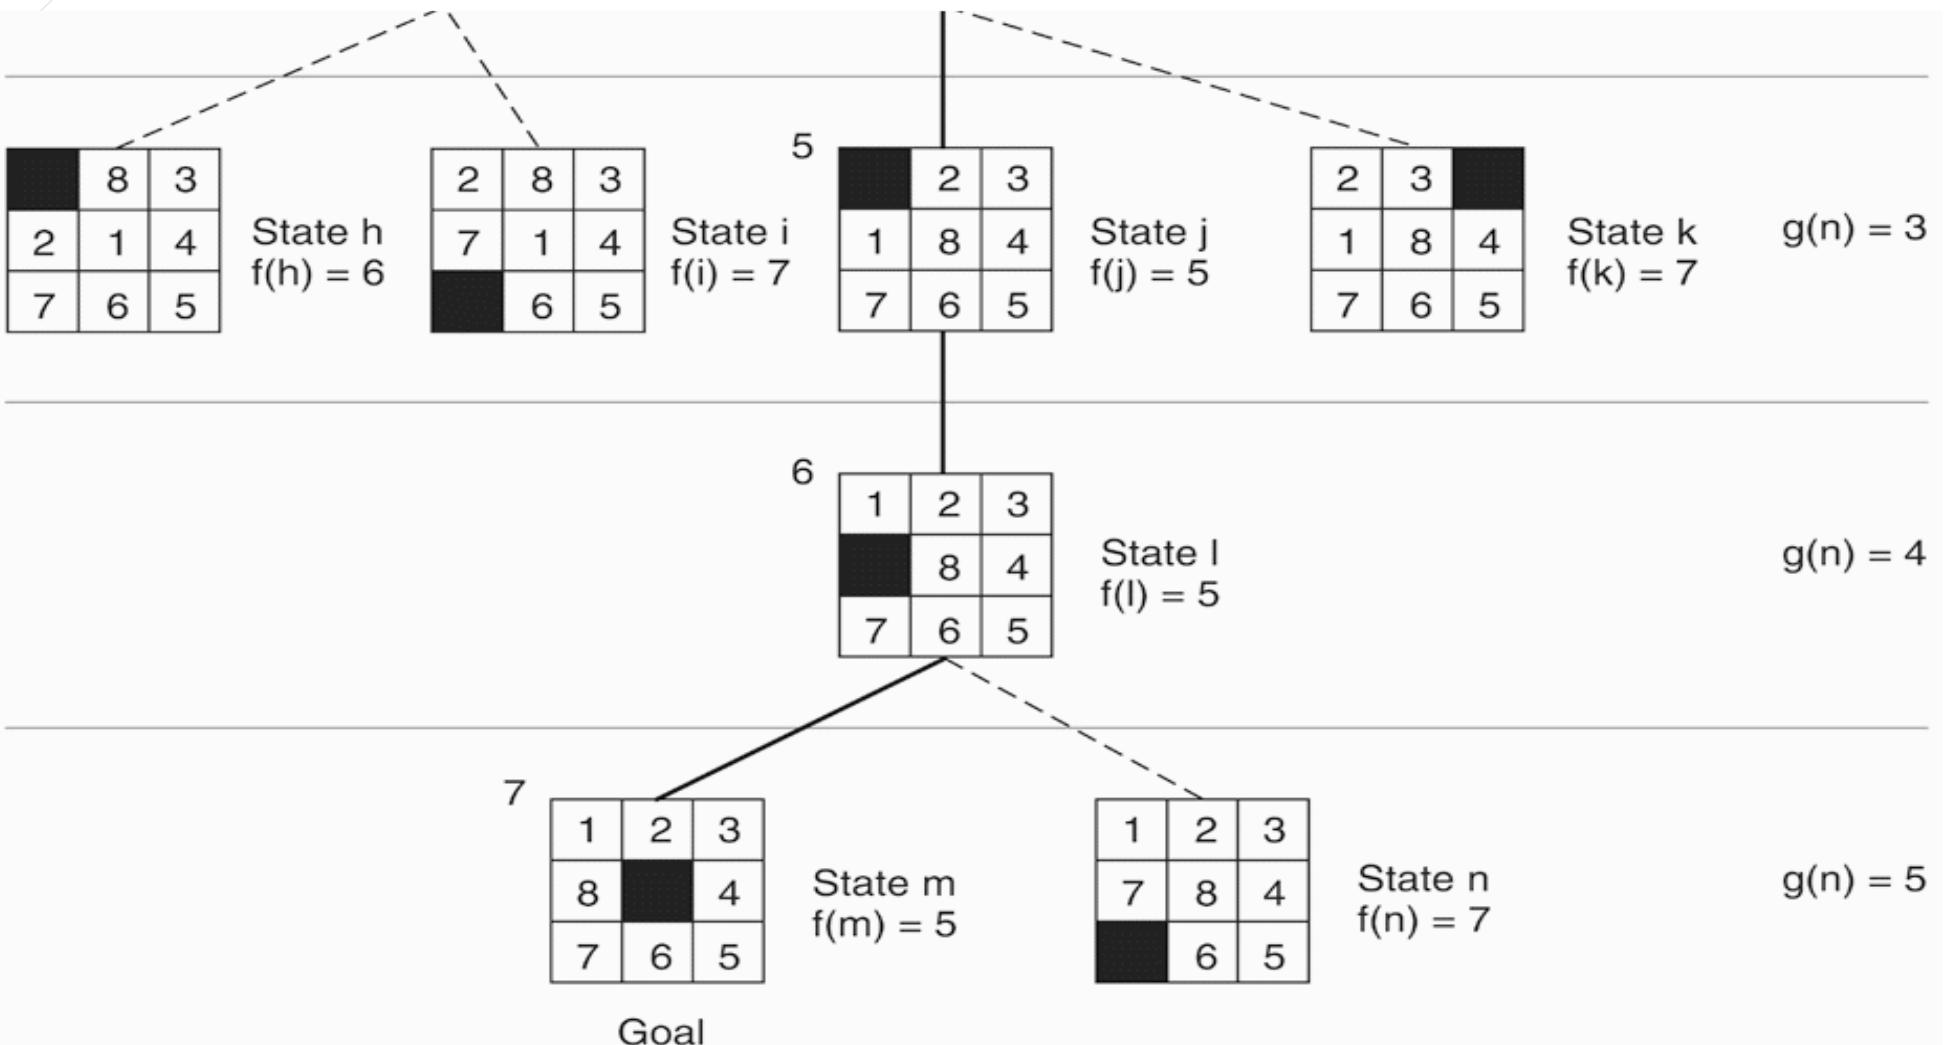


Goal

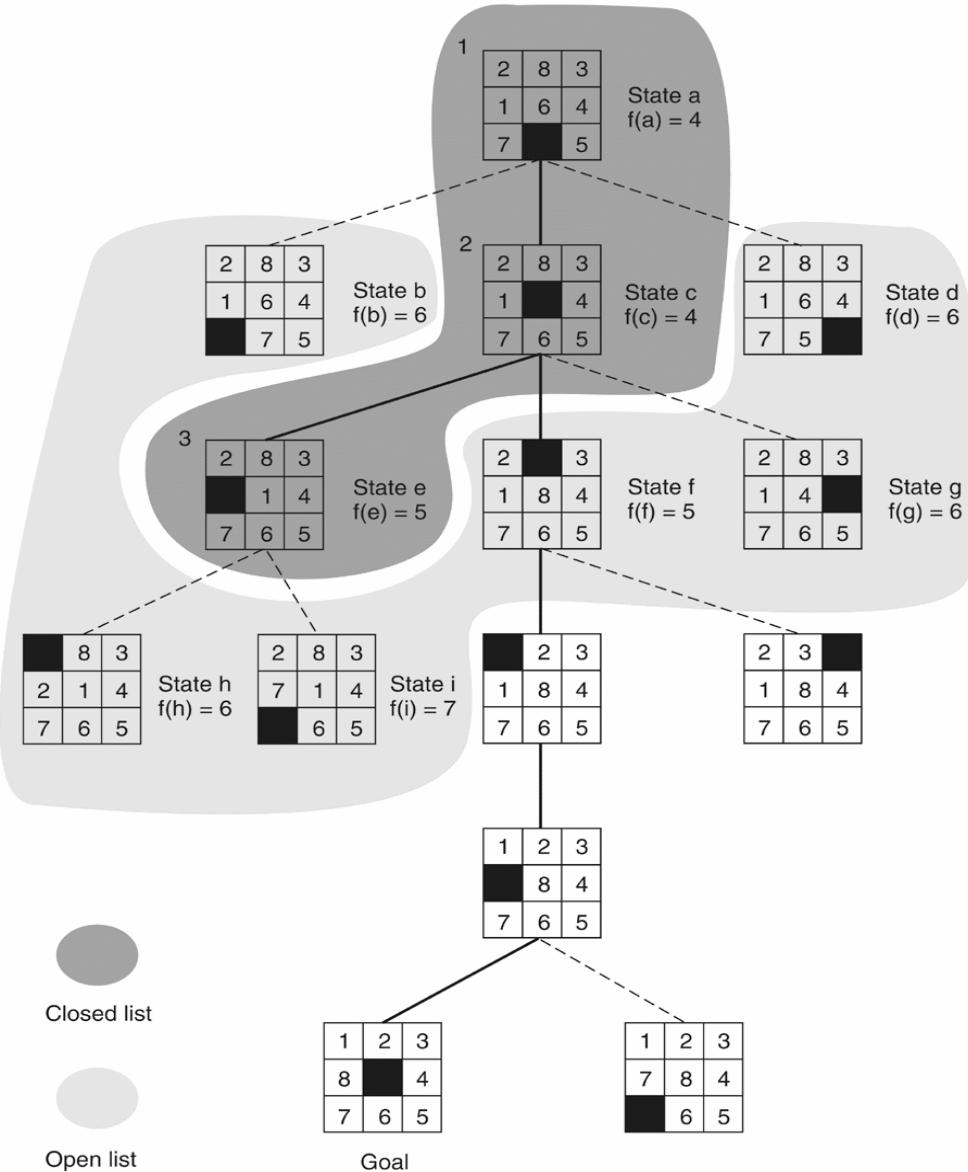
Algorithm A Example: 8-Puzzle



Algorithm A Example: 8-Puzzle



Algorithm A Example: 8-Puzzle

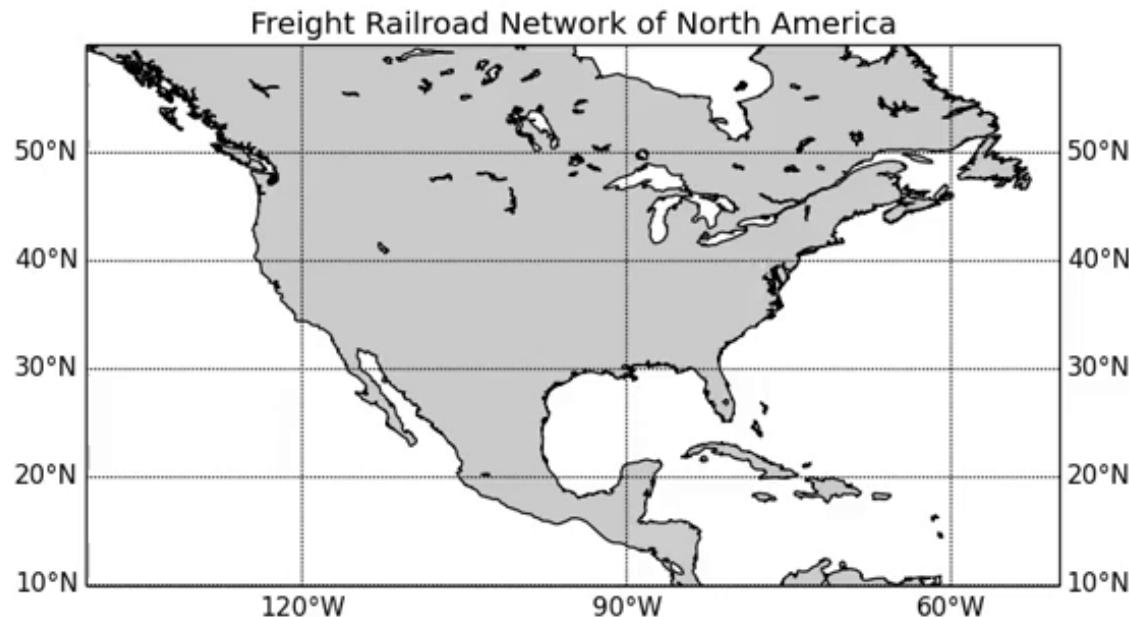


Algorithm A Example: 8-Puzzle

- ▶ if $g(n) \geq g^*(n)$ for all n
 - ▶ best-first used with such a $g(n)$ is called “algorithm A”
- ▶ if $h(n) \leq h^*(n)$ for all n
 - ▶ i.e. $h(n)$ never overestimates the true cost from n to a goal
 - ▶ algorithm A used with such an $h(n)$ is called “algorithm A*”
 - ➔ **an A* algorithm is admissible**
 - ➔ i.e it guarantees to find the lowest cost solution path from the initial state to the goal
- ▶ $g^*(n)$ cost of shortest path from start to node n
- ▶ $h^*(n)$ actual cost of shortest path from n to goal

Algorithm A* Example

The A* algorithm also has real-world applications. In this example, edges are railroads and $h(x)$ is the great-circle distance (the shortest possible distance on a sphere) to the target. The algorithm is searching for a path between Washington, D.C. and Los Angeles.



Types of State-Space Search

- ▶ **State-Space Search**
 - ▶ **Uninformed Search**
 - ▶ *Breadth-first and Depth-first*
 - ▶ *Depth-limit Search*
 - ▶ *Iterative Deepening*
 - ▶ *Uniform Cost*
 - ▶ **Informed Search**
 - ▶ *Hill Climbing*
 - ▶ *Best – Frist*
 - ▶ *Design Heuristics*
 - ▶ *A**

Summary

Search	Uses $h(n)$?	Open is a...
Breadth-first	No	Queue
Depth-first	No	Stack
Depth-limited	No	Stack
Iterative Deepening	No	Stack
Uniform Cost	No	Priority queue sorted by $g(n)$
Hill climbing	Yes	none
Best-First	Yes	Priority queue sorted by $h(n)$
Algorithm A - no constraints on $h(n)$	Yes	Priority queue sorted by $f(n)$ $f(n) = g(n) + h(n)$
Algorithm A* - same as A, but $h(n)$ must be admissible	Yes	Priority queue sorted by $f(n)$ $f(n) = g(n) + h(n)$

The End

