



Chapter 3 Adversarial Search

COMP 6721 Introduction of AI

Russell & Norvig – Section 5.1 & 5.4

Some slides from: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Adversarial Search

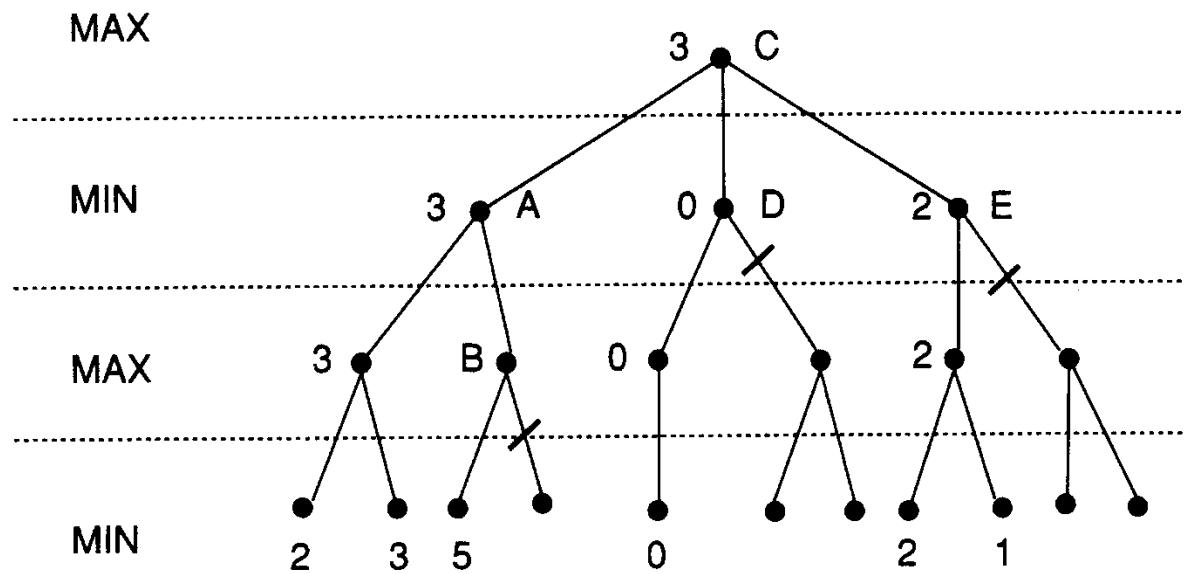
- ▶ **Minimax Search**
- ▶ Alpha – beta pruning

Alpha-Beta Pruning

- ▶ Optimization over minimax, that:
 - ▶ ignores (cuts off, prunes) branches of the tree that cannot possibly lead to a better solution
 - ▶ reduces branching factor
 - ▶ allows deeper search with same effort

Alpha-Beta Pruning Example 1

- With minimax, we look at all possible nodes at the n-ply depth
- With a-β pruning, we ignore branches that could not possibly contribute to the final decision



B will be ≥ 5
 So we can ignore B's right branch, because A must be 3

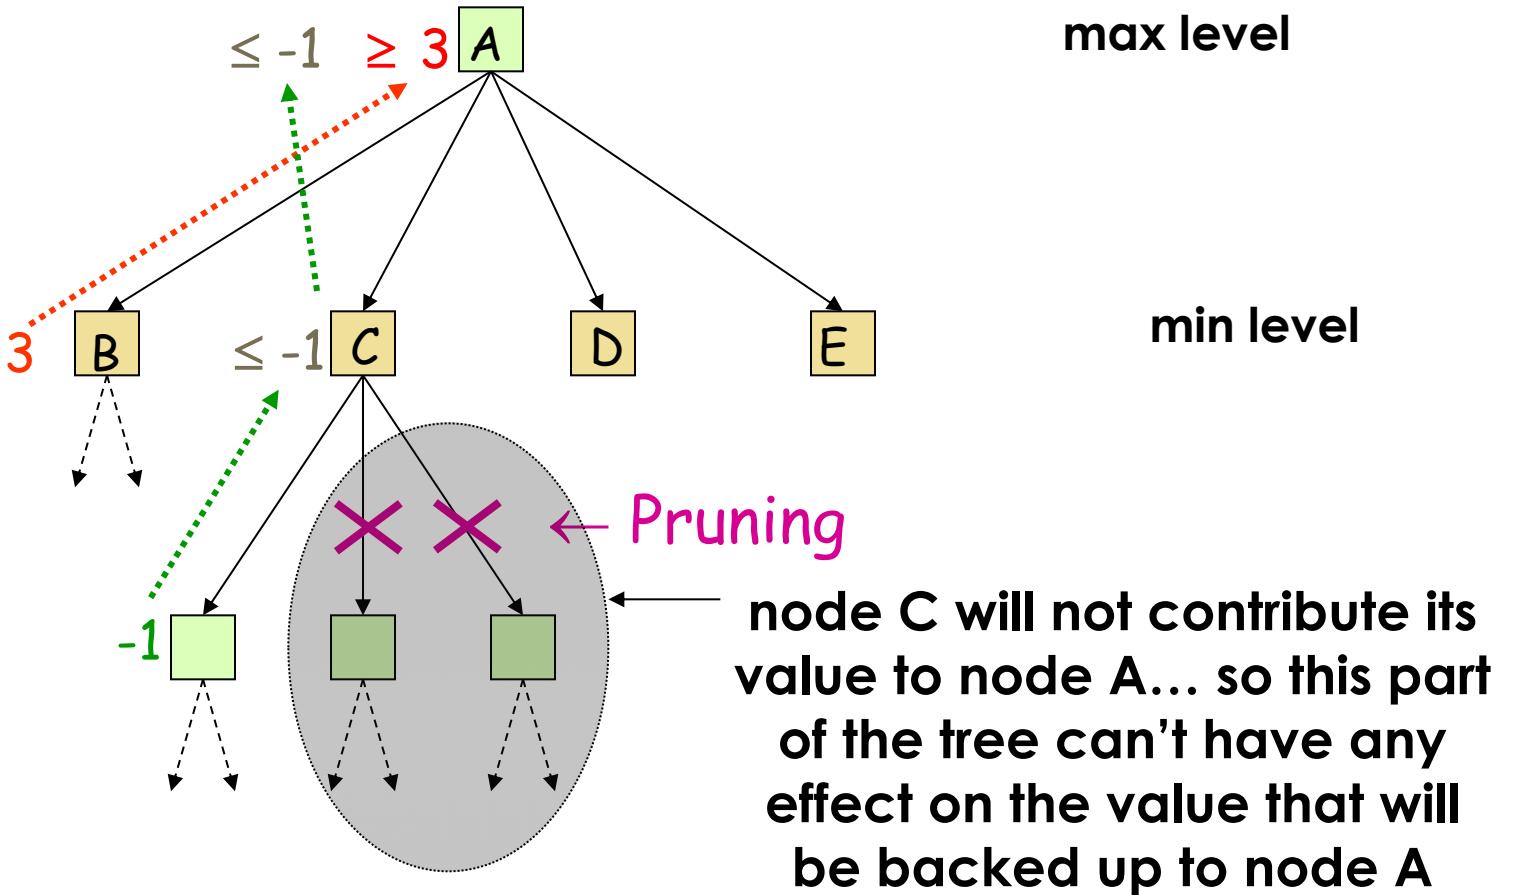
D will be ≤ 0
 But C will be ≥ 3

So we can ignore D's right branch

E will be ≤ 2 .
 So we can ignore E's right branch

Because C will be 3.

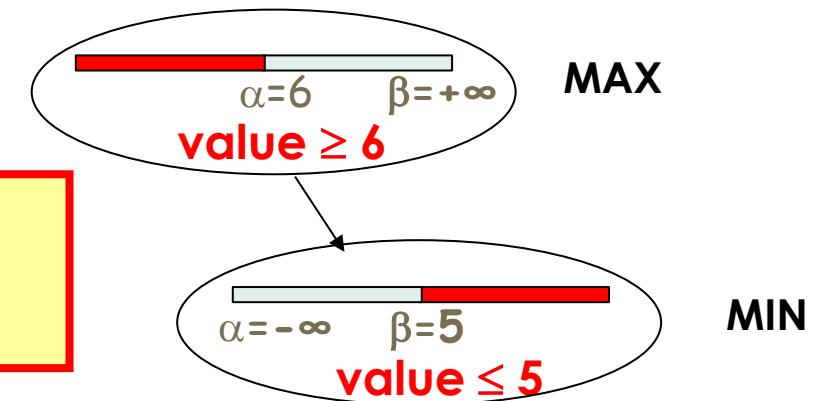
A Closer Look



Alpha-Beta Pruning Algorithm

- ▶ α : lower bound on the final backed-up value.
- ▶ β : upper bound on the final backed-up value.
- ▶ **Alpha pruning:**
 - ▶ eg. if MAX node's $\alpha = 6$, then the search can prune branches from a MIN descendant that has a $\beta \leq 6$.
 - ▶ if child $\beta \leq$ ancestor $\alpha \rightarrow$ prune

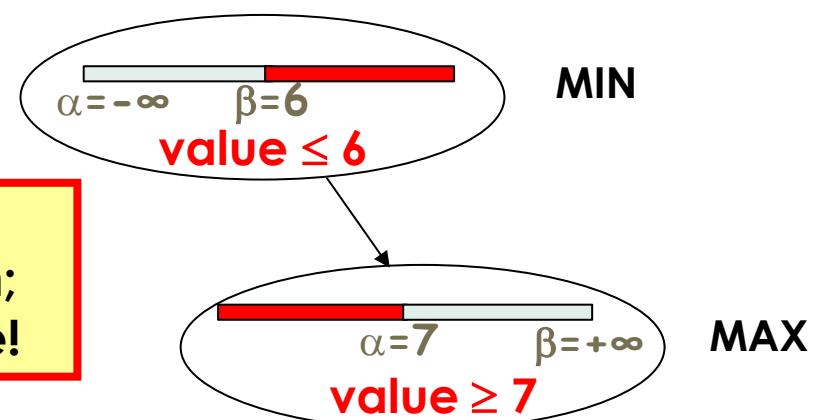
incompatible...
so stop searching the right branch;
the value cannot come from there!



Alpha-Beta Pruning Algorithm

- ▶ α : lower bound on the final backed-up value.
- ▶ β : upper bound on the final backed-up value.
- ▶ **Beta pruning:**
 - ▶ eg. if a MIN node's $\beta = 6$, then the search can prune branches from a MAX descendant that has an $\alpha >= 6$.
 - ▶ if ancestor $\beta \leq$ child $\alpha \rightarrow$ prune

incompatible...
so stop searching the right branch;
the value cannot come from there!



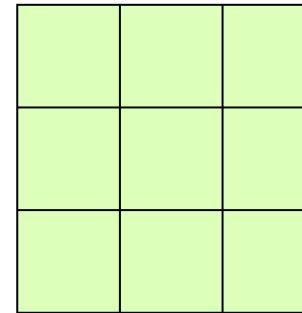
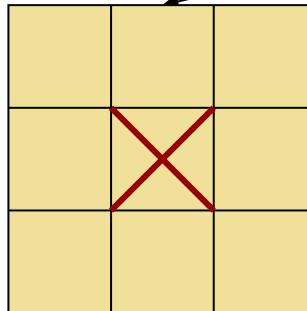
Alpha-Beta Pruning Algorithm

```
01 function alphabeta(node, depth, a, β, maximizingPlayer)
02     if depth = 0 or node is a terminal node
03         return the heuristic value of node
04     if maximizingPlayer
05         v := -∞
06         for each child of node
07             v := max(v, alphabeta(child, depth - 1, a, β, FALSE))
08             a := max(a, v)
09             if β ≤ a
10                 break (* β cut-off *)
11         return v
12     else
13         v := ∞
14         for each child of node
15             v := min(v, alphabeta(child, depth - 1, a, β, TRUE))
16             β := min(β, v)
17             if β ≤ a
18                 break (* a cut-off *)
19         return v
```

Initial call:

alphabeta(origin, depth, -∞, +∞, TRUE)

Example: tic-tac-toe

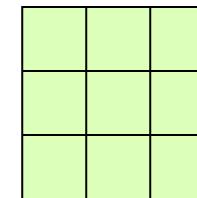


max level

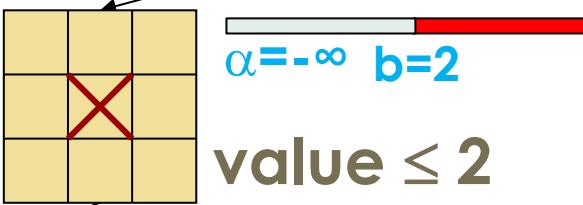
min level

Example: tic-tac-toe

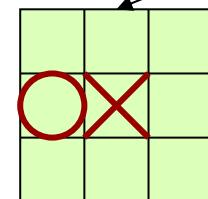
max level



min level

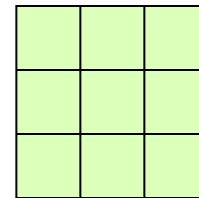


$e(n) = 2$

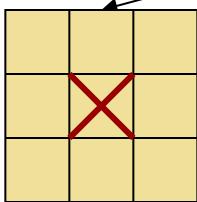


Example: tic-tac-toe

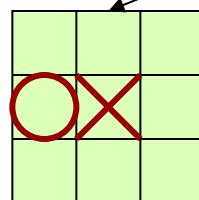
max level



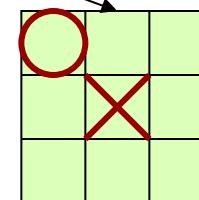
min level



$\alpha = -\infty$ $b = 2$ 1
value ≤ 2 1



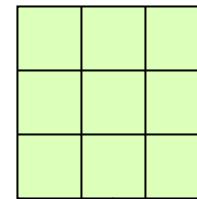
$e(n) = 2$



$e(n) = 1$

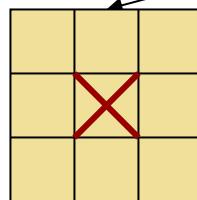
Example: tic-tac-toe

max level

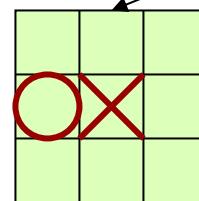


$\alpha=1$ $b=+\infty$
value ≥ 1

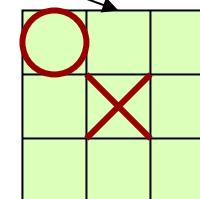
min level



value = 1



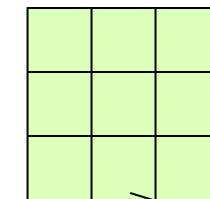
$e(n) = 2$



$e(n) = 1$

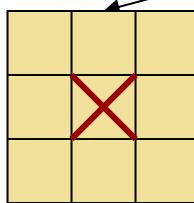
Example: tic-tac-toe

max level

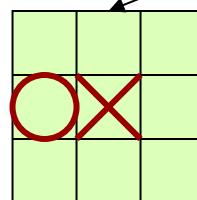


value ≥ 1

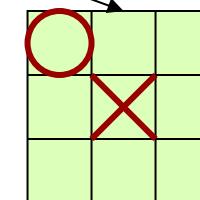
min level



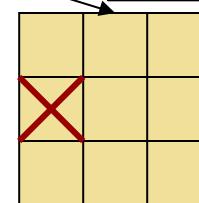
value = 1



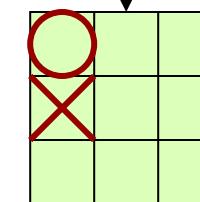
$e(n) = 2$



$e(n) = 1$



value ≤ -1

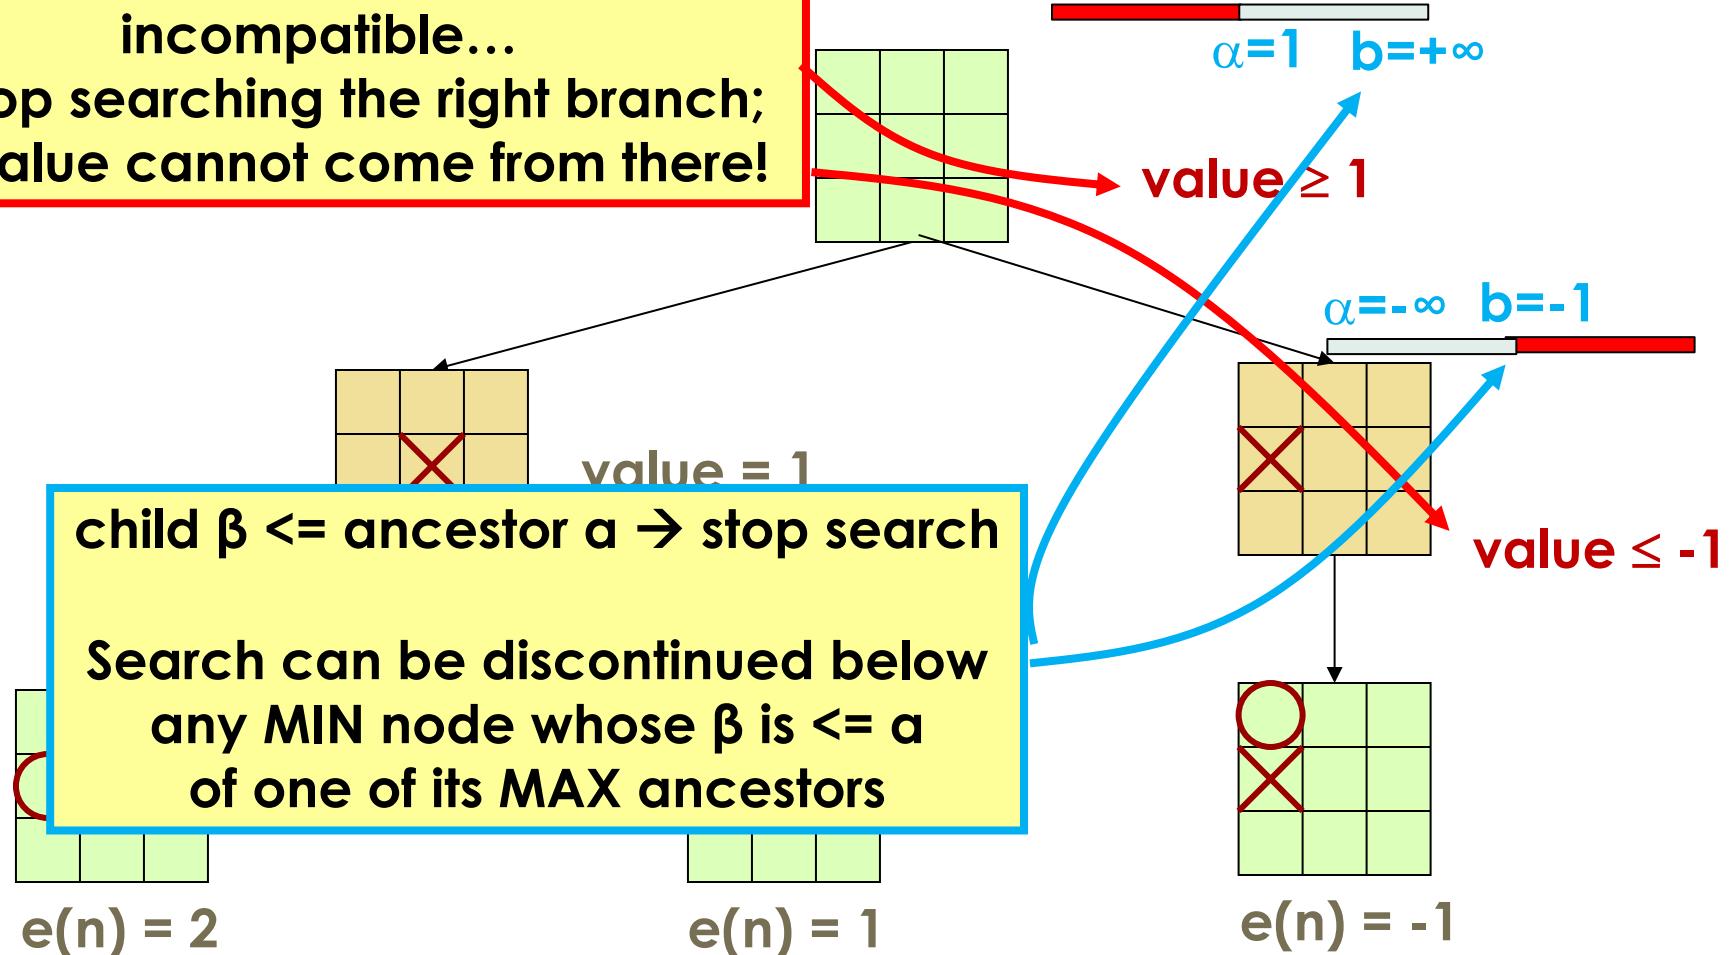


$e(n) = -1$

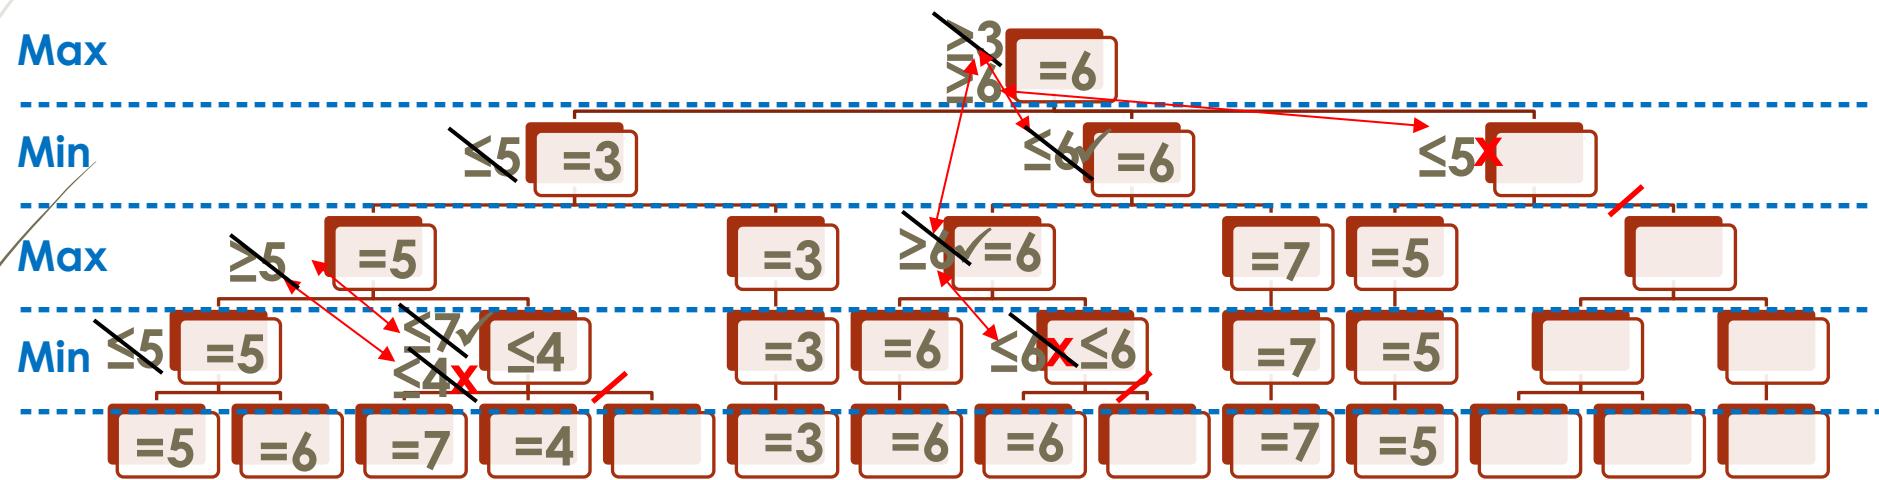
Example: tic-tac-toe

max level

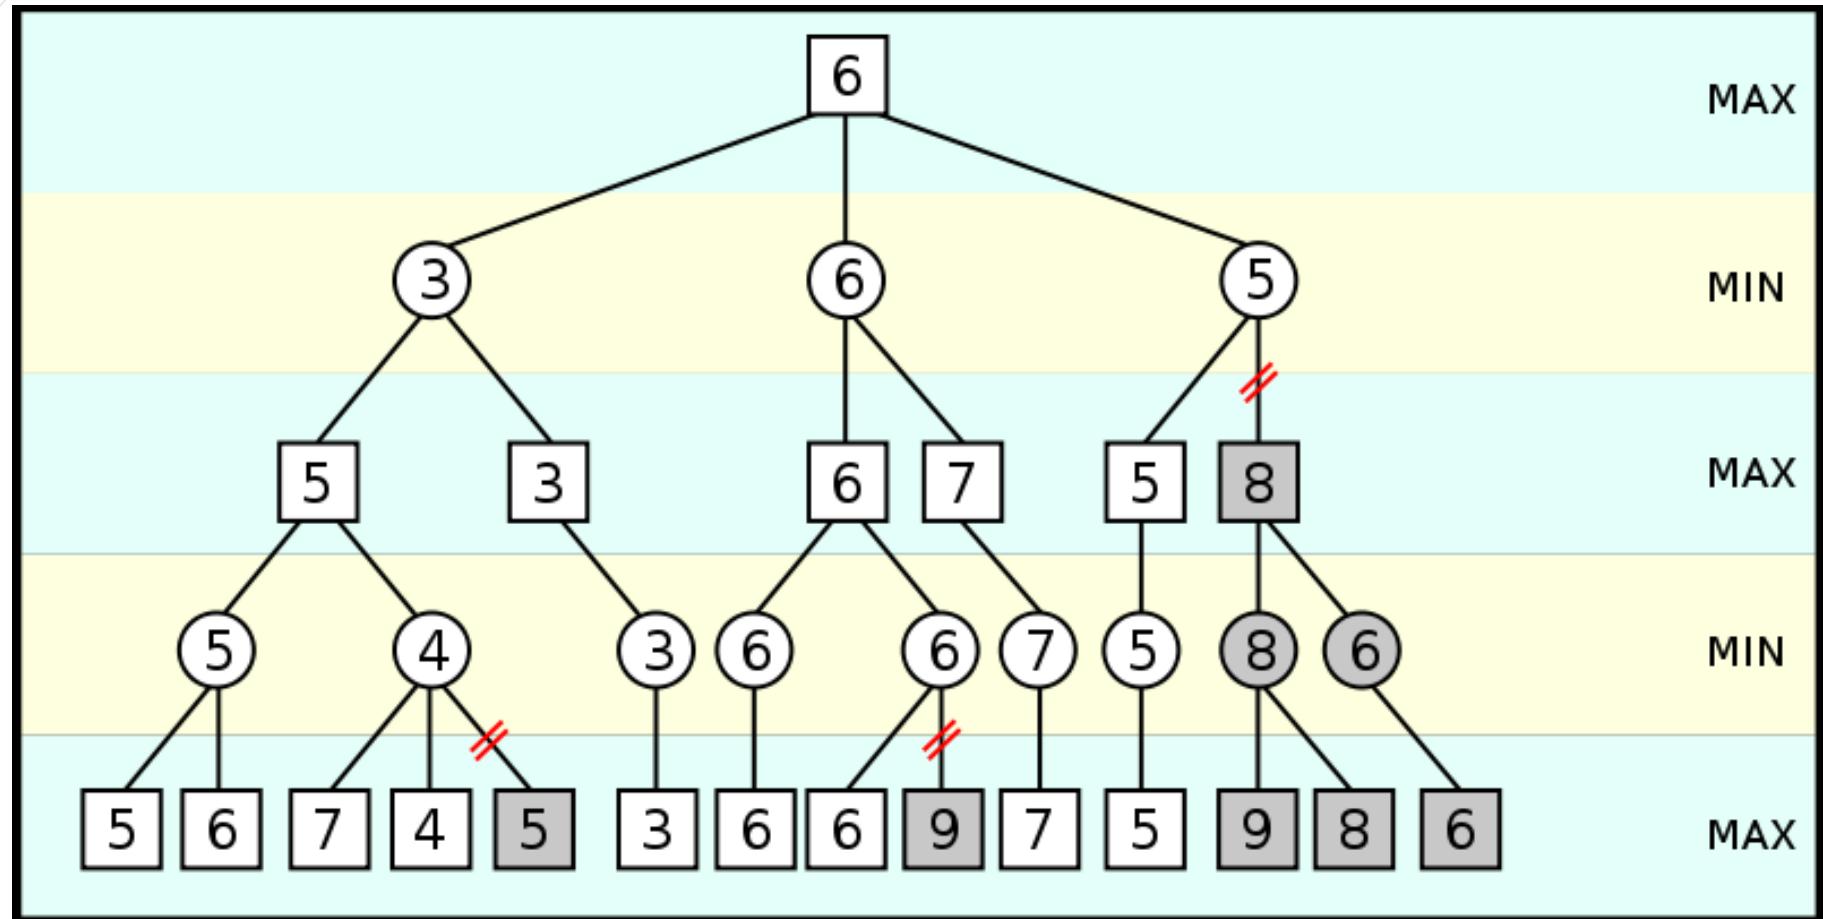
min level



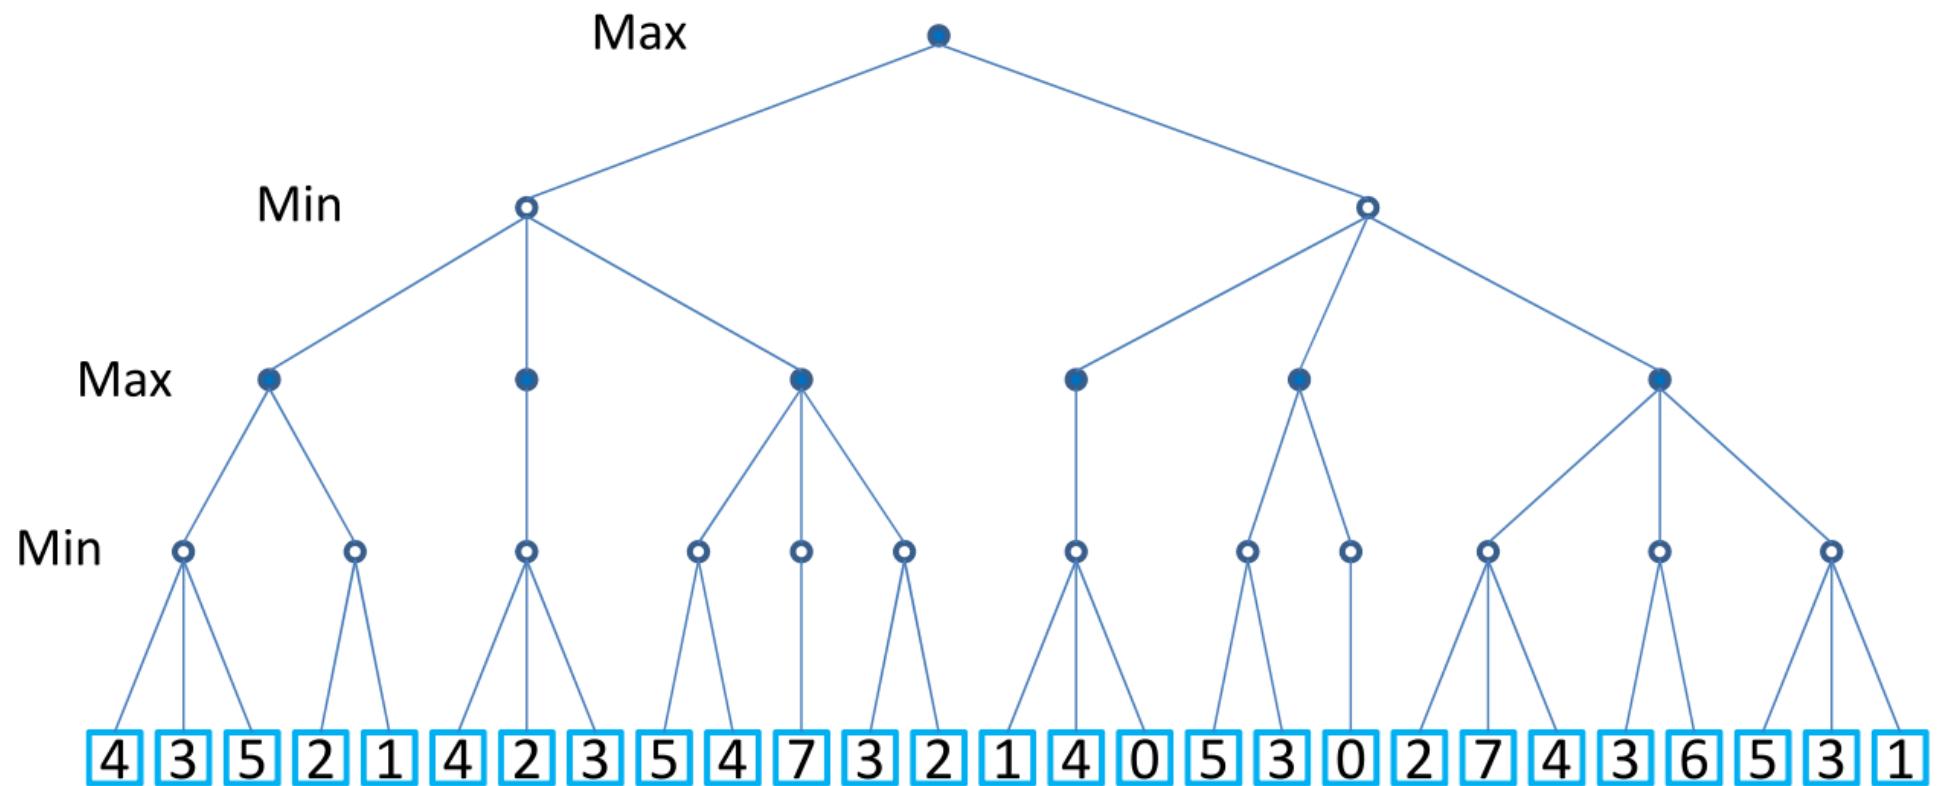
Alpha – Beta Pruning Example 2



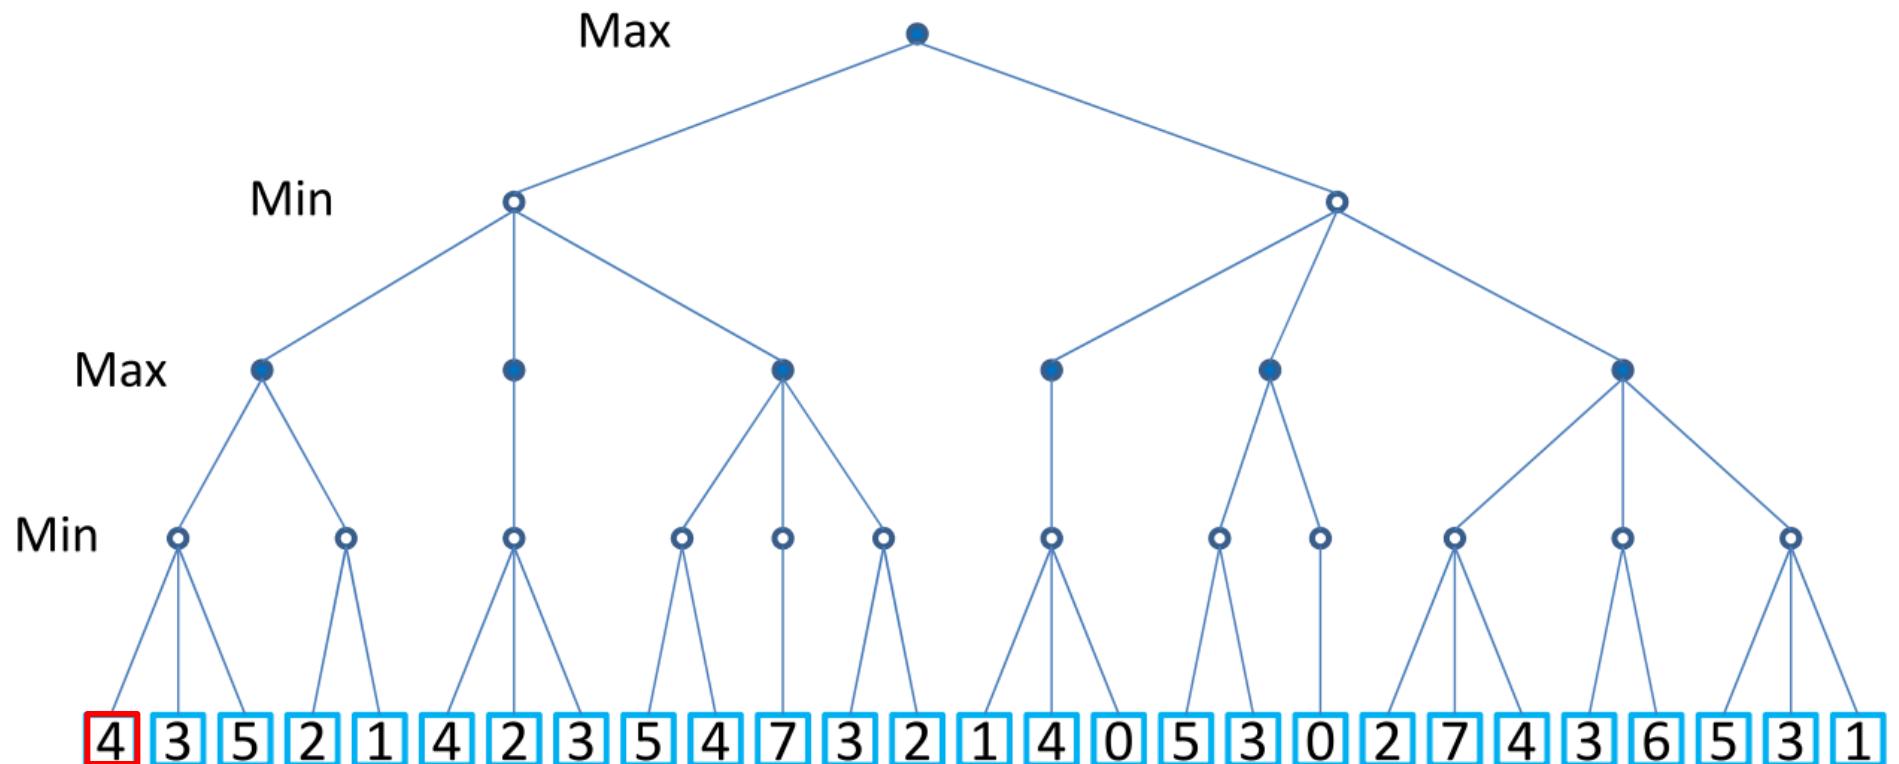
Alpha – Beta Pruning Example 2



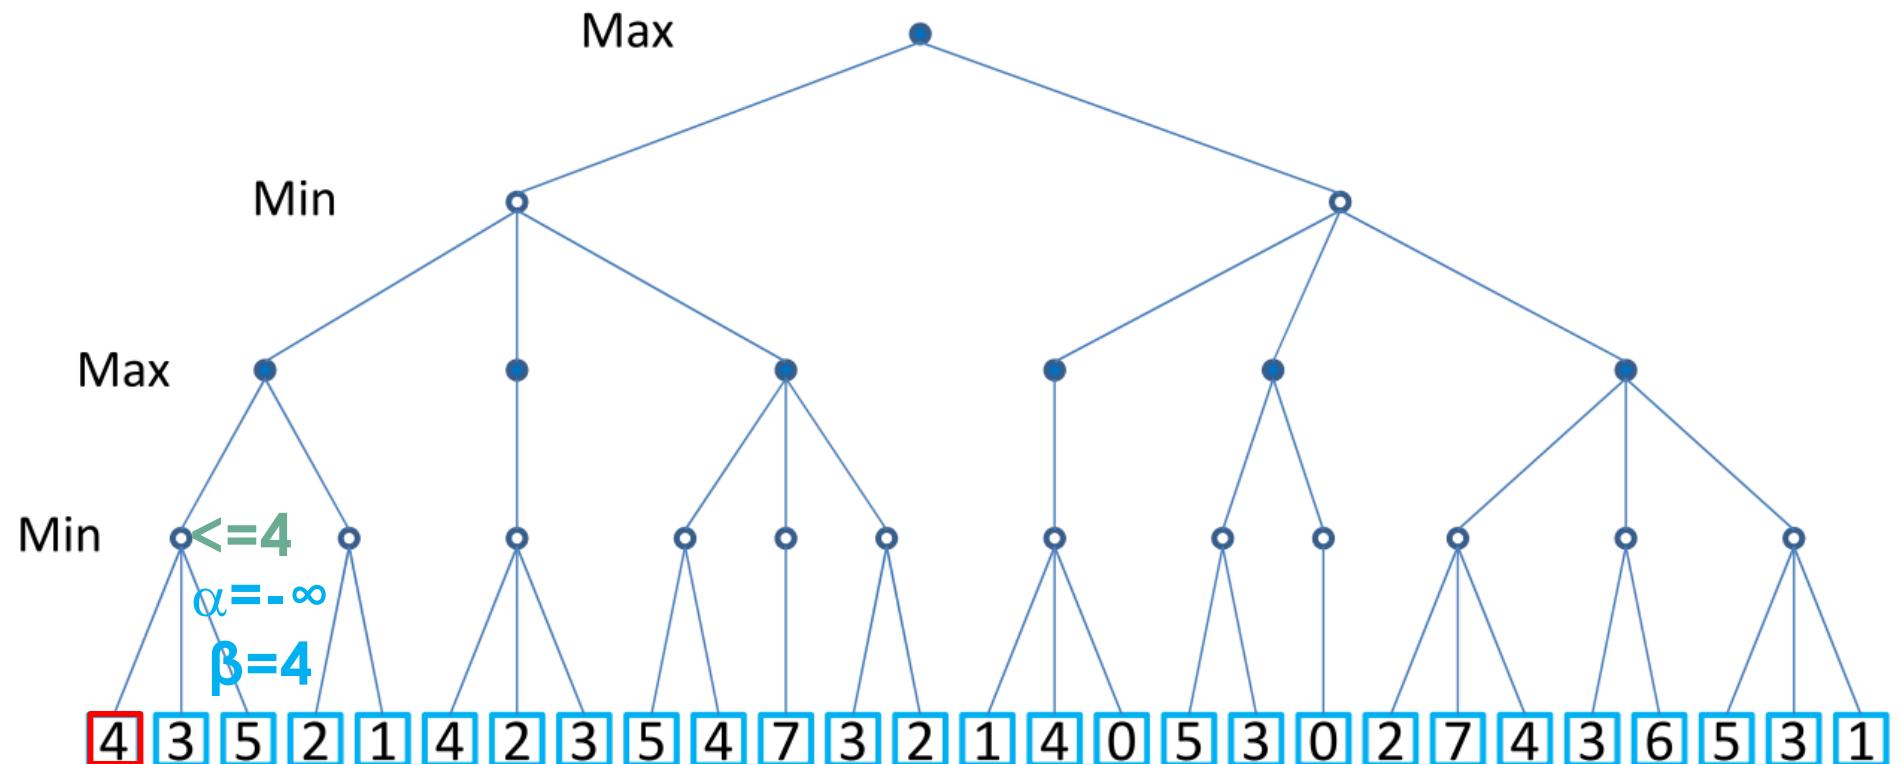
Alpha – Beta Pruning Example 3



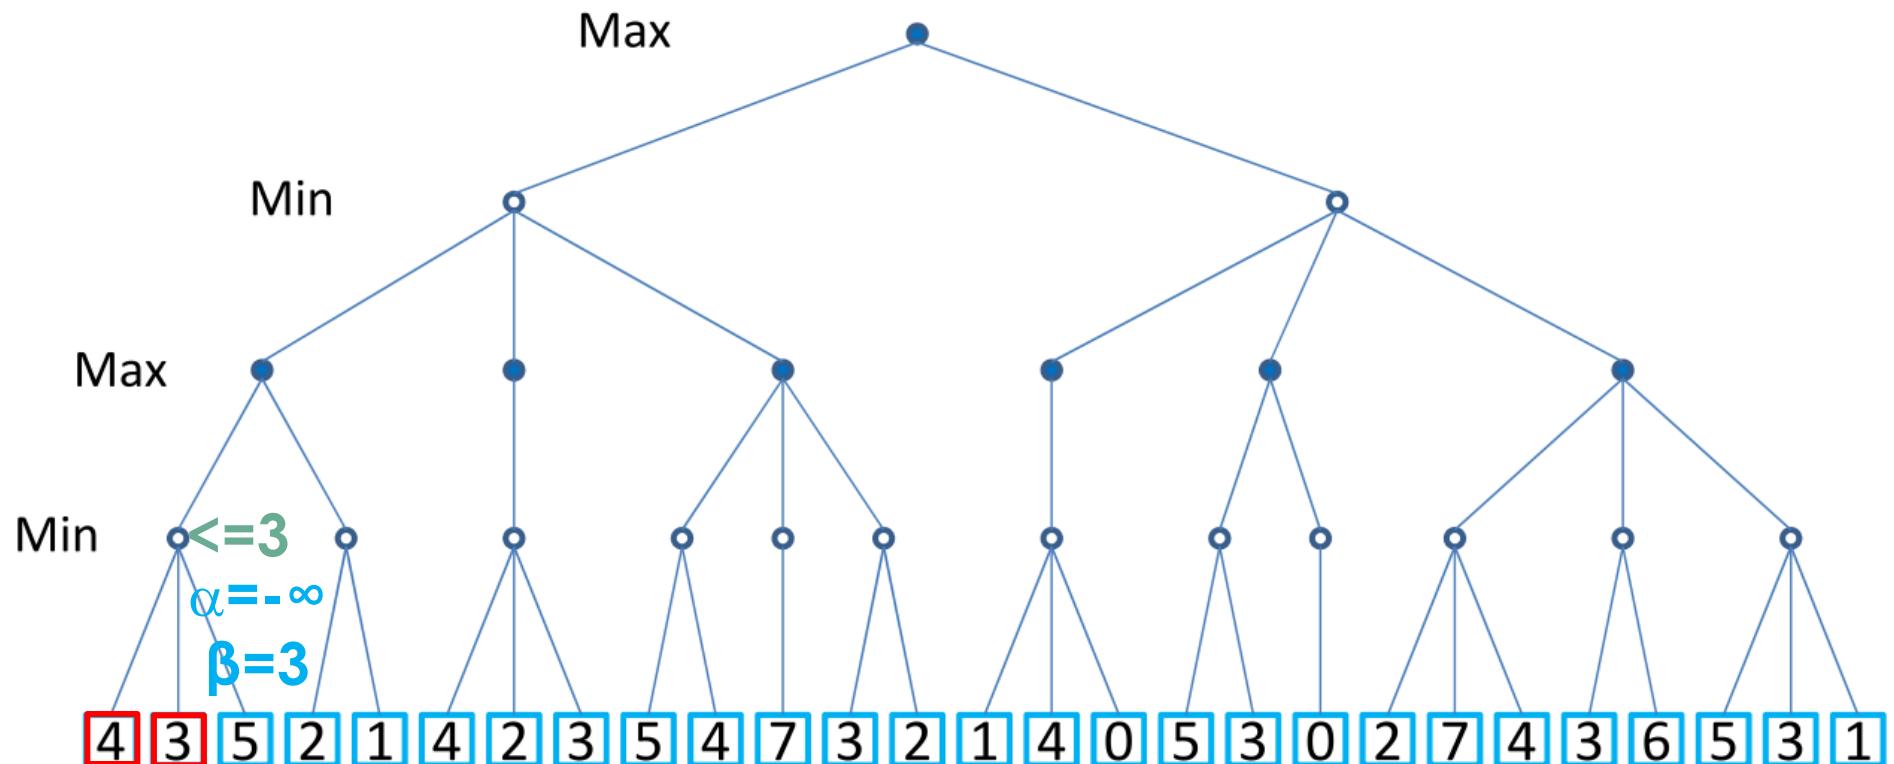
Alpha – Beta Pruning Example 3



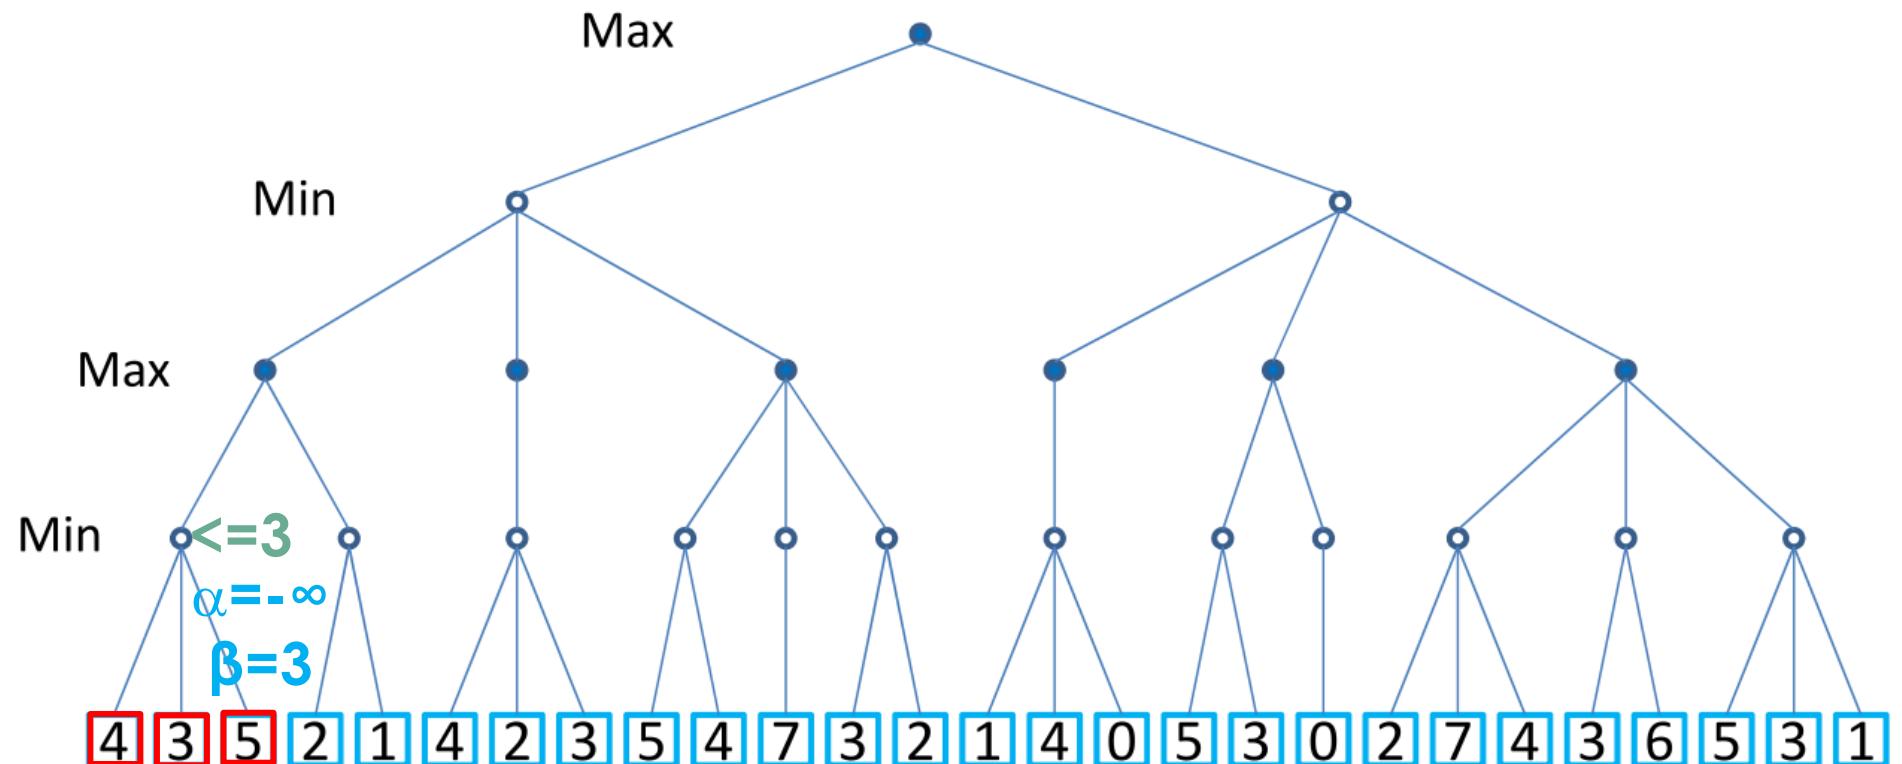
Alpha – Beta Pruning Example 3



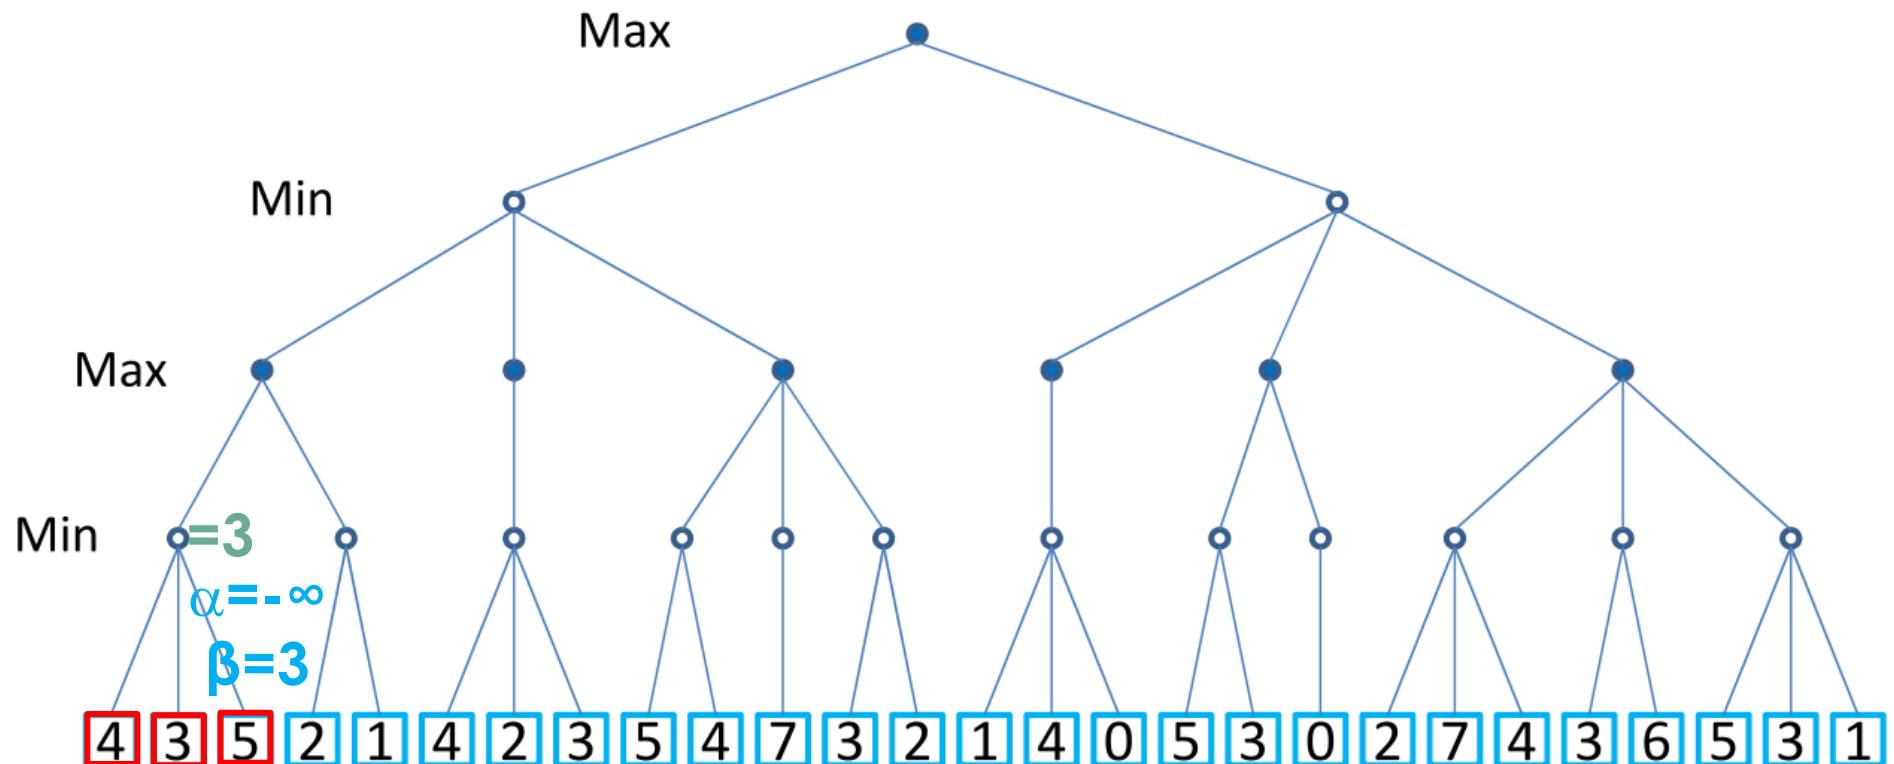
Alpha – Beta Pruning Example 3



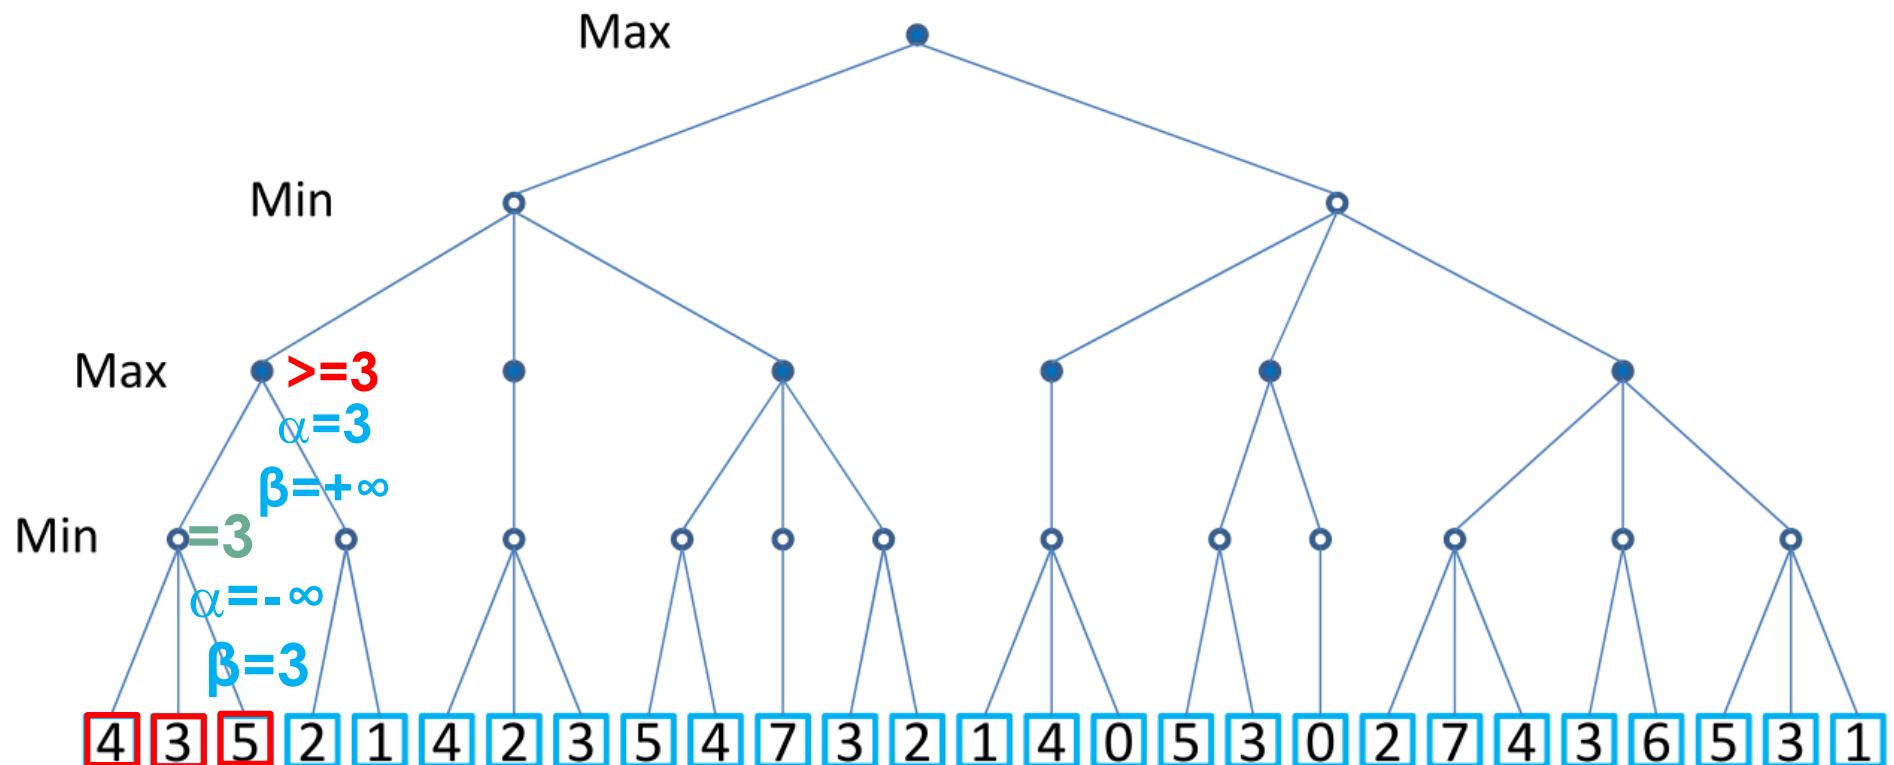
Alpha – Beta Pruning Example 3



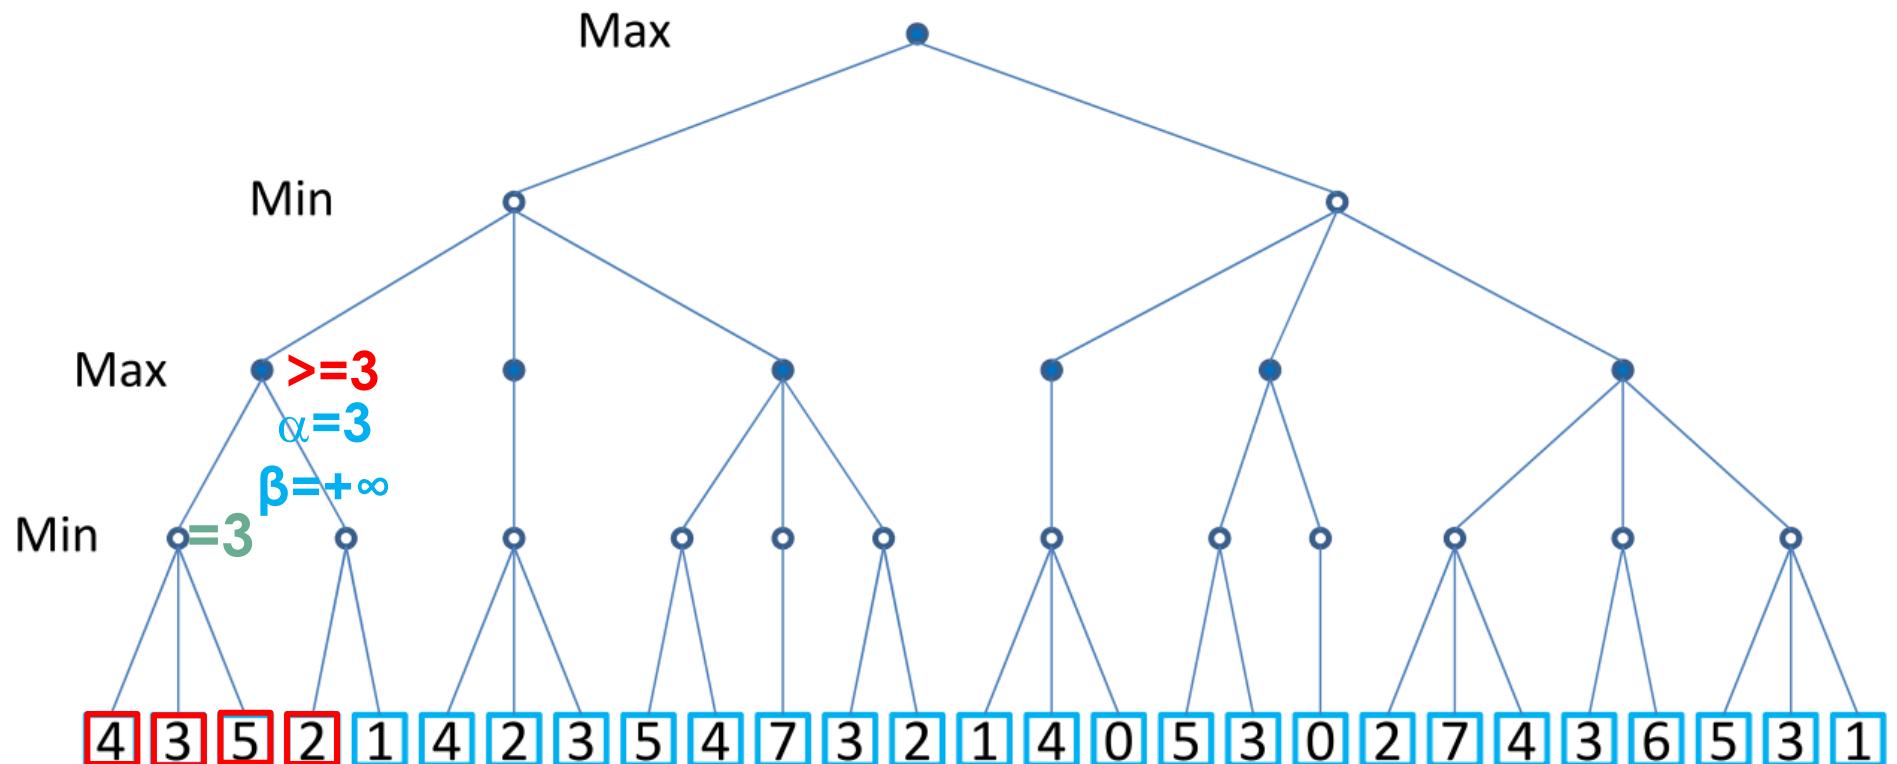
Alpha – Beta Pruning Example 3



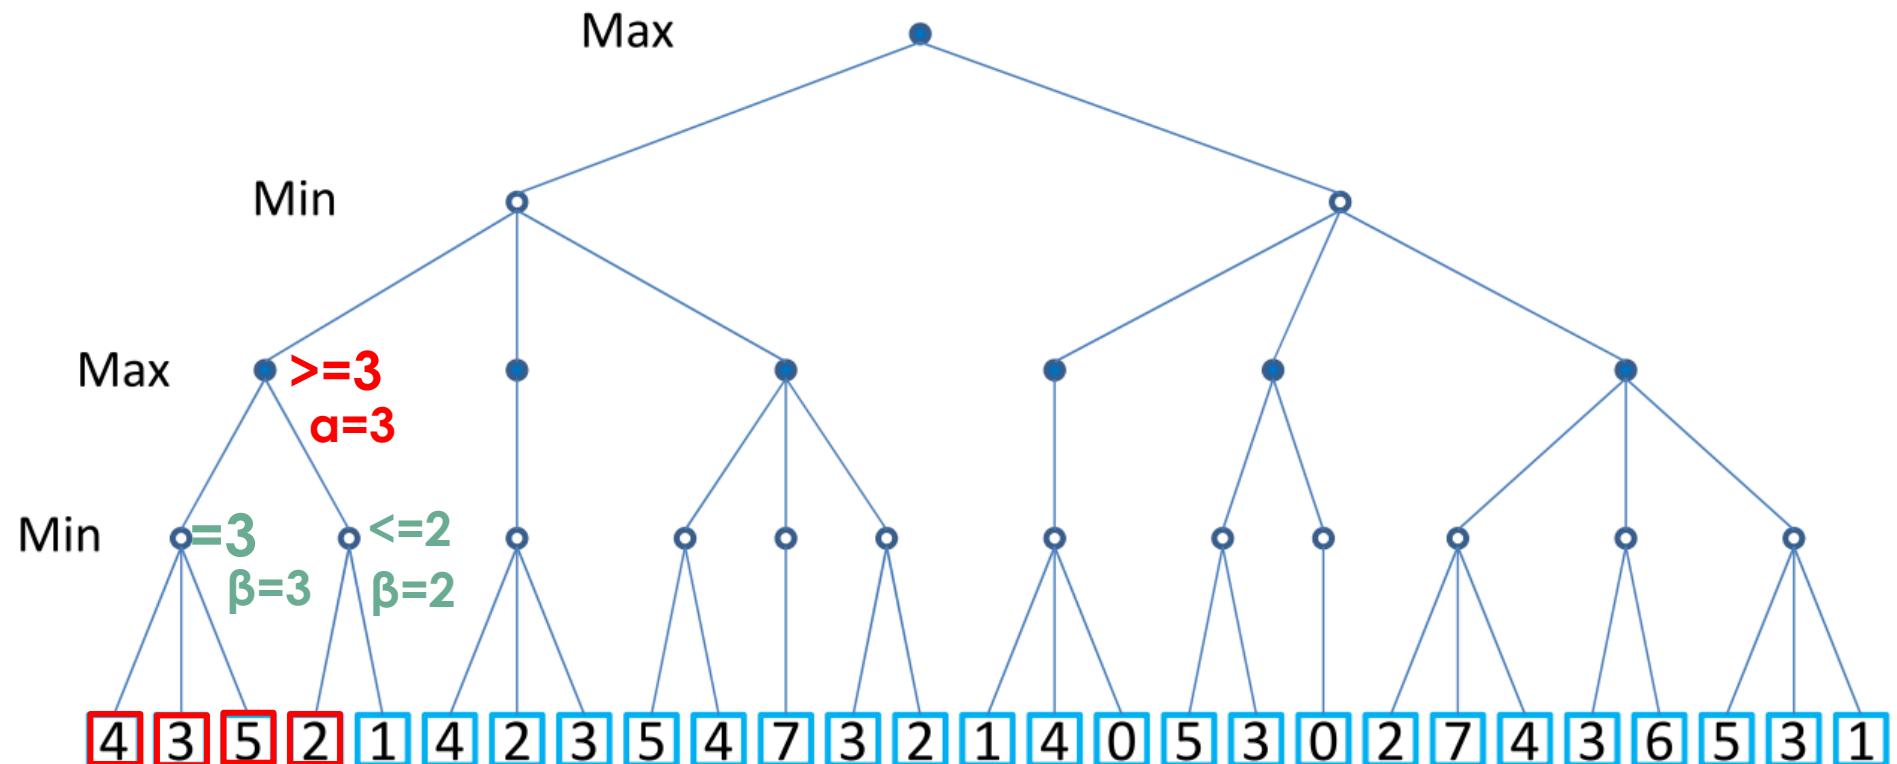
Alpha – Beta Pruning Example 3



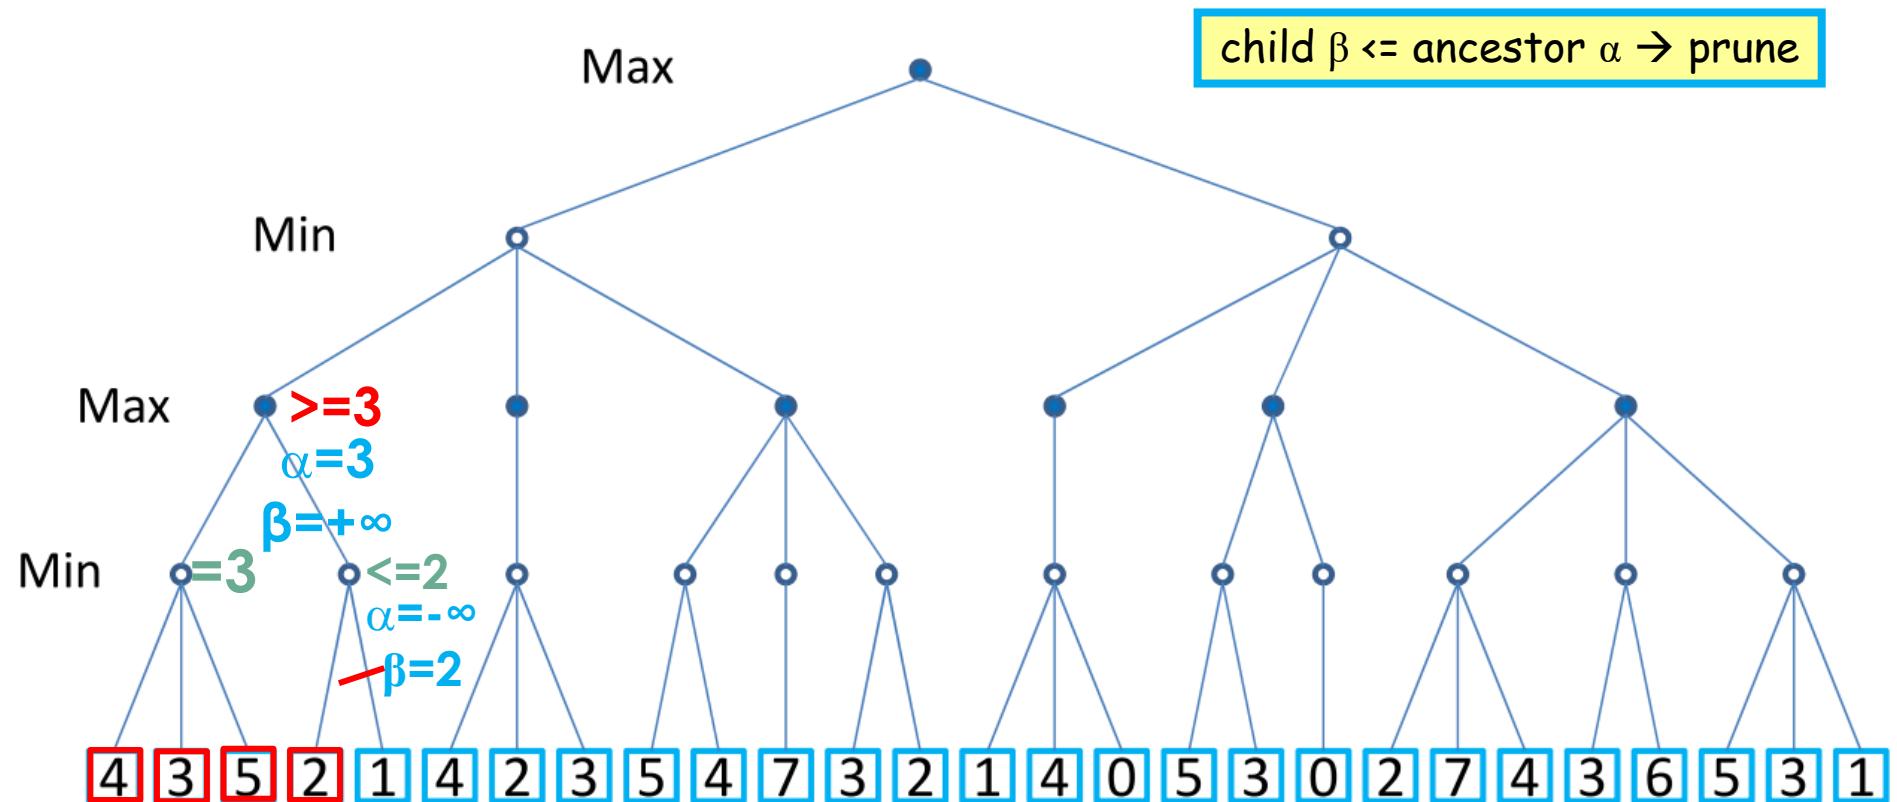
Alpha – Beta Pruning Example 3



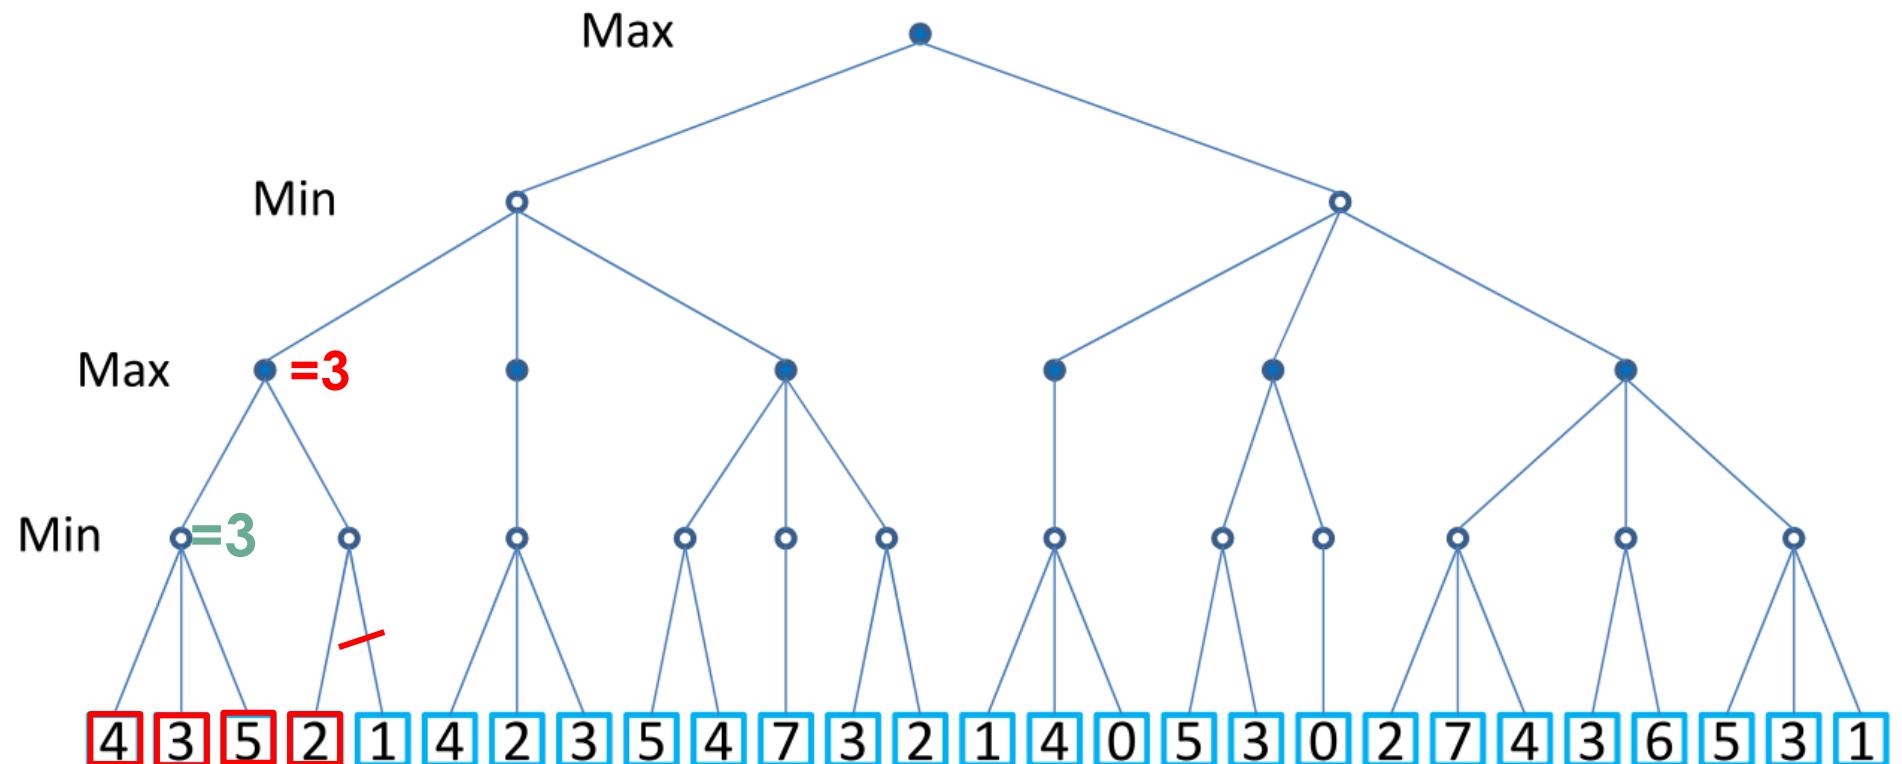
Alpha – Beta Pruning Example 3



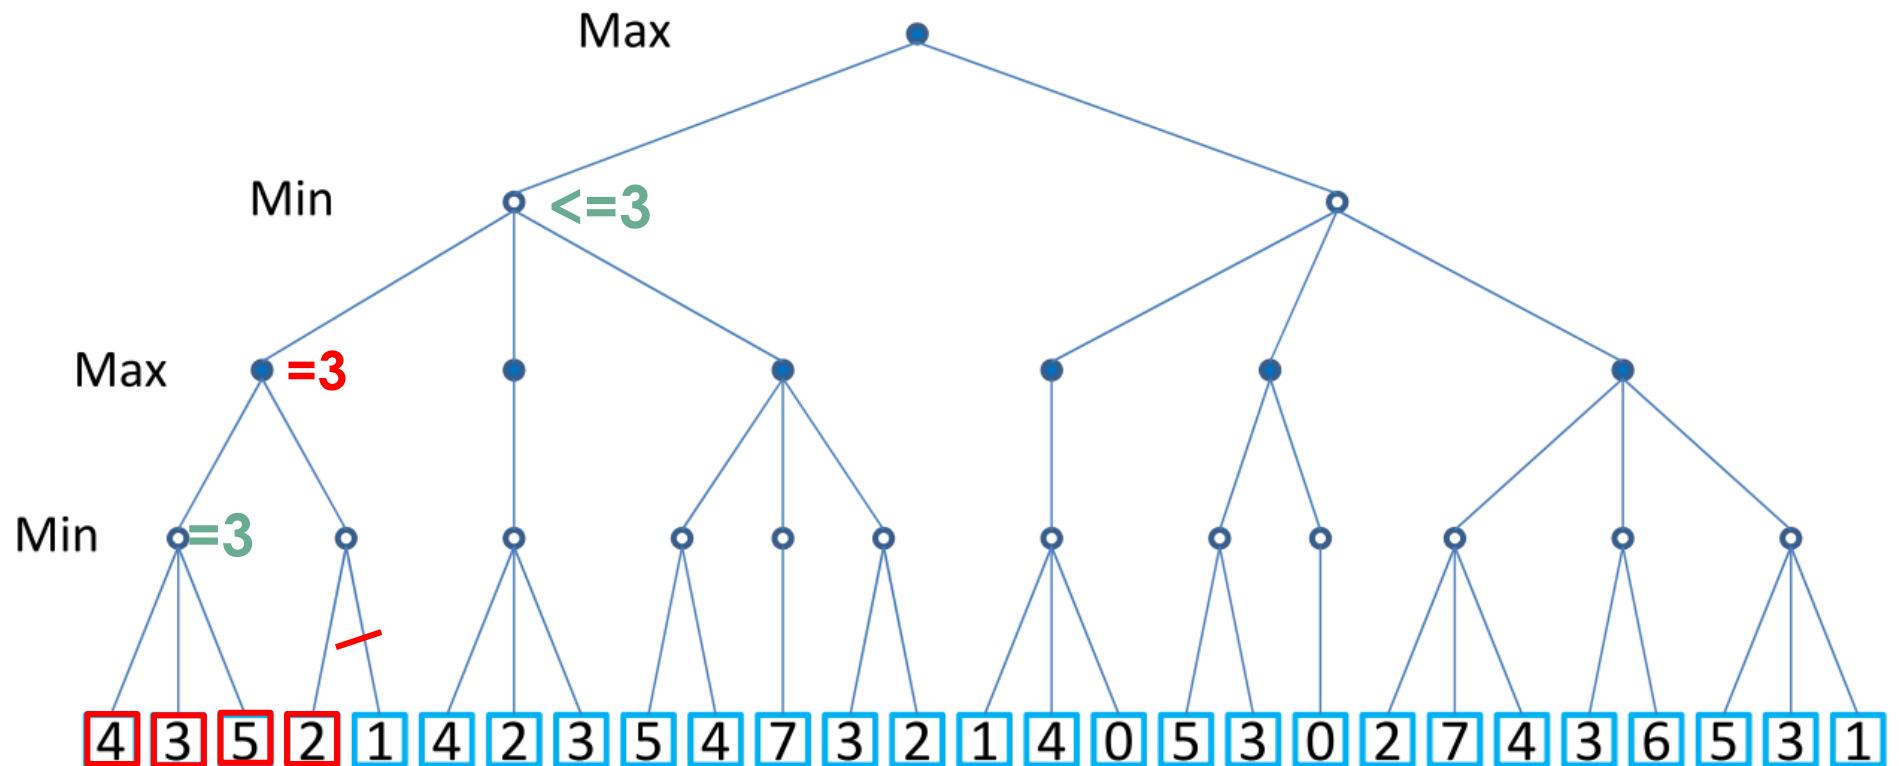
Alpha – Beta Pruning Example 3



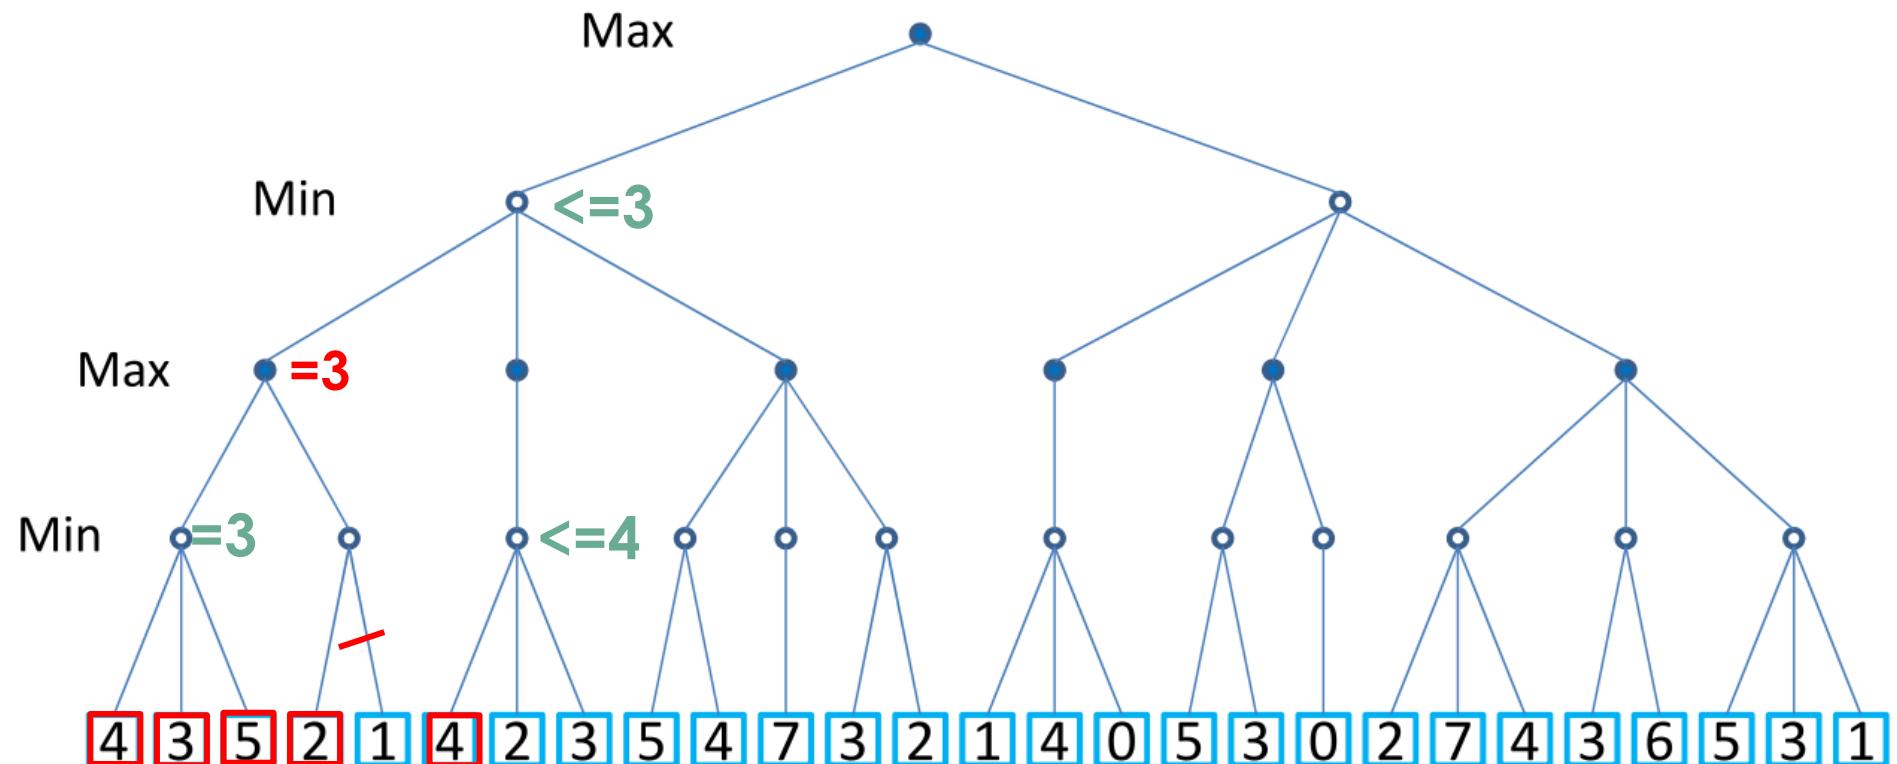
Alpha – Beta Pruning Example 3



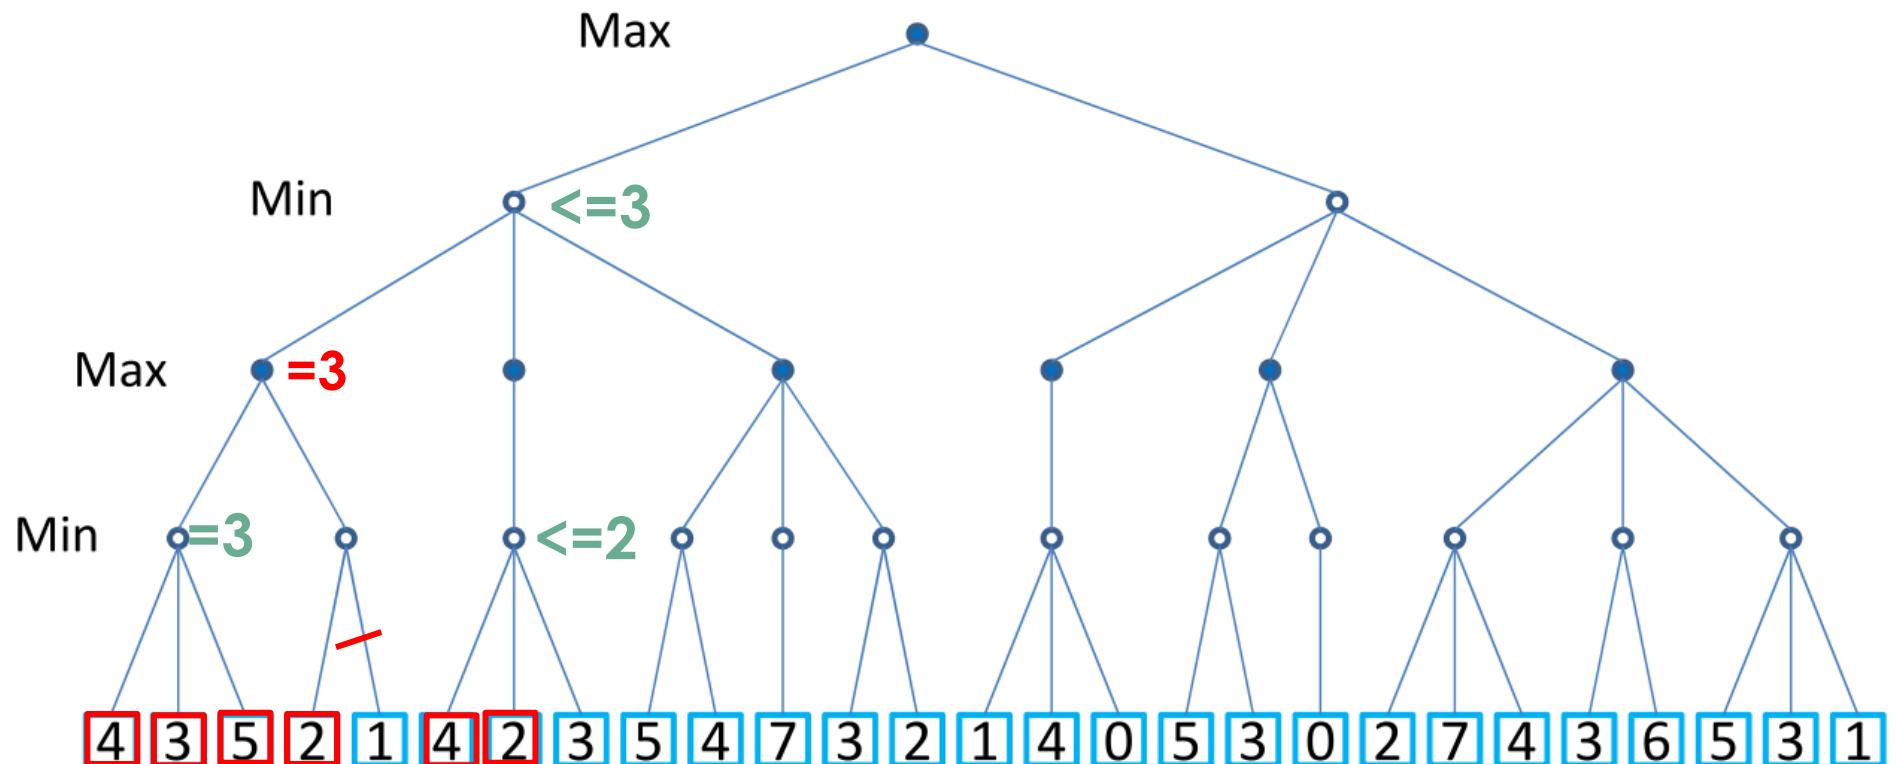
Alpha – Beta Pruning Example 3



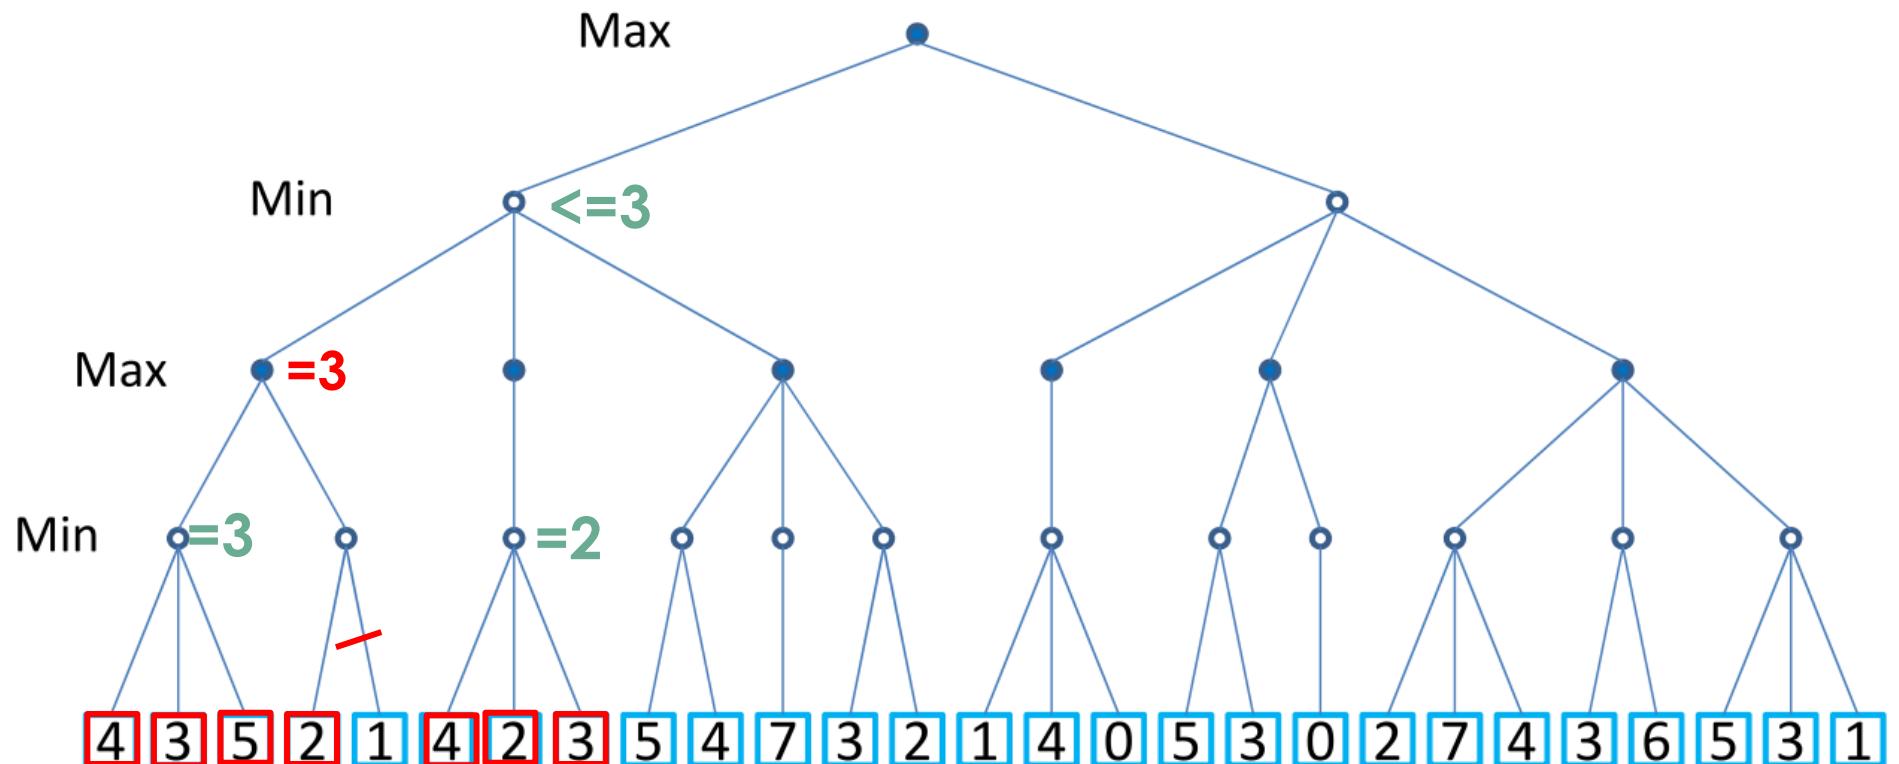
Alpha – Beta Pruning Example 3



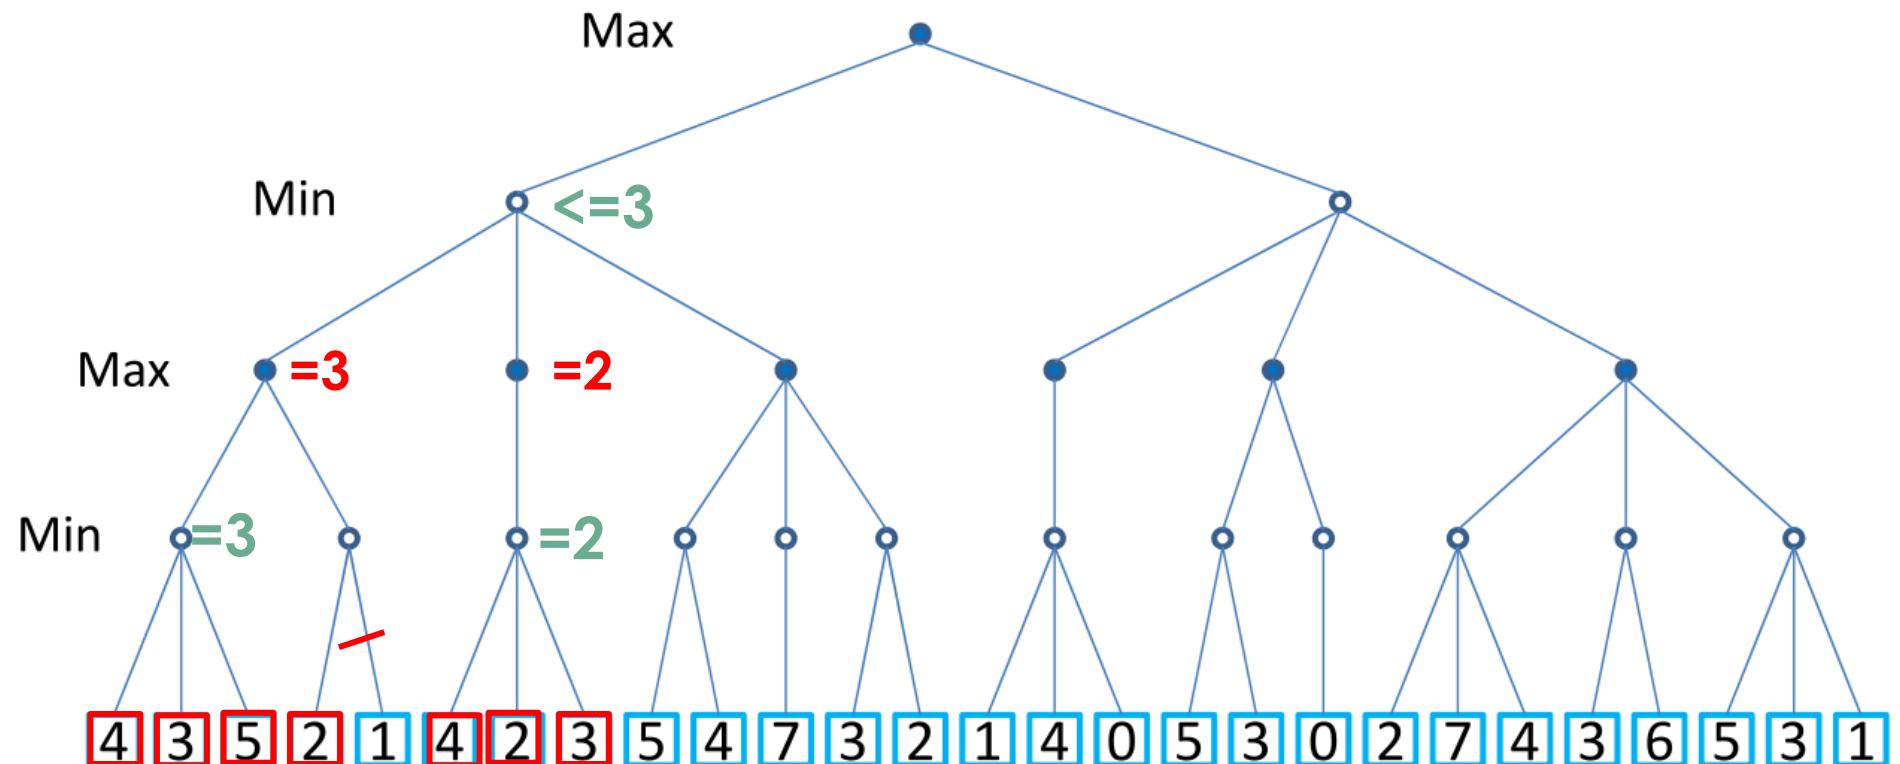
Alpha – Beta Pruning Example 3



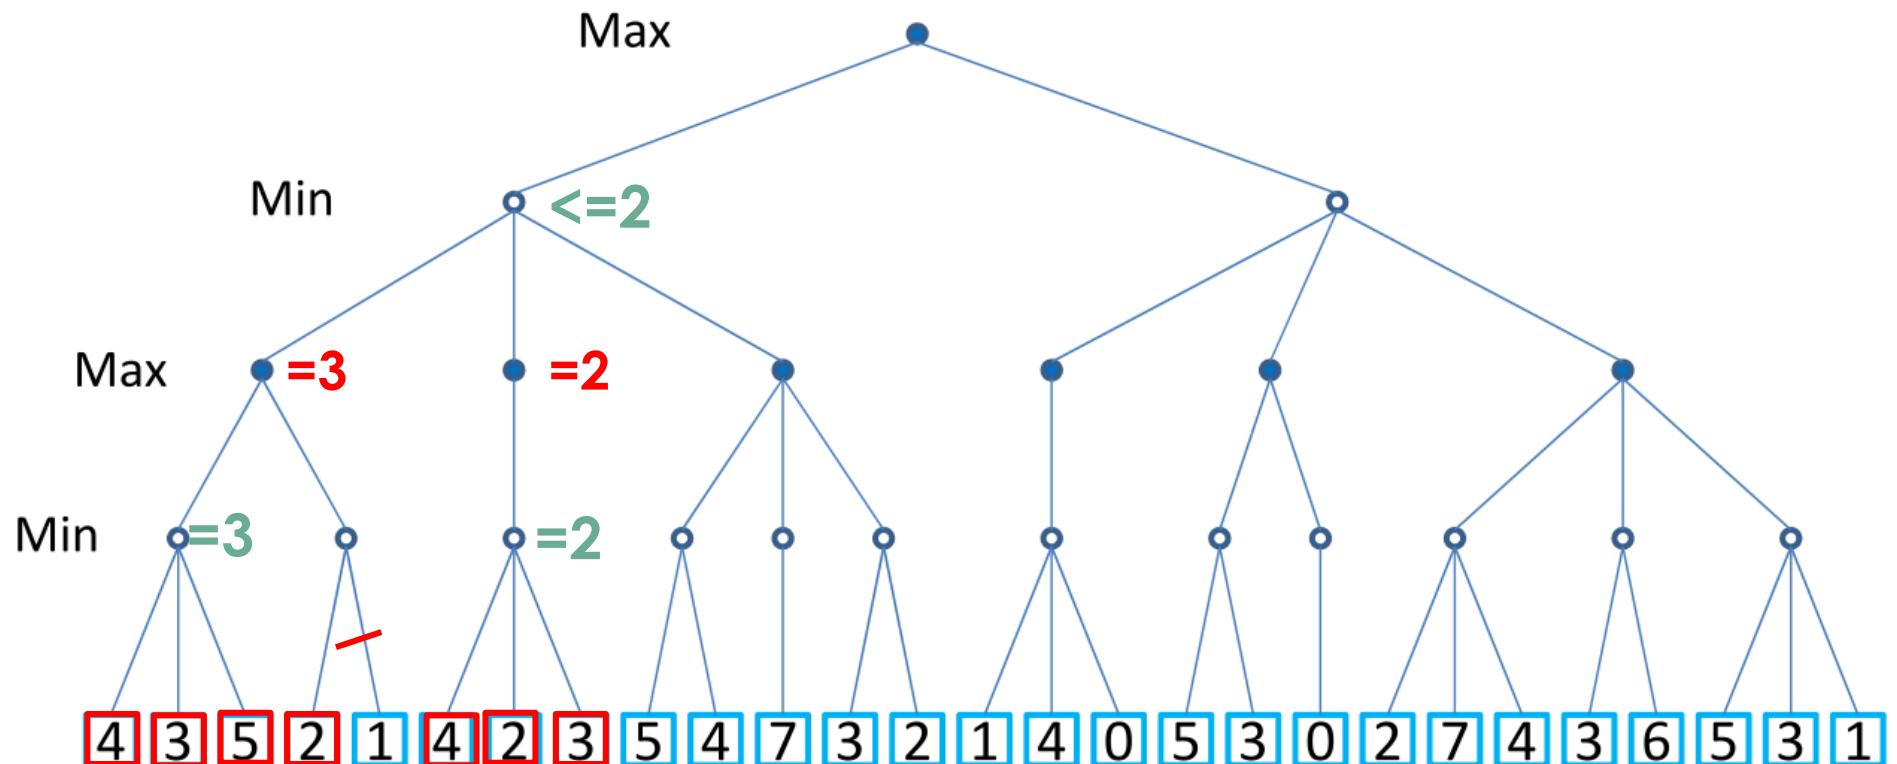
Alpha – Beta Pruning Example 3



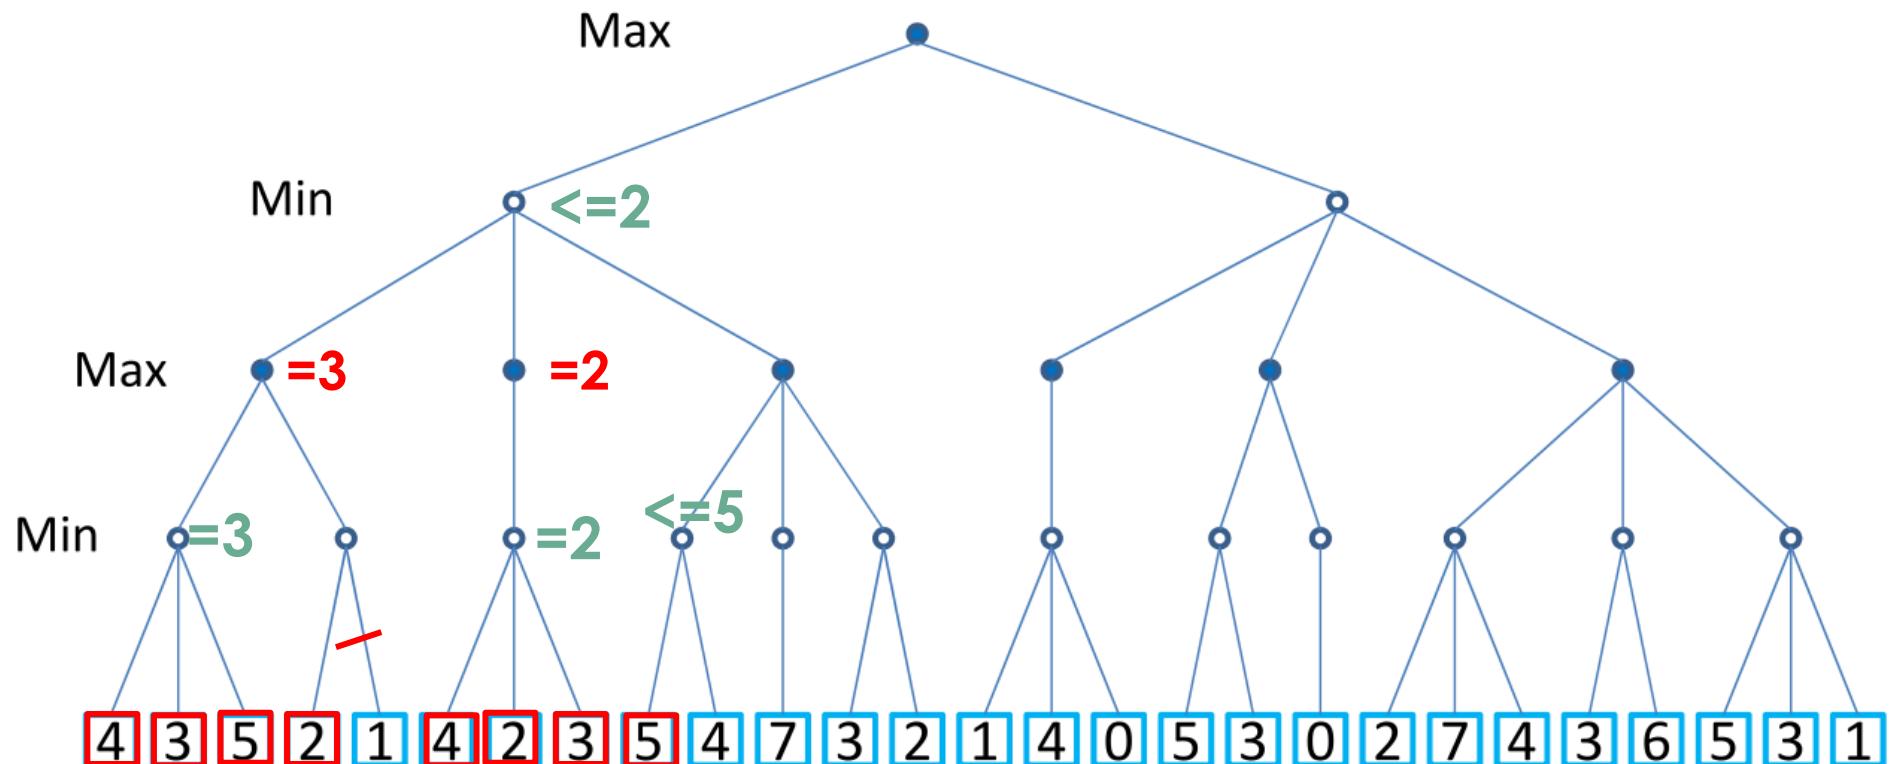
Alpha – Beta Pruning Example 3



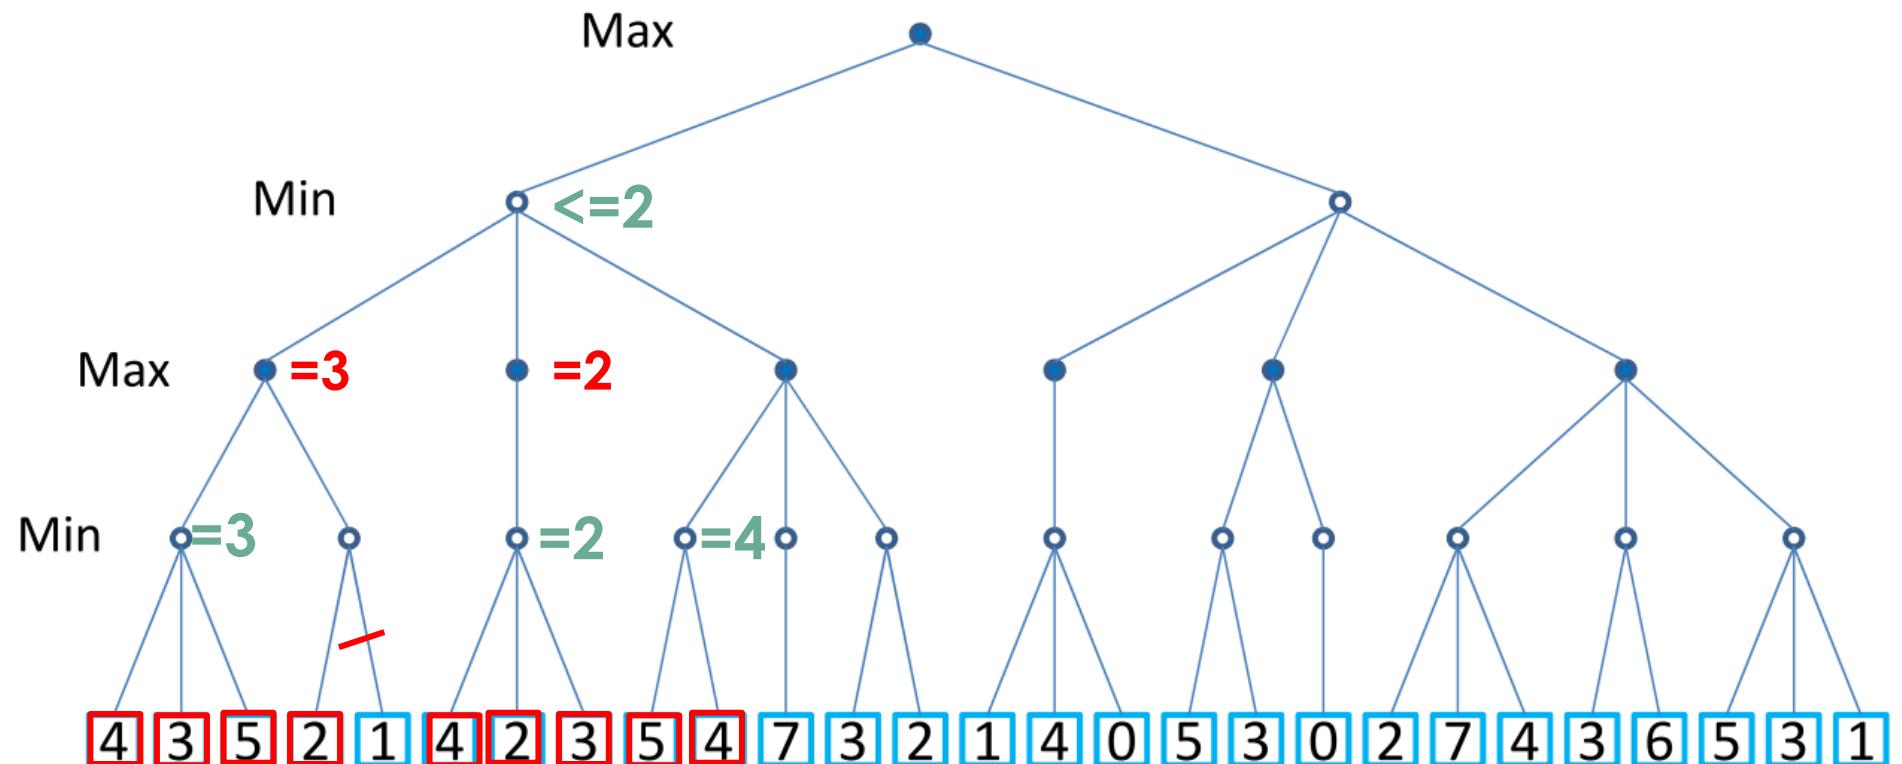
Alpha – Beta Pruning Example 3



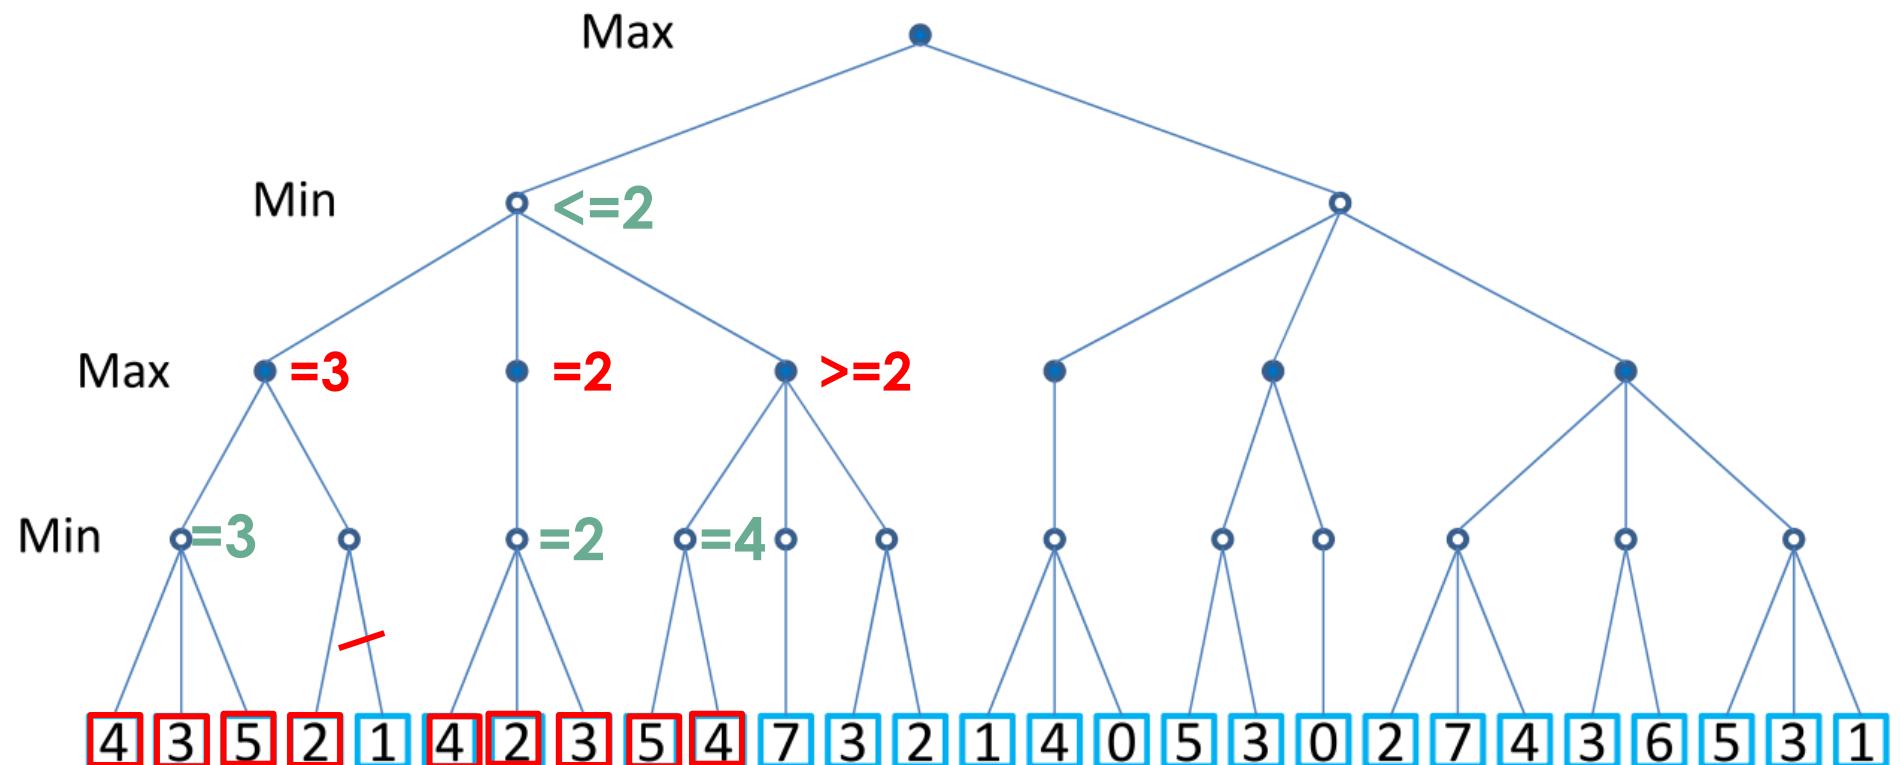
Alpha – Beta Pruning Example 3



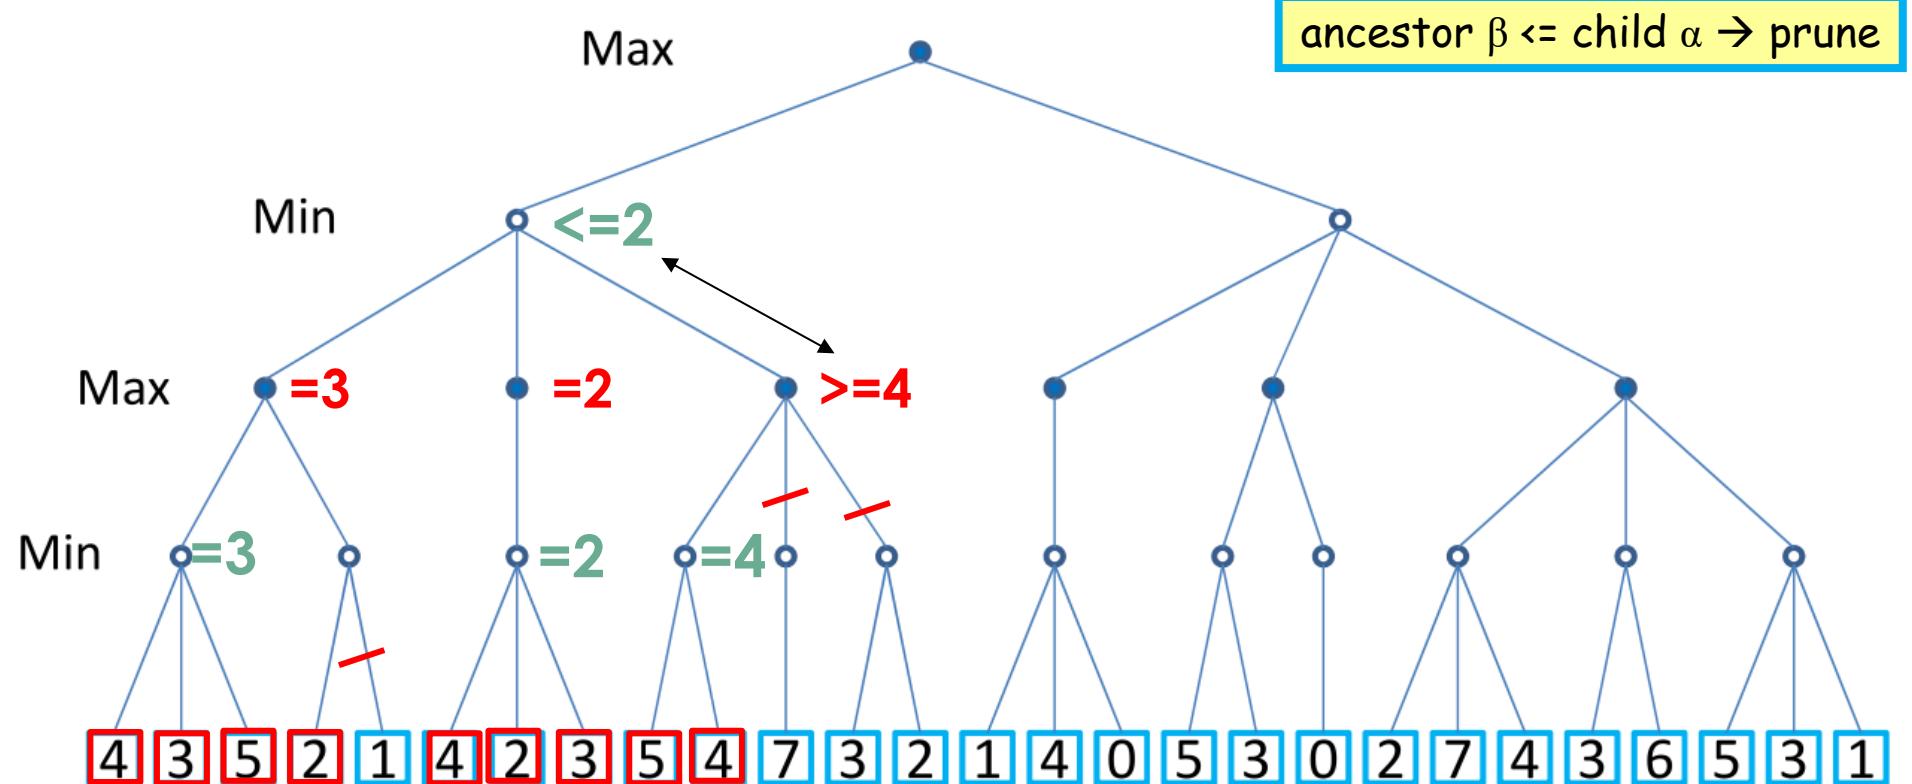
Alpha – Beta Pruning Example 3



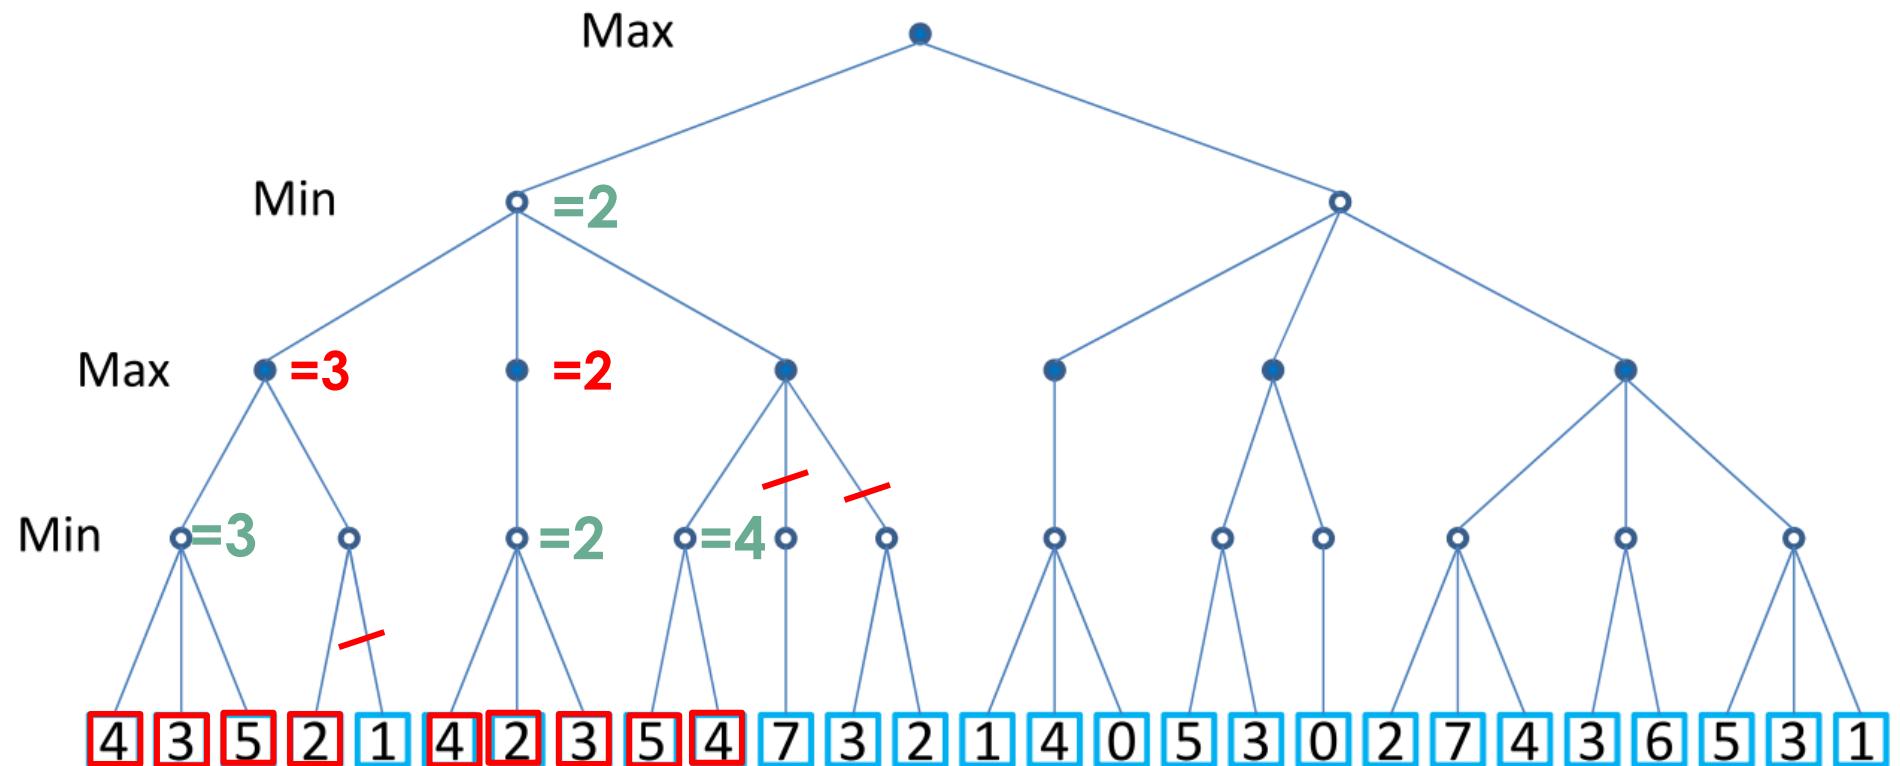
Alpha – Beta Pruning Example 3



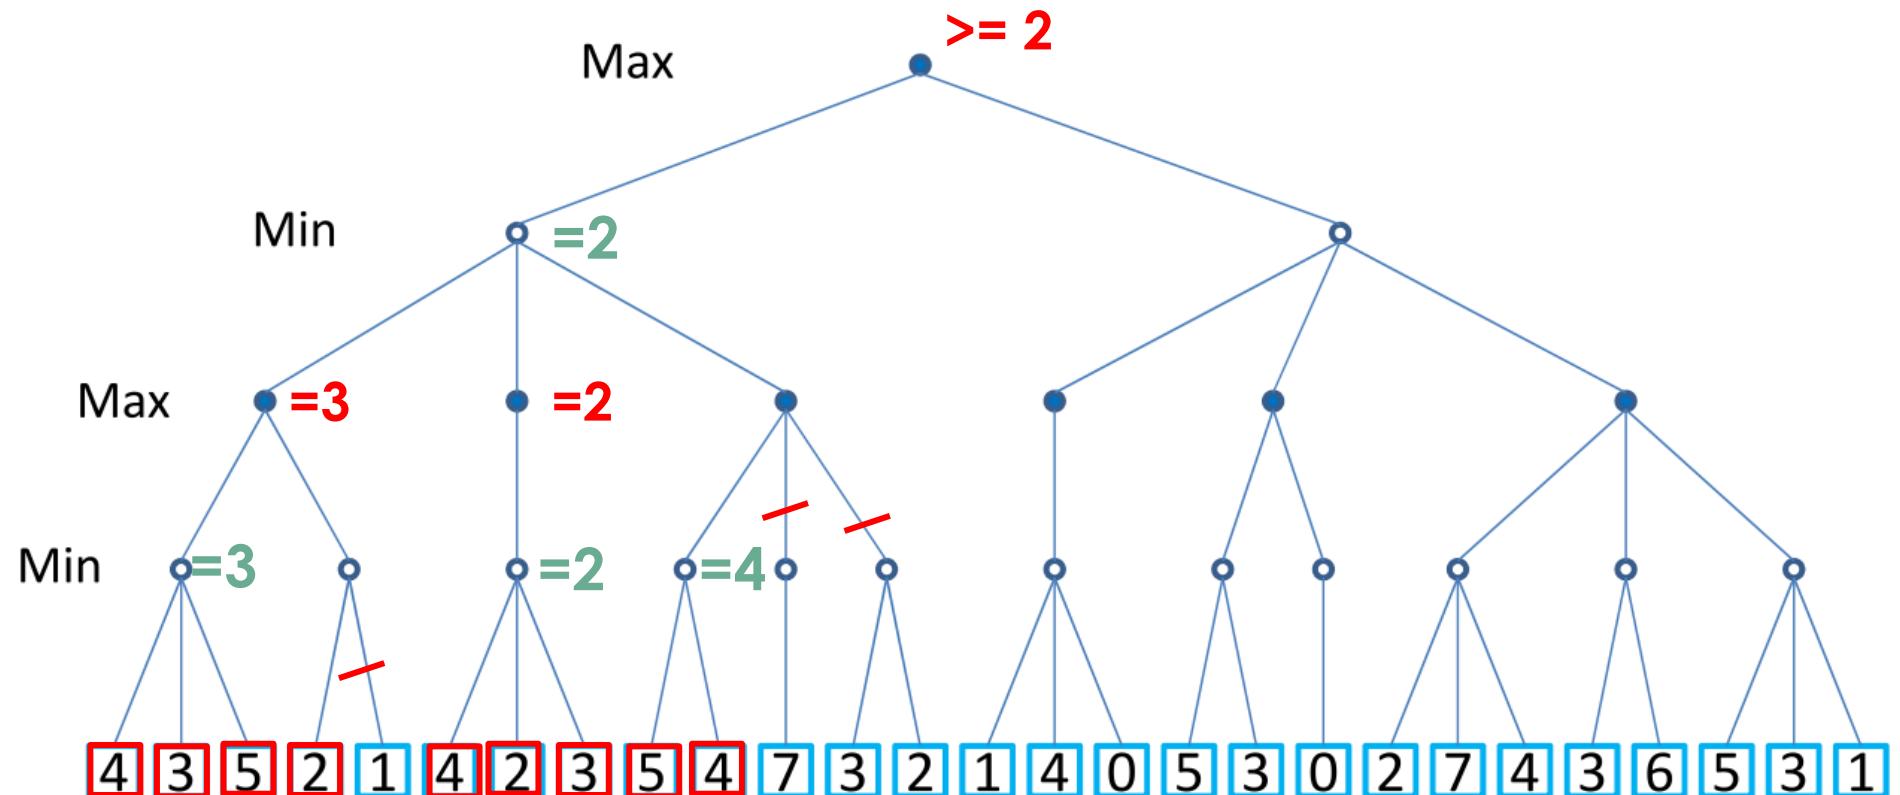
Alpha – Beta Pruning Example 3



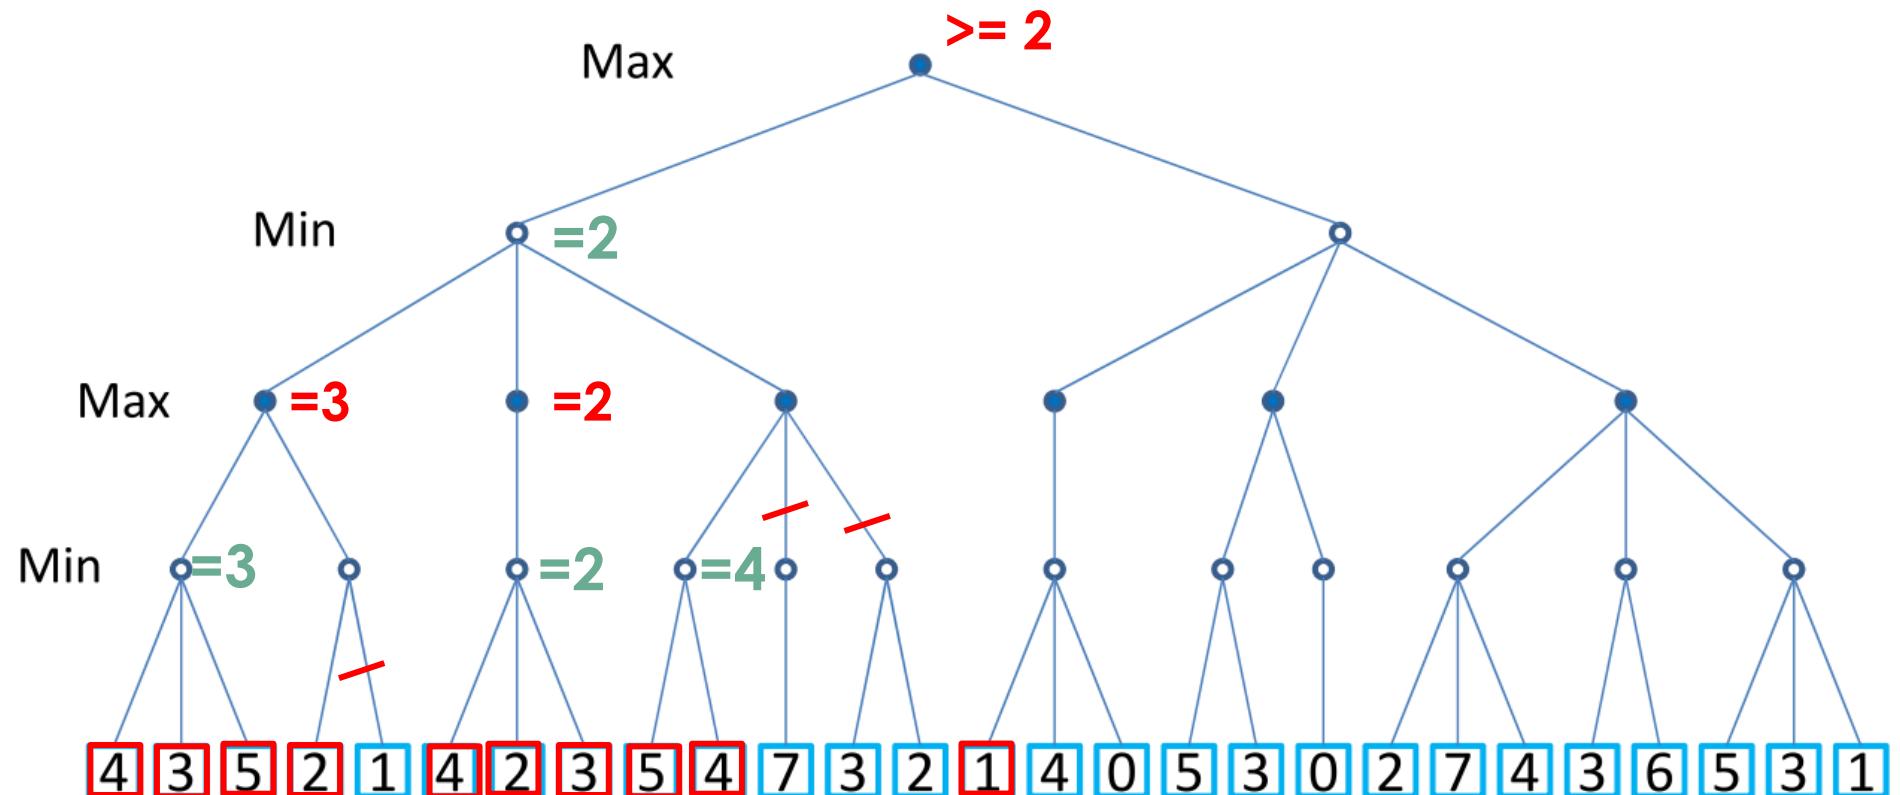
Alpha – Beta Pruning Example 3



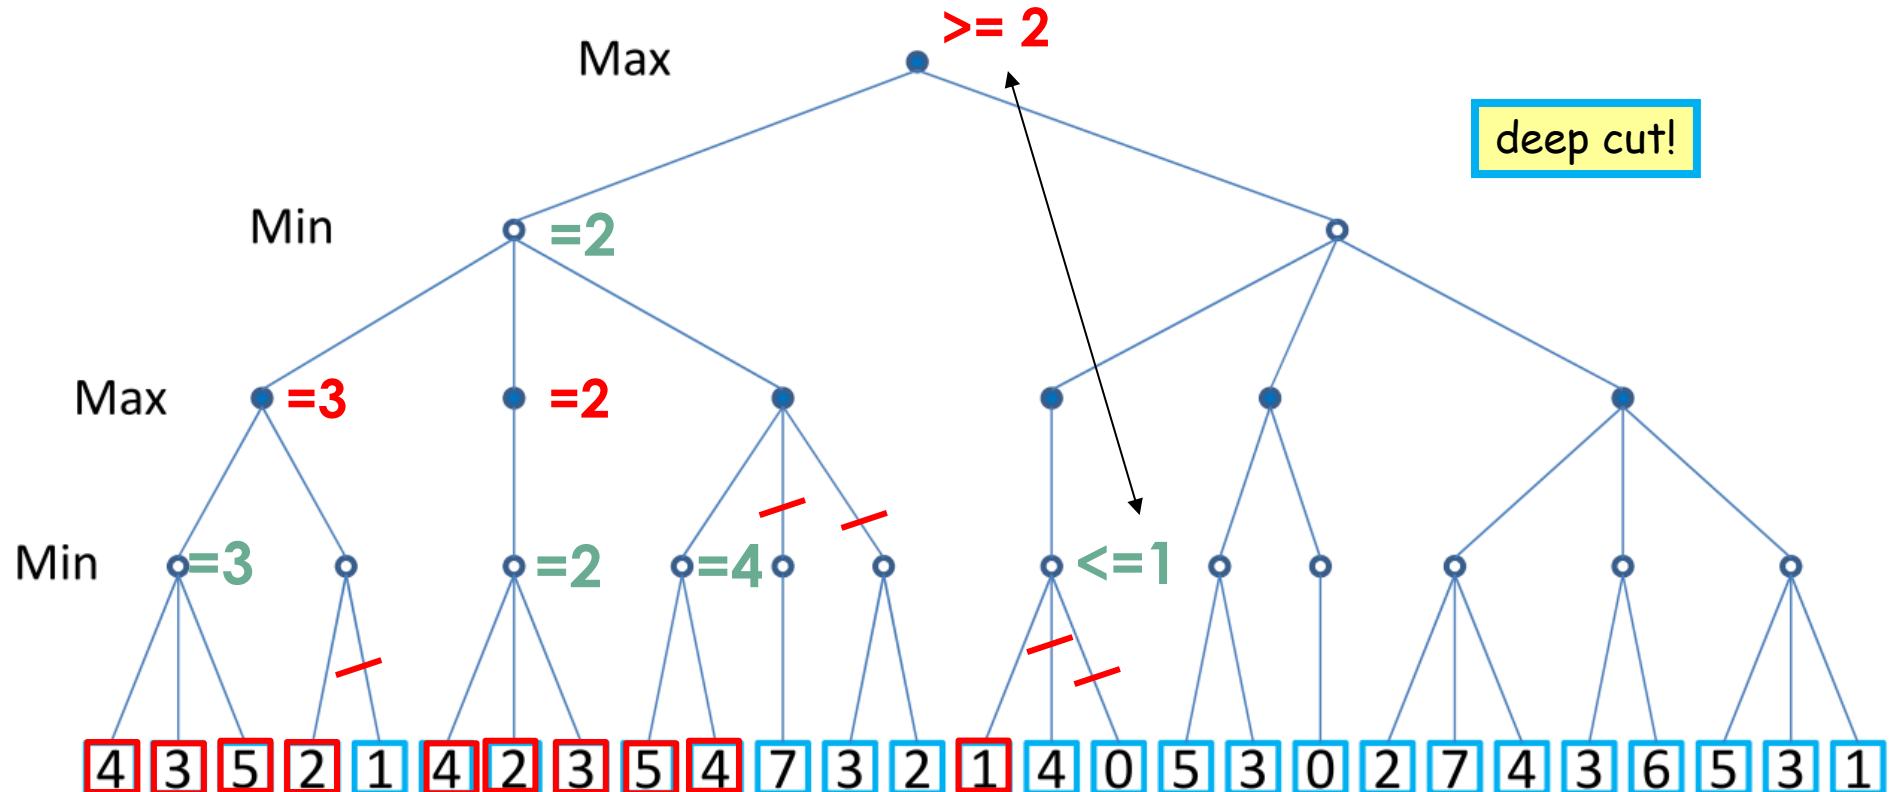
Alpha – Beta Pruning Example 3



Alpha – Beta Pruning Example 3

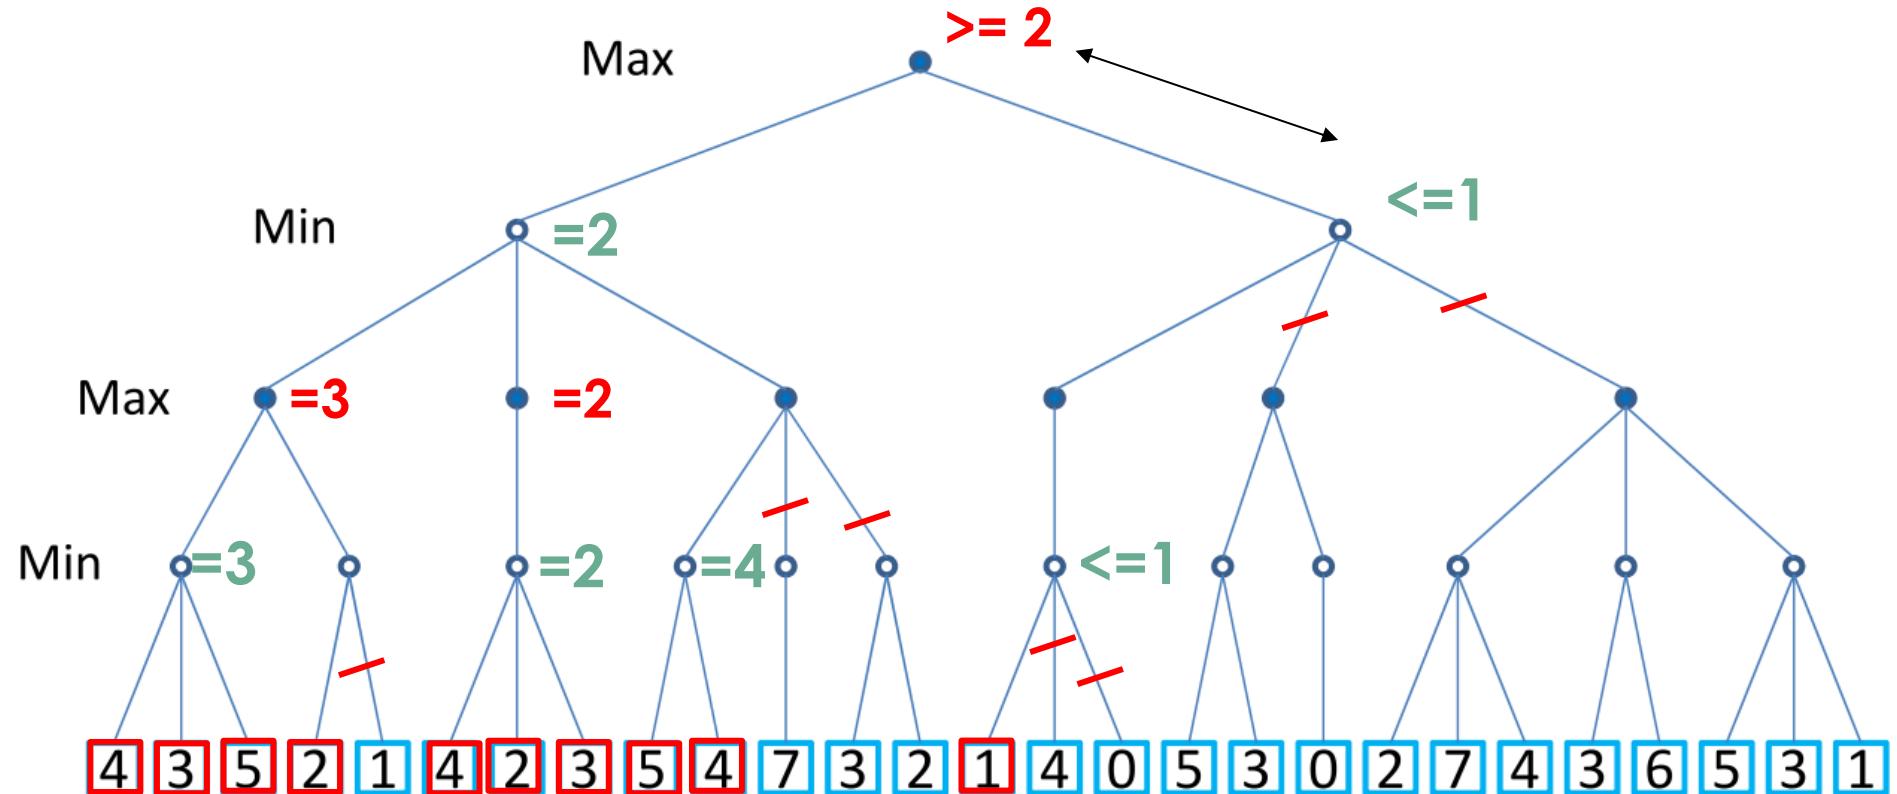


Alpha – Beta Pruning Example 3

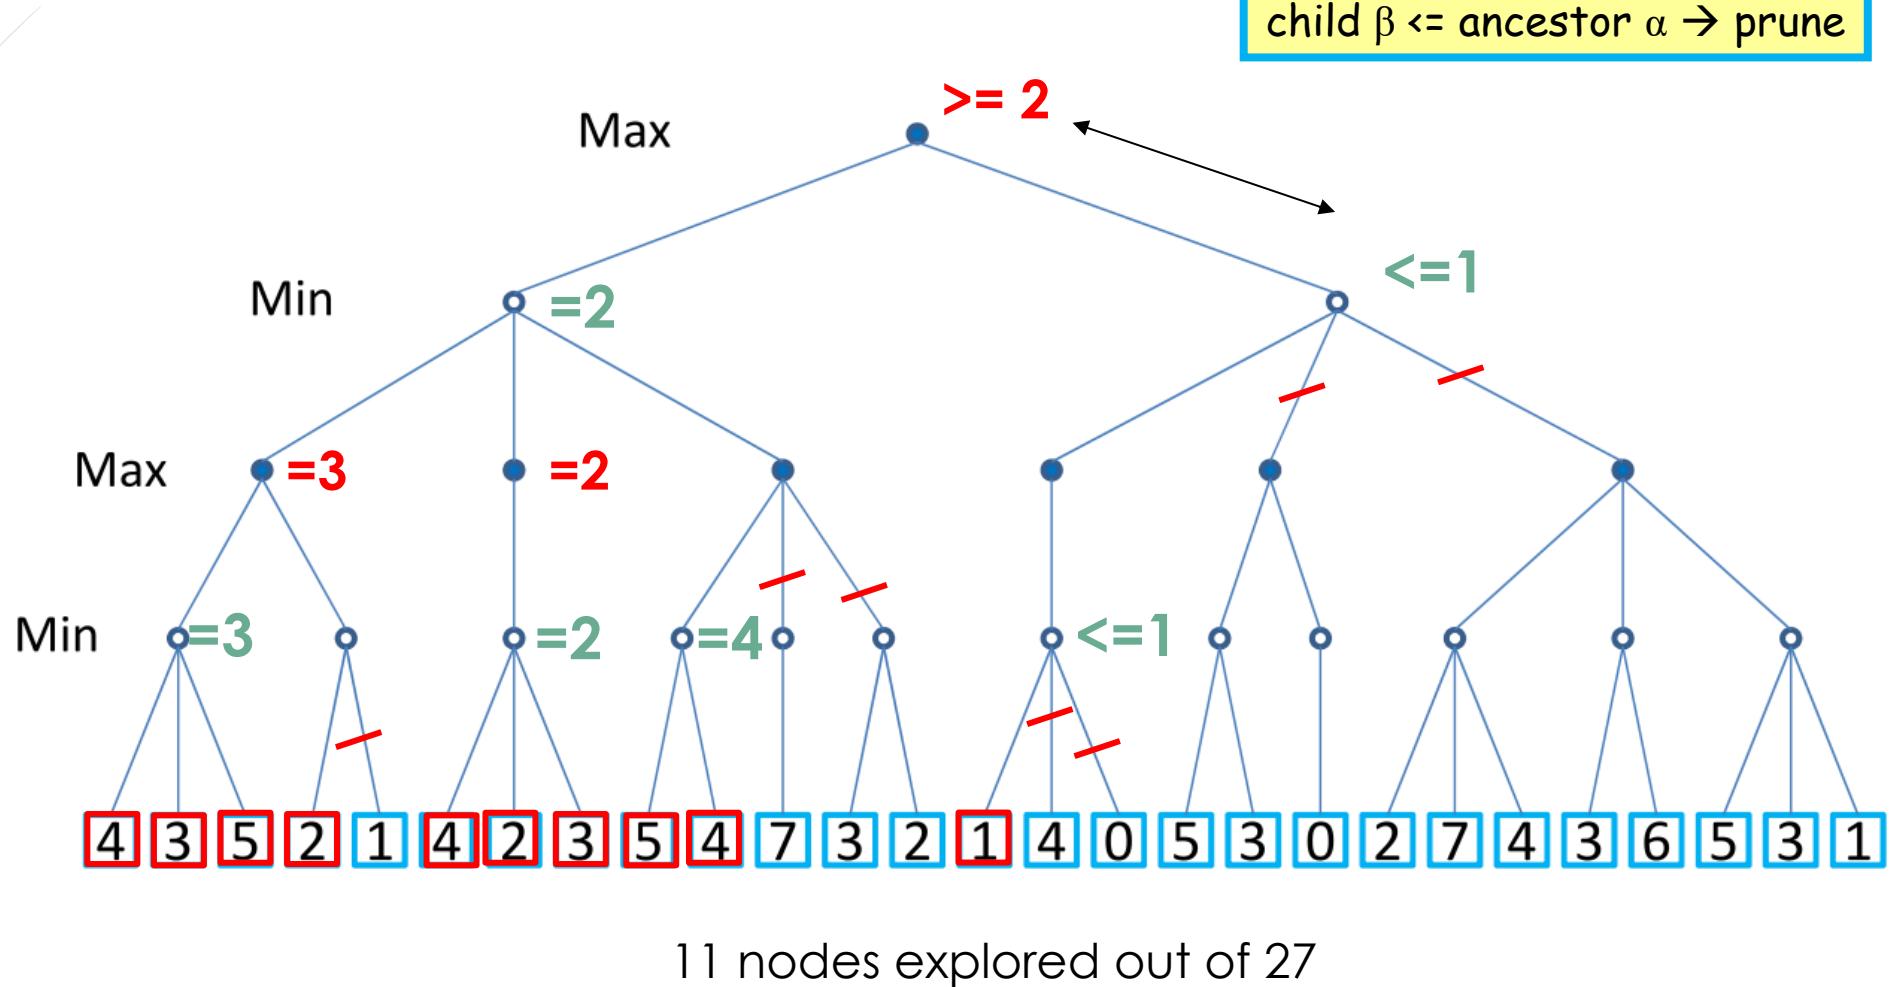


Alpha – Beta Pruning Example 3

child $\beta \leq \text{ancestor } \alpha \rightarrow \text{prune}$



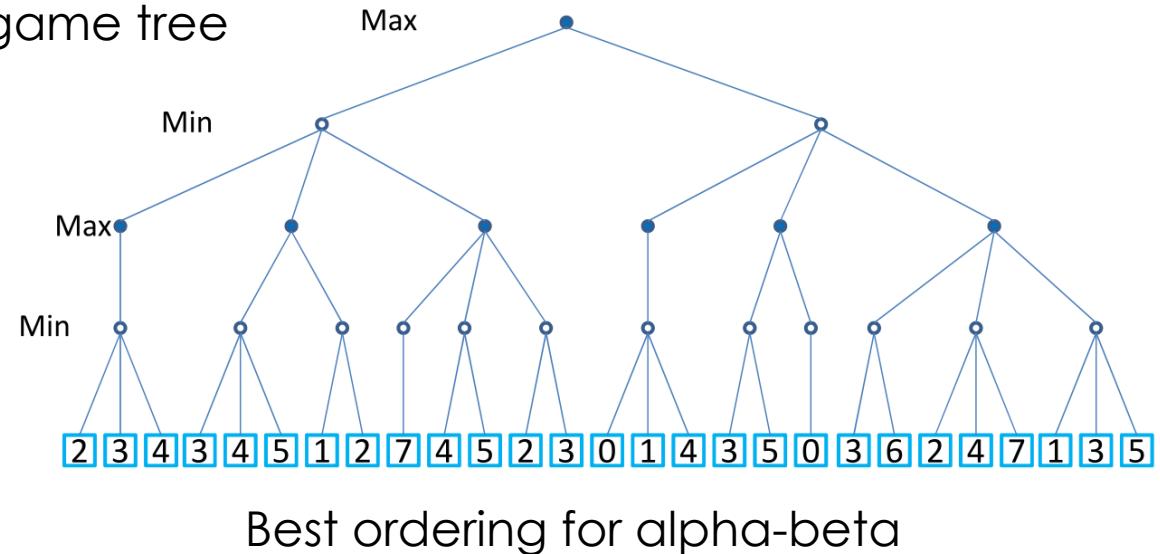
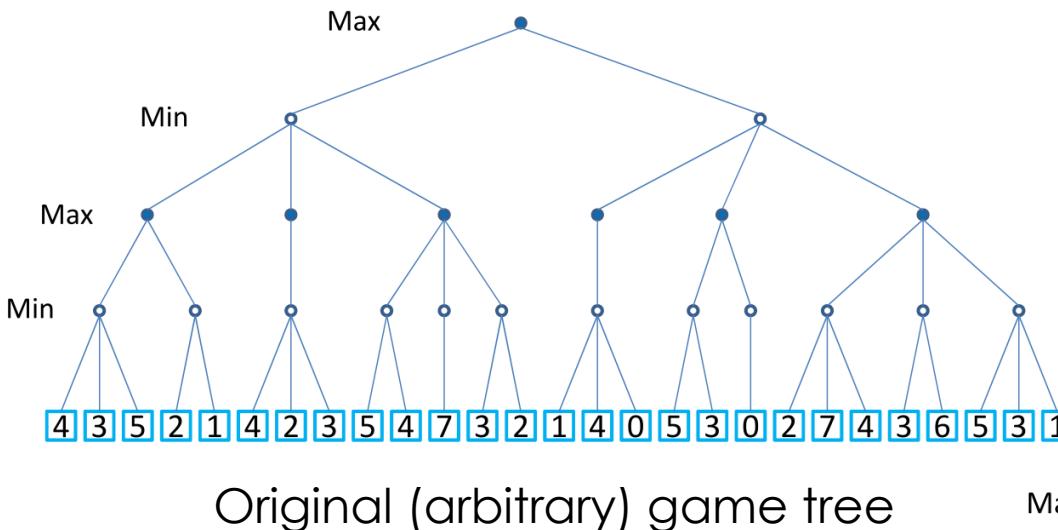
Alpha – Beta Pruning Example 3



Efficiency

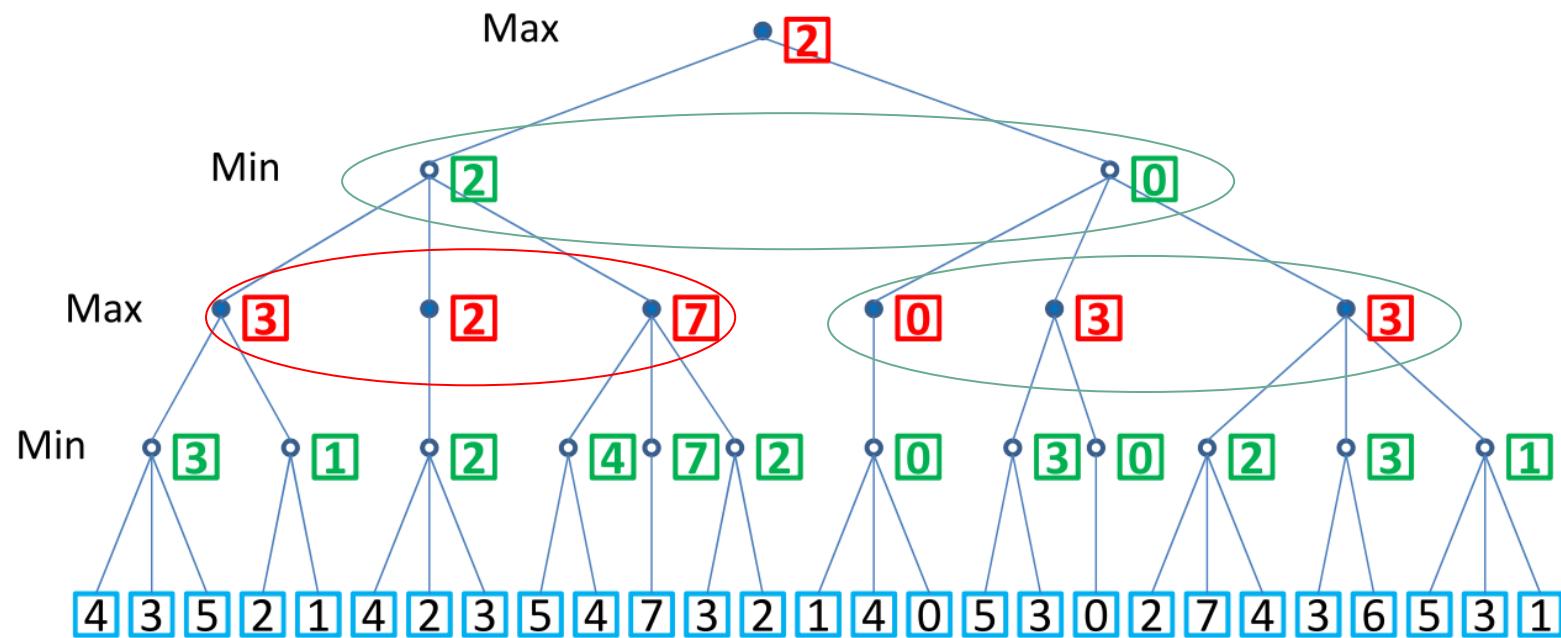
- ▶ Depends on the order the siblings
- ▶ In worst case:
 - ▶ alpha-beta provides no pruning
- ▶ In best case:
 - ▶ branching factor is reduced to its square root
 - ▶ so can search can go twice as deep with the same amount of computation

Alpha-Beta: Best ordering

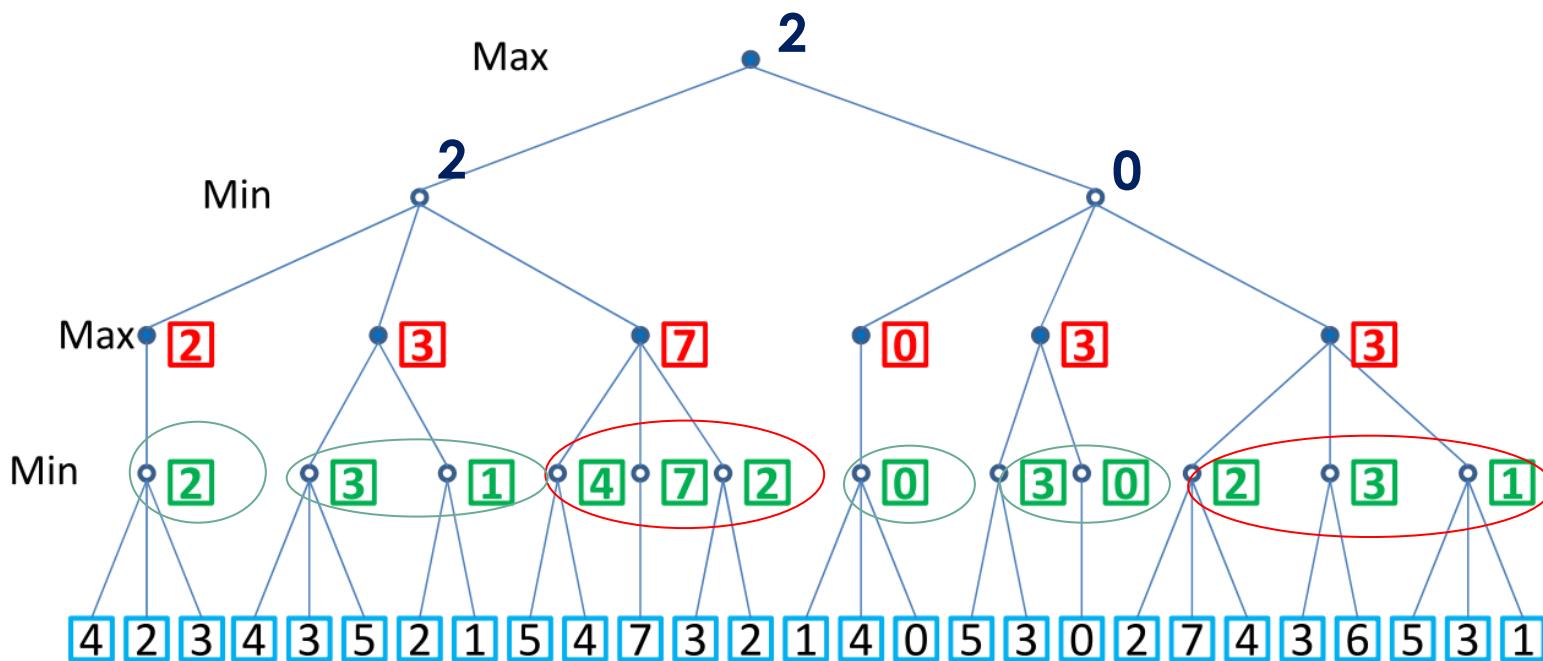


Alpha-Beta: Best ordering

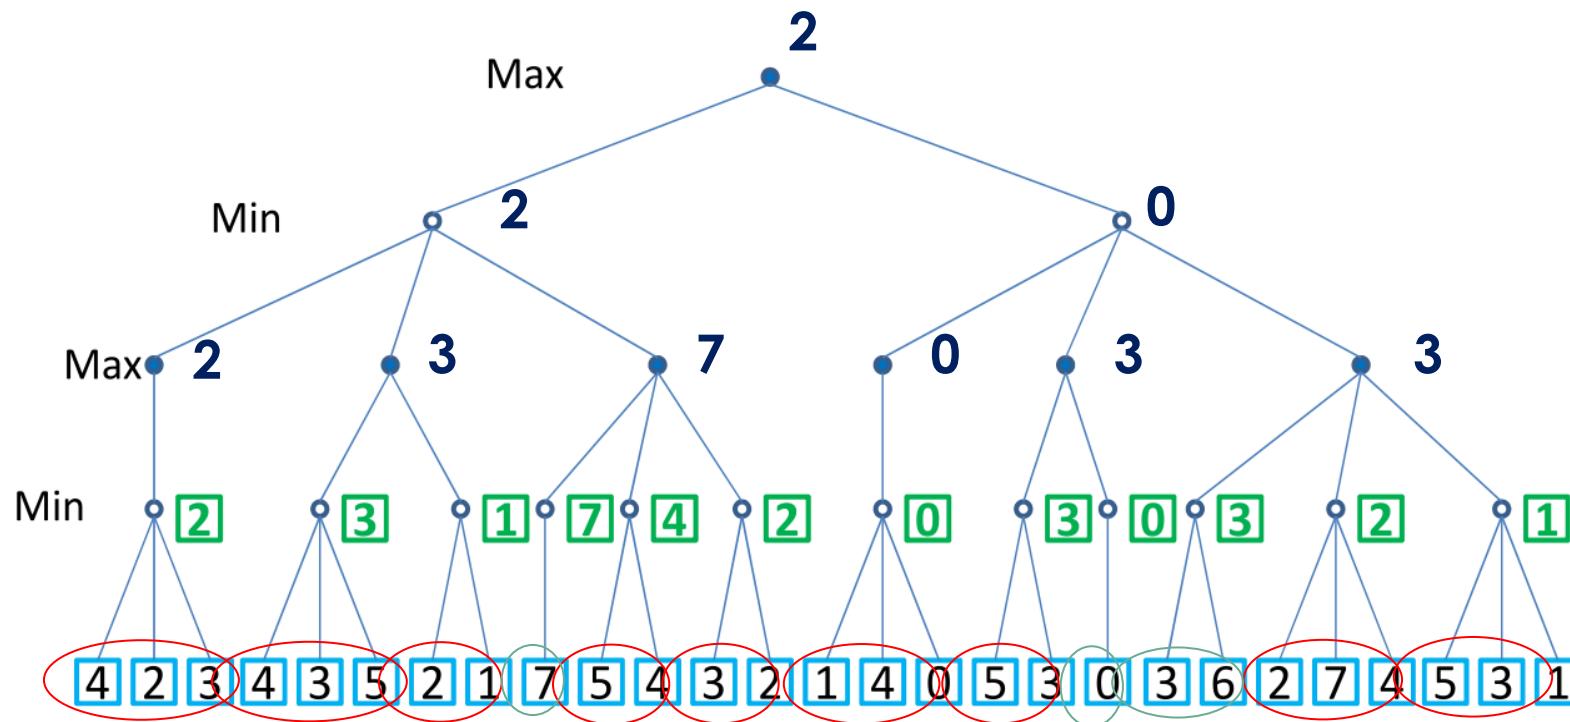
- ▶ best ordering:
 1. children of MIN : smallest node first
 2. children of MAX: largest node first



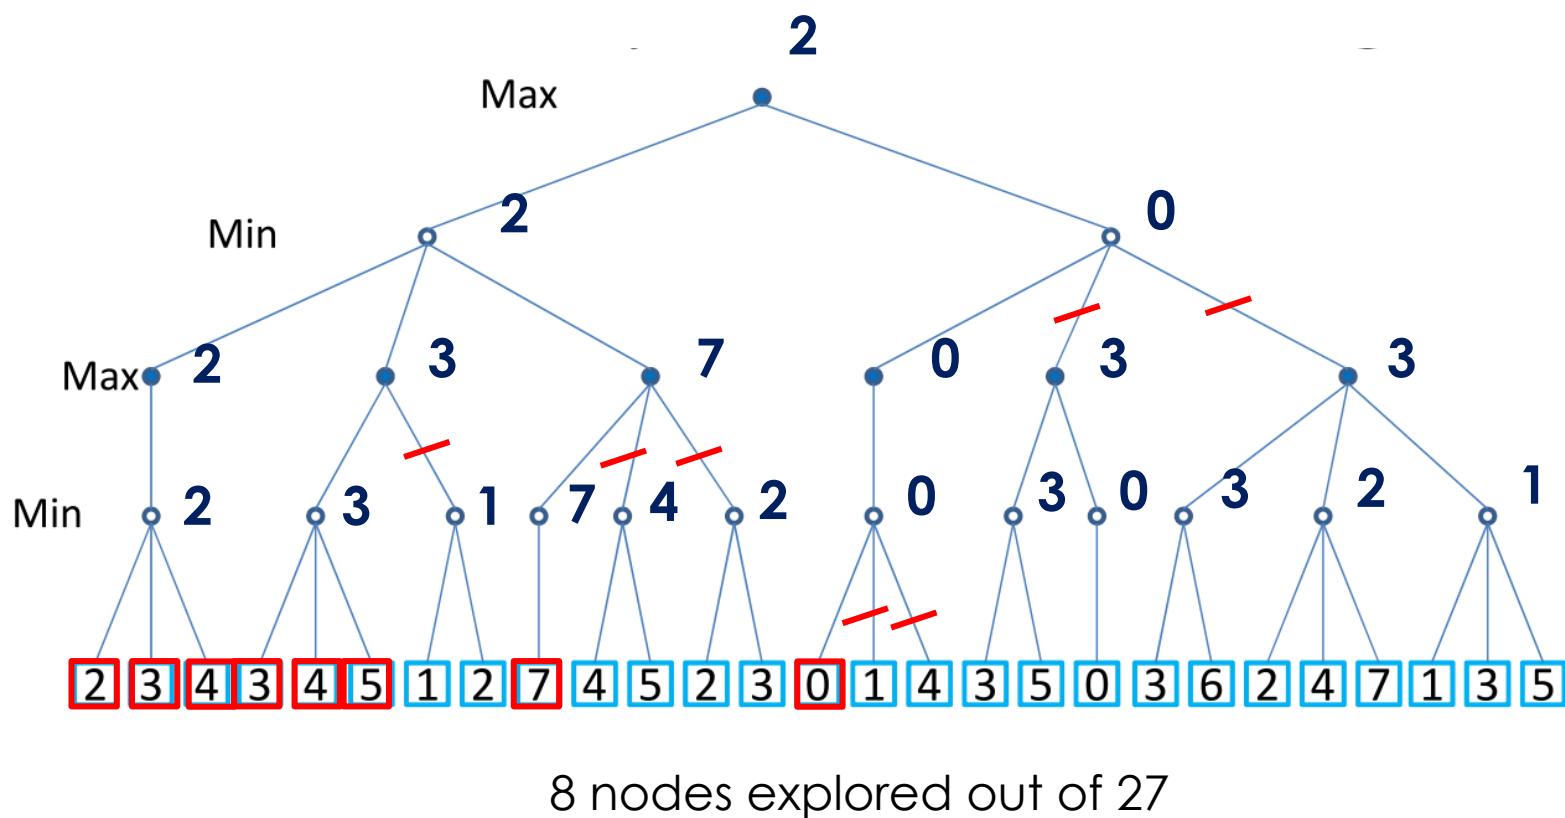
Alpha-Beta: Best ordering



Alpha-Beta: Best ordering

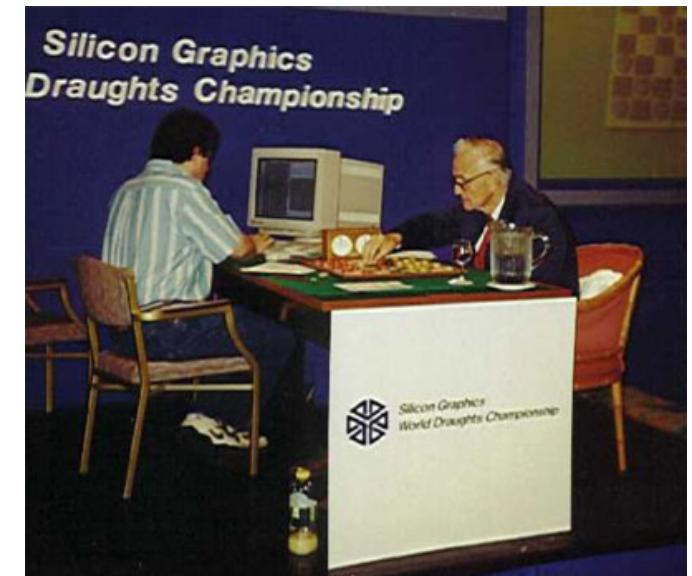


Alpha-Beta: Best ordering



Applications

- ▶ **1992 – 1994 Checkers: Tinsley VS Chinook**
- ▶ Marion Tinsley – World Champion for over 40 years
- ▶ Chinook – Developed by Jonathan Schaeffer, professor at University of Alberta
- ▶ In 1992, Tinsley beat Chinook in 4 games to 2, with 33 draws.
- ▶ In 1996: 6 draws
- ▶ In 2007, Schaeffer announced that checkers was solved, and anyone playing against Chinook would only be able to draw, never win.



<https://webdocs.cs.ualberta.ca/~chinook/play/>

Applications

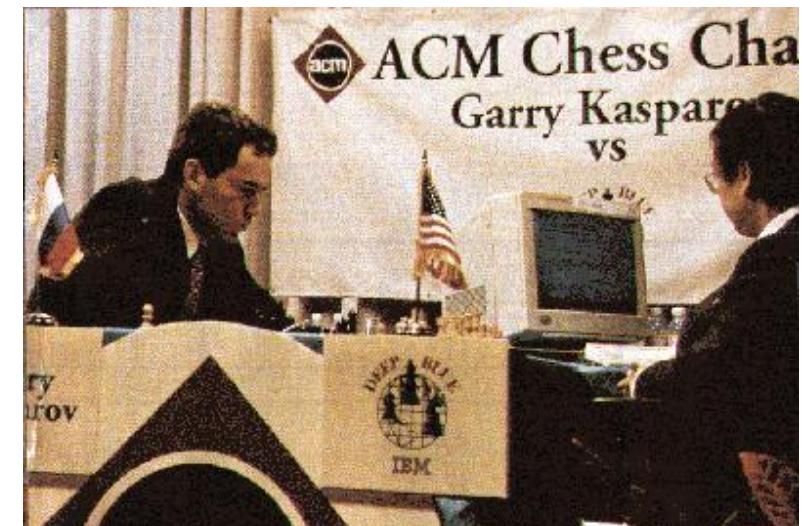
- ▶ **1997 Othello: Murakami VS Logistello**
- ▶ Takeshi Murakami – World Othello Champion
- ▶ Logistello – Developed by Michael Buro runs on a standard PC
- ▶ Logistello beat Murakami by 6 games to 0



<https://skatgame.net/mburo/log.html>
(including source code)

Applications

- ▶ **1997 Chess: Kasparov VS Deep Blue**
- ▶ Garry Kasparov – 50 billion neurons 2 positions/sec
- ▶ Deep Blue – 32 RISC (Reduced Instruction-Set Computer) processors + 256 VLSI (Very Large-Scale Integration) chess engines 200,000,000 pos/sec
- ▶ Deep Blue wins by 3 wins, 1 loss, and 2 draws



Applications

- ▶ **1997 Chess: Kasparov VS Deep Junior**
- ▶ Garry Kasparov – still 50 billion neurons 2 positions/sec
- ▶ Deep Junior – 8 CPU, 8 GB RAM, Win 2000 2,000,000 pos/sec Available at \$100
- ▶ Match ends in a 3/3 tie!



Applications

- ▶ **2016 Go: AlphaGo VS Lee Se-dol**
- ▶ GO was always considered a much harder game to automate than chess because of its very high branching factor (35 for chess vs 250 for Go!)
- ▶ In 2016, AlphaGo beat Lee Sedol in a five-game match of GO.
- ▶ In 2017 AlphaGo beat Ke Jie, the world No.1 ranked player at the time
- ▶ uses a Monte Carlo tree search algorithm to find its moves based on knowledge previously "learned" by deep learning



The End

